

Building Serverless MCP Servers on AWS

(pun intended)



Anton Aleksandrov
Principal Solutions Architect
AWS Serverless



MCP

Model Context Protocol

Because even the smartest models need additional context

Model Context Protocol

Because even the smartest models need additional context

Why do we need MCP?

Building a local
MCP Server

Building a remote
MCP Server on AWS

Agentic AI vs RAG

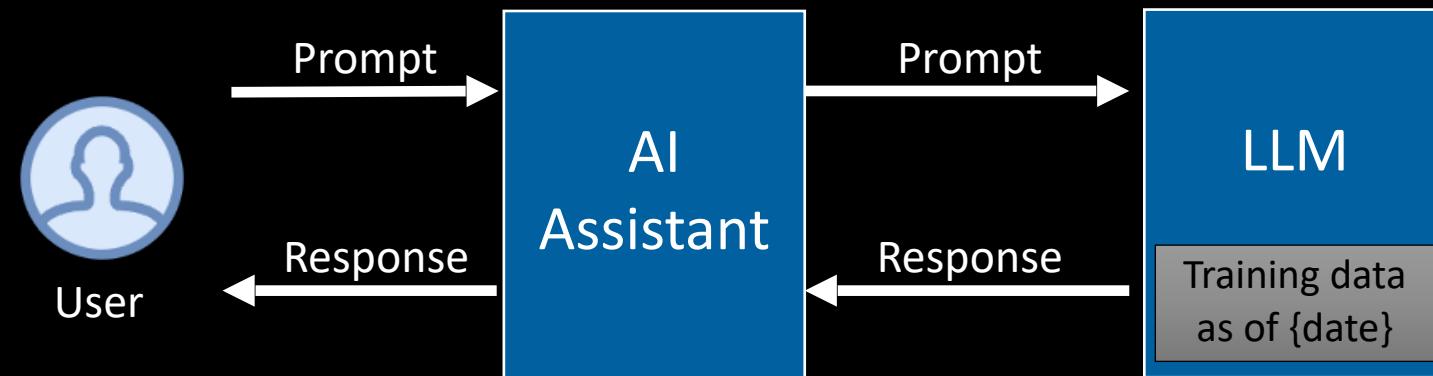
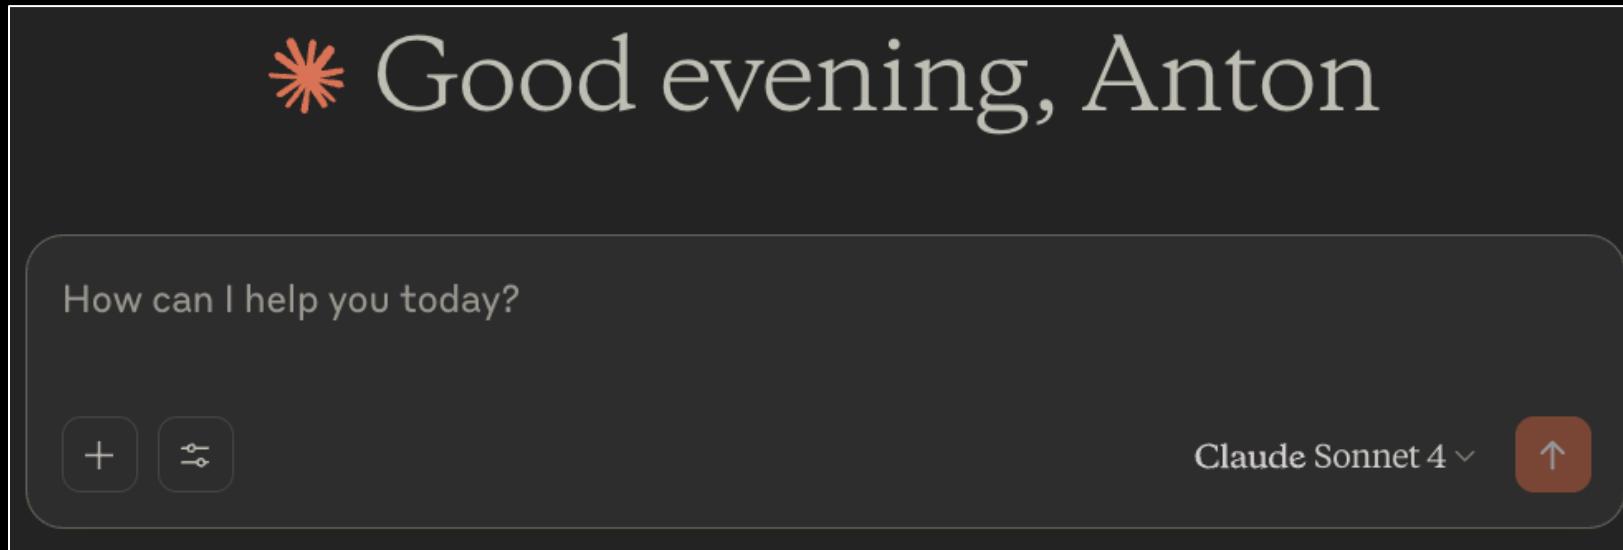
MCP building
blocks

Advanced patterns

The Problem

Why do we need MCP in the first place?

The Genesis



The Genesis



User

Write me a
poem about
serverless

"Events gently flow,
Your apps simply run.
No servers to tend—
It's cloud and it's fun"

- Claude 3.7 Sonnet, 2025



- DALL-E, 2025

Realizing the limitations

A

What's the weather right now?

Realizing the limitations

A

What's the weather right now?

I don't have access to current weather data or your location, so I cannot tell you the current weather conditions. To get accurate up-to-date weather information, I'd recommend checking a weather app on your phone, browsing to a local weather service website, or simply looking out the window for immediate local conditions.

Realizing the limitations

A Get me sales data for my department last week

I don't have access to your company's sales data or internal systems. I'm an AI assistant without the ability to connect to your organization's databases, CRM systems, or sales platforms.

To get your department's sales data from last week, you would need to:

- Check your company's sales dashboard or CRM system (like Salesforce, HubSpot, etc.)
- Contact your sales operations team or data analytics department
- Access your business intelligence tools if your company uses them
- Review reports from your sales management system

If you have specific sales data that you'd like help analyzing or formatting once you obtain it, I'd be happy to assist with that.

Realizing the limitations

A fix bugs in my code

I'd be happy to help you fix bugs in your code! However, I don't see any code in your message.

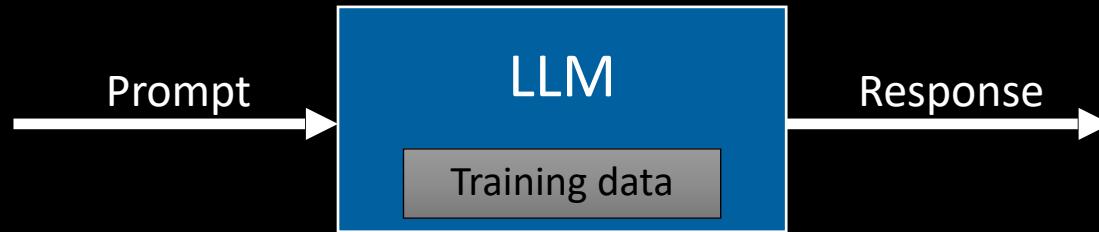
Could you please share the code that you're having issues with?

When you provide the code, it would be helpful if you could also include:

- What programming language you're using
- What the code is supposed to do
- What specific errors or unexpected behavior you're experiencing
- Any error messages you're seeing

Once you share the code, I'll take a look and help identify and fix any bugs.

Realizing the limitations



**LLMs don't know
what they don't
know.**

**LLMs can generate
content, but they
cannot take actions.**

Injecting the context

how much is 1+2+3?

1 + 2 + 3 = 6 ✓

Injecting the context

how much is 1+2+3?

1 + 2 + 3 = 6 ✓

assume 1 = apple, 2=peach, and 3 = grape. How
much is 1+2+3?

If 1 = apple, 2 = peach, and 3 = grape, then:

1 + 2 + 3 = apple + peach + grape

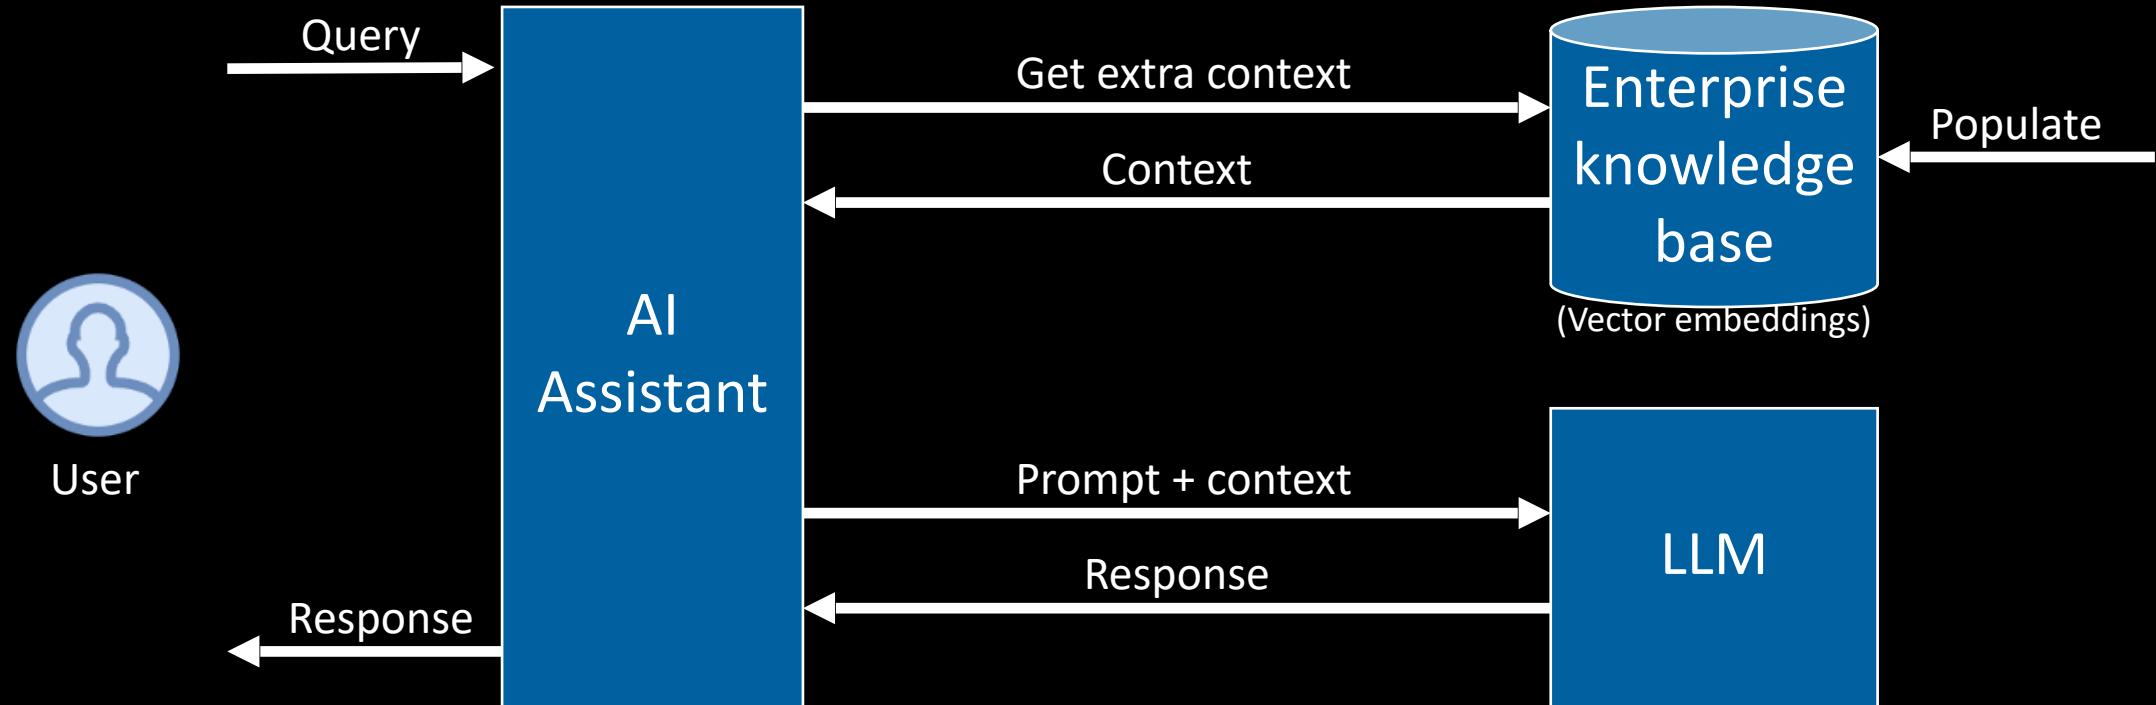
So the answer is a delicious fruit salad 😊🍎🍑🍇

Realizing the limitations

So... can we inject the context
into prompts to make LLMs
smarter!?

Absolutely!!!!

A simple RAG workflow



Retrieval-Augmented Generation (RAG)

A

Need final sales data for my department for the last week

Retrieval-Augmented Generation (RAG)

A

Need final sales data for my department for the last week

Based on your department's sales data for the last week, here's a summary of the final figures:

Total Sales: \$247,835.92

Units Sold: 1,482

Average Transaction Value: \$167.23

Conversion Rate: 24.8%

Top performing products:

- Premium Package: \$82,450 (33.3% of total sales)
- Standard Bundle: \$64,290 (25.9% of total sales)
- Basic Plan: \$43,870 (17.7% of total sales)

The department saw a 12.3% increase compared to the previous week, largely driven by the new promotion on Premium Packages. The eastern region continues to be the strongest performer, accounting for 41% of total sales.

Would you like more specific details about any particular aspect of these sales figures?

Adding real-time context dynamically

What's the weather forecast for tomorrow in my location?

Adding real-time context dynamically

What's the weather forecast for tomorrow in my location?

Tomorrow, August 15th, 2025, Expect a **hot, partly sunny day** in Austin, Texas with temperatures climbing into the **mid-90s °F (around 35 °C)** in the early afternoon. A **chance of thunderstorms** may develop late in the afternoon—especially around 5 PM—so it's wise to keep an eye on the skies. The evening will bring some **relief** with skies clearing and temperatures cooling into the **low 80s °F (around 27 °C)** by late evening.

Sources

Adding real-time context dynamically

What's the weather forecast for tomorrow in my location?

Tomorrow, August 15th, 2025, Expect a **hot, partly sunny day** in Austin, Texas with temperatures climbing into the **mid-90s °F (around 35 °C)** in the early afternoon. A **chance of thunderstorms** may develop late in the afternoon—especially around 5 PM—so it's wise to keep an eye on the skies. The evening will bring some **relief** with skies clearing and temperatures cooling into the **low 80s °F (around 27 °C)** by late evening.

Sources

 National Weather Service

7-Day Forecast 30.27N 97.74W - National Weather Service

Monday Night: Partly cloudy, with a low around 76. South southeast wind around. Low: 76 °F. Partly...

 AccuWeather

Austin, TX Weather Forecast - AccuWeather

Looking Ahead. A thunderstorm tomorrow afternoon; storms will be slow moving and bring heavy rain....

 The Weather Channel

10-Day Weather Forecast for Austin, TX

10 Day Weather-Austin, TX. As of 9:16 pm CDT. Tonight. --/77°. 7%. Night. 77°. 7%. S 8 mph. Partly...

 KEYE

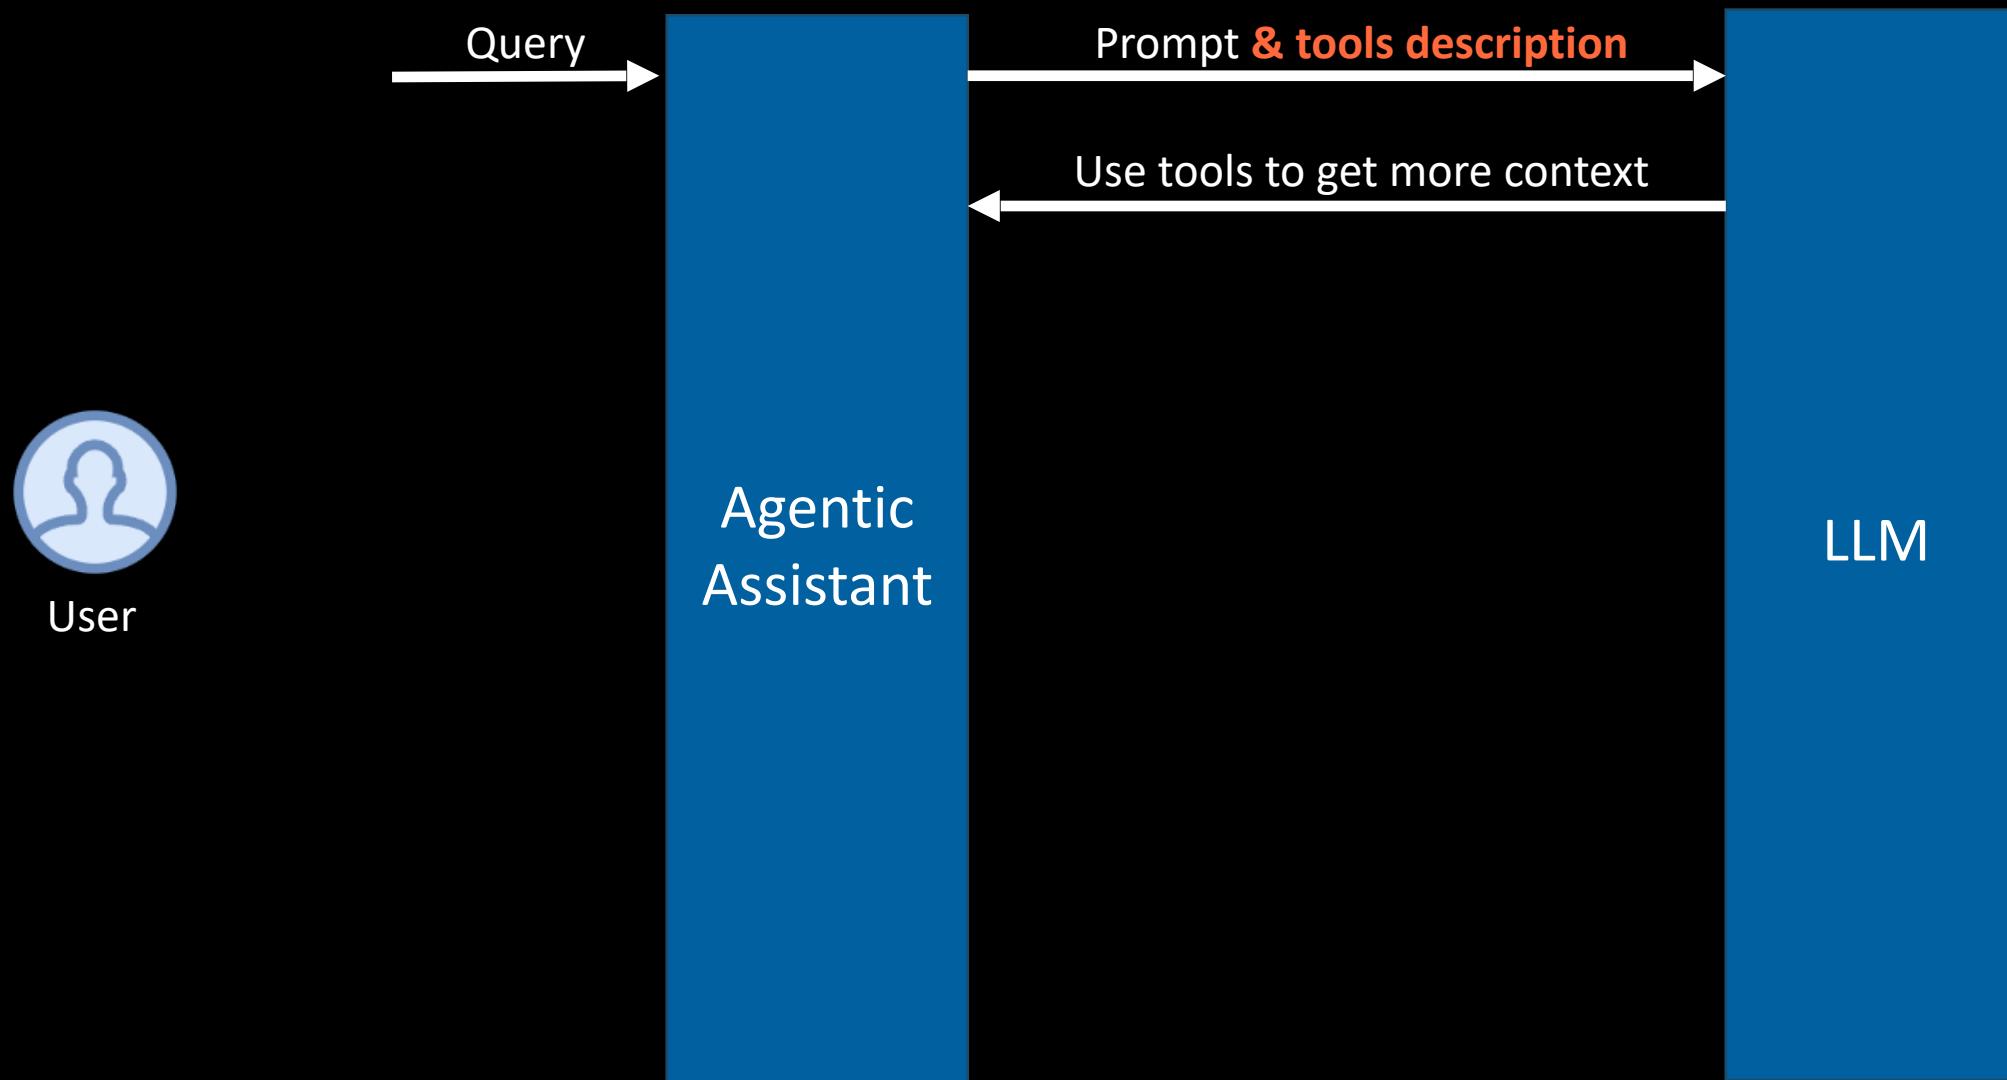
Austin Weather | News, Weather, Sports, Breaking News

Monday and Tuesday, we may see the sea breeze a little more successful, possibly bringing spotty rain a...

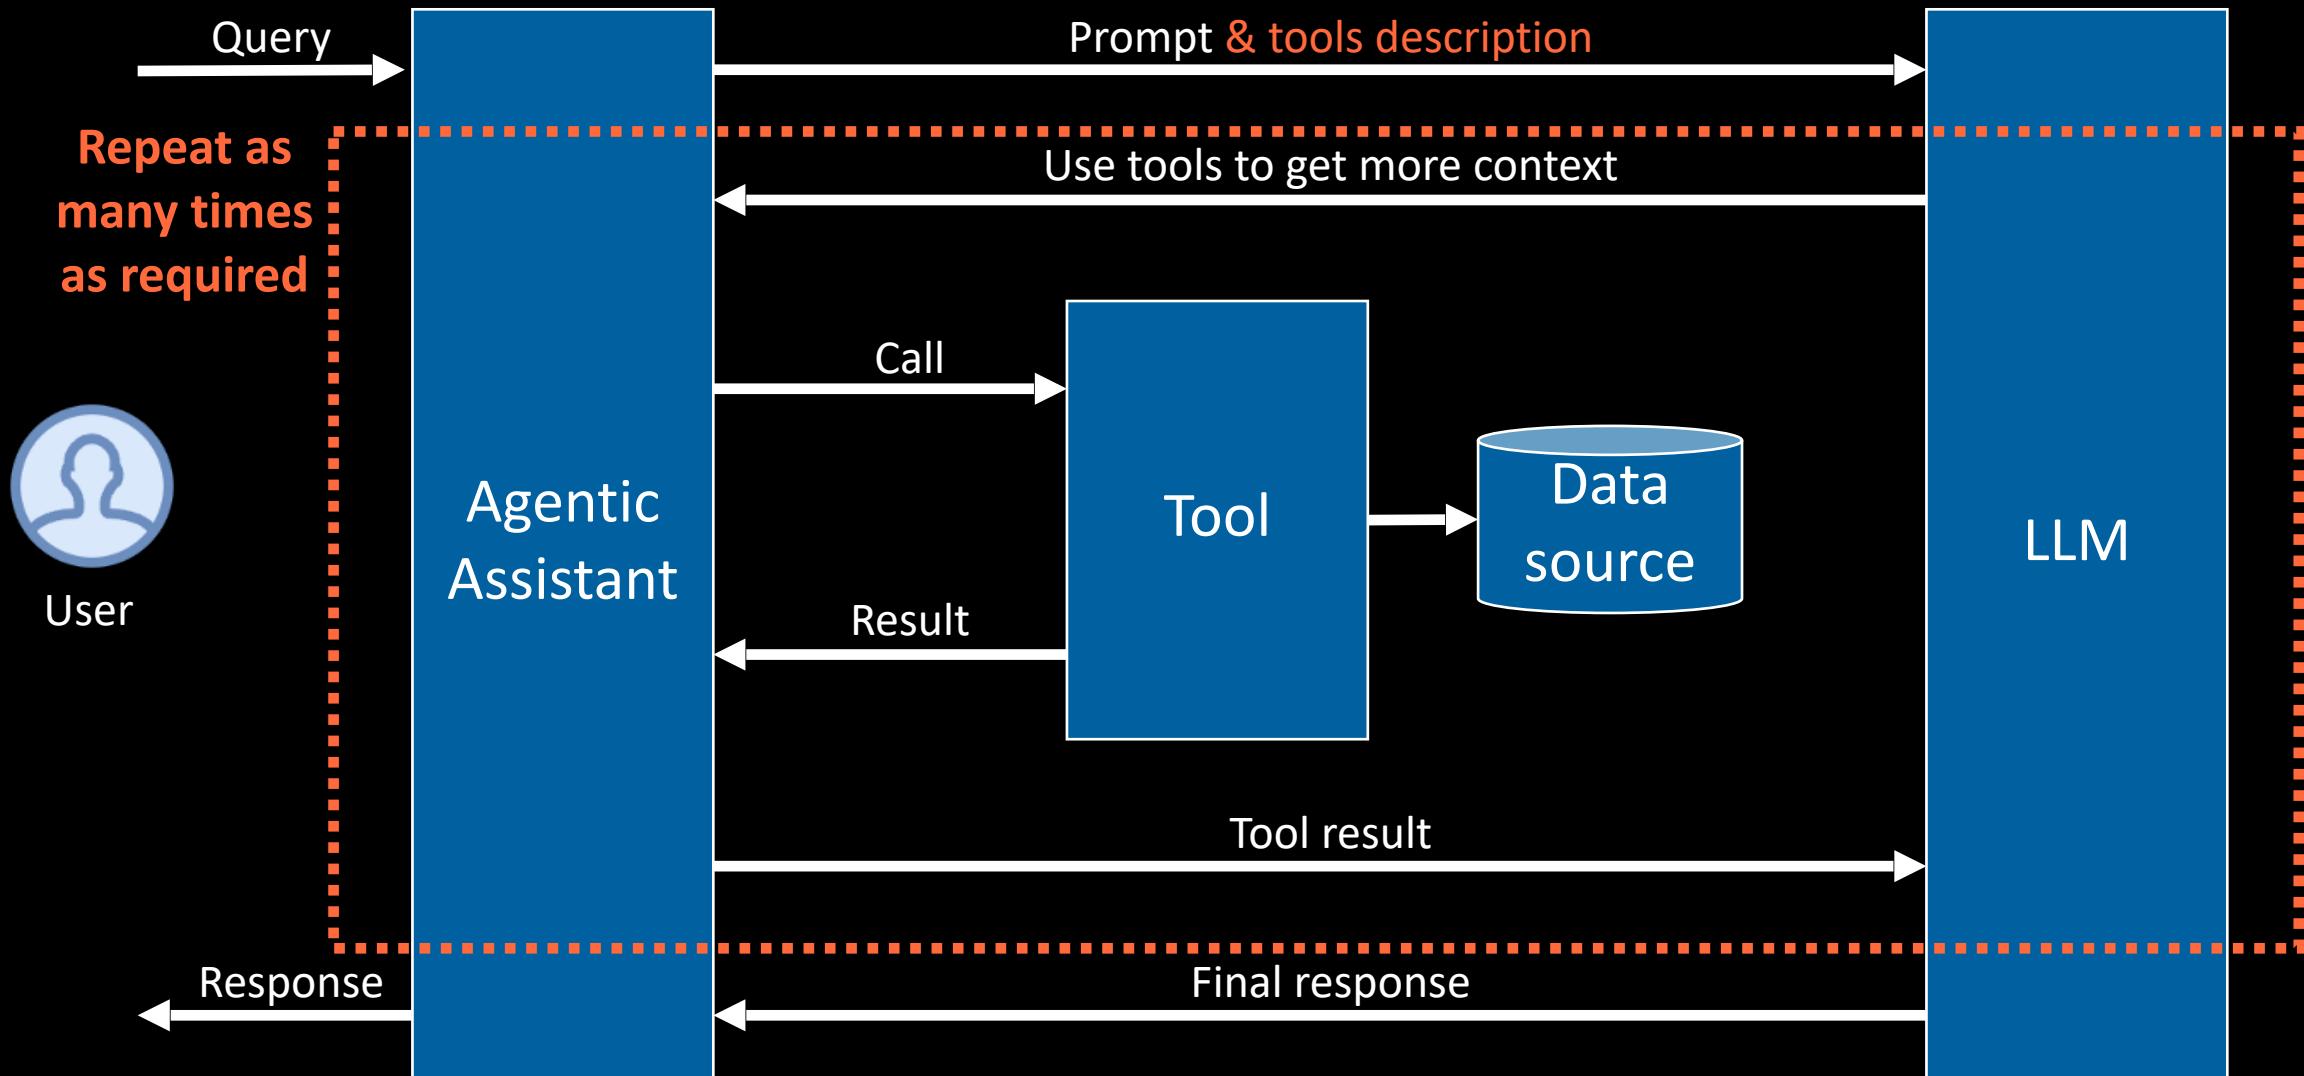
 statesman.com

Austin bakes in summer heat this week, but why

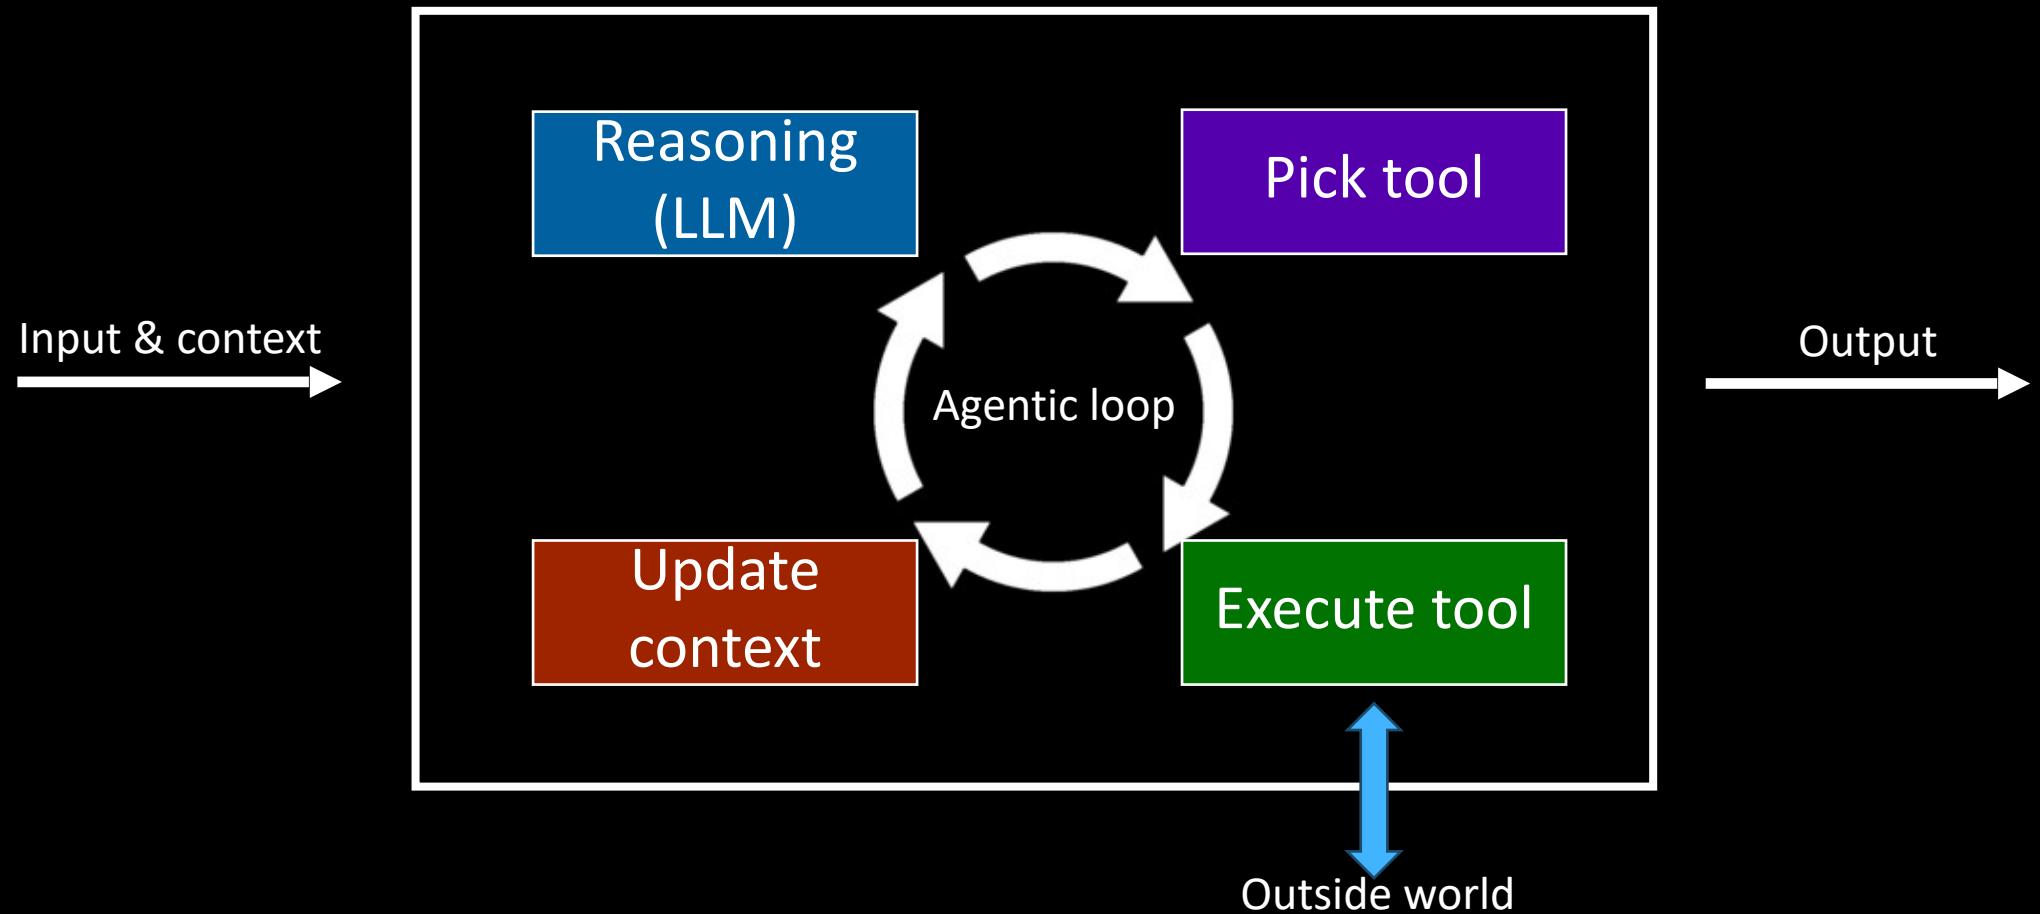
A simple agentic workflow



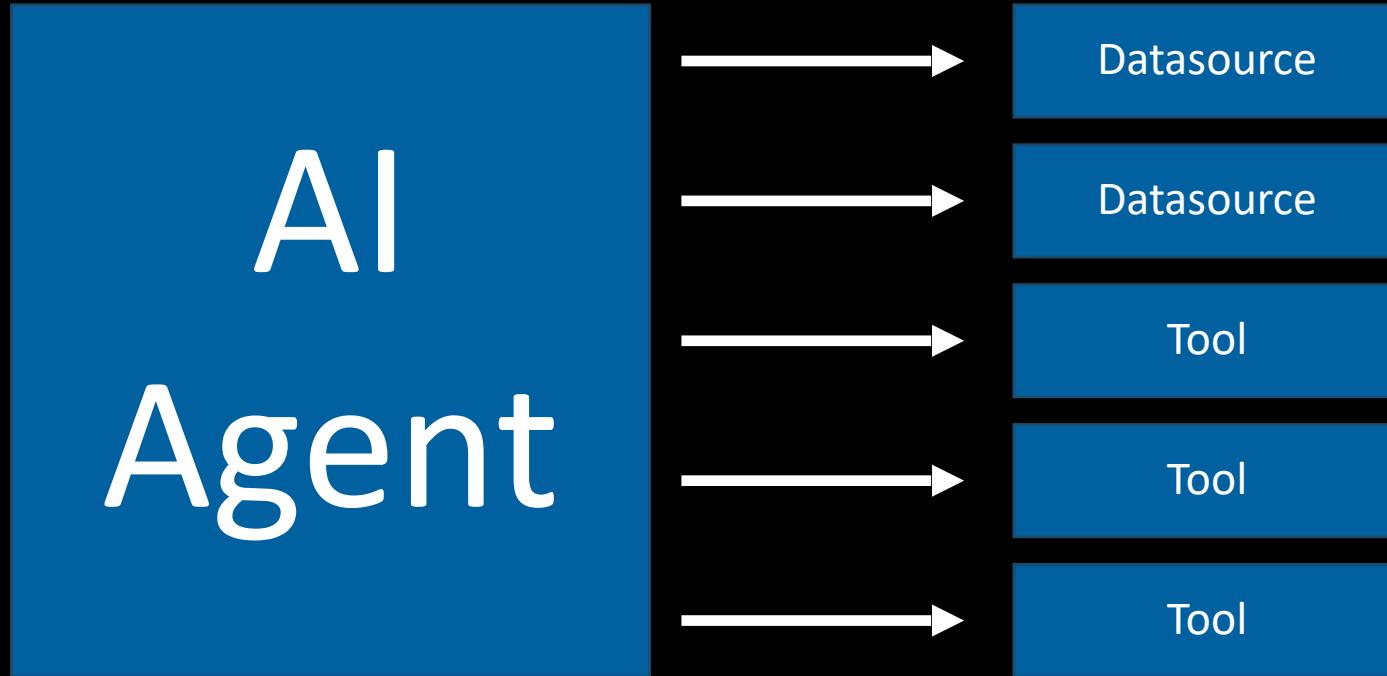
A simple agentic workflow



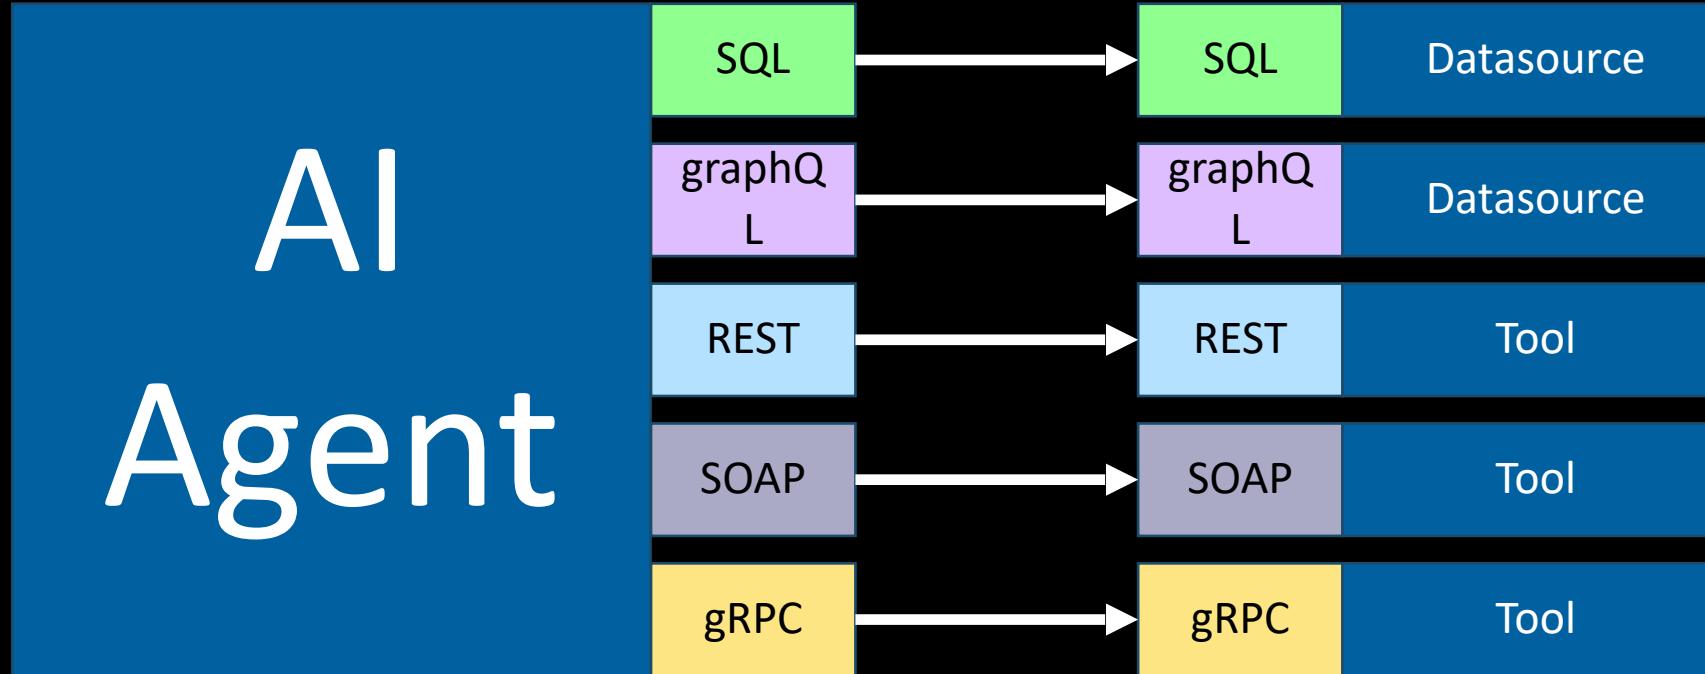
Agentic AI system in a nutshell



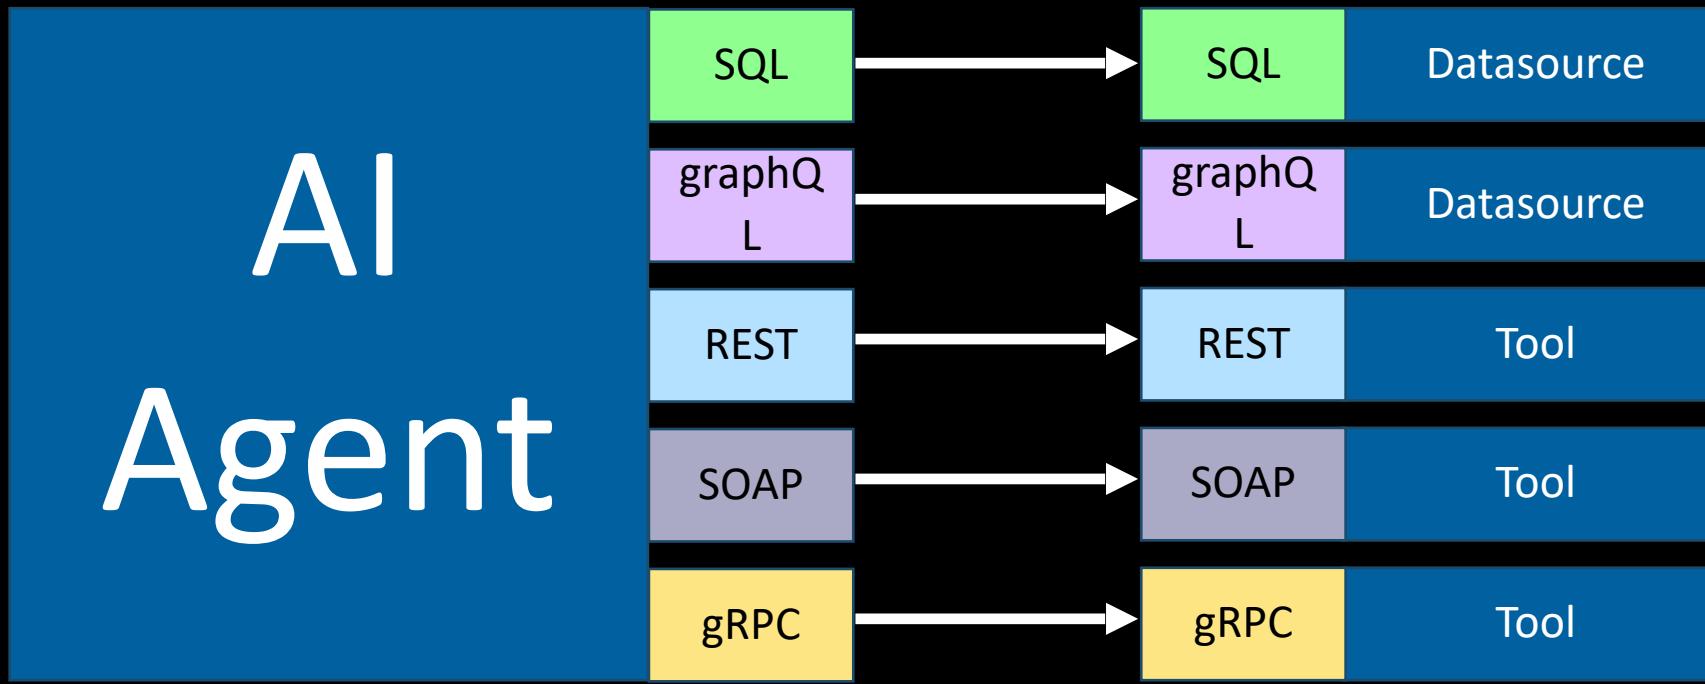
Communicating with multiple tools and resources



Communicating with multiple tools and resources



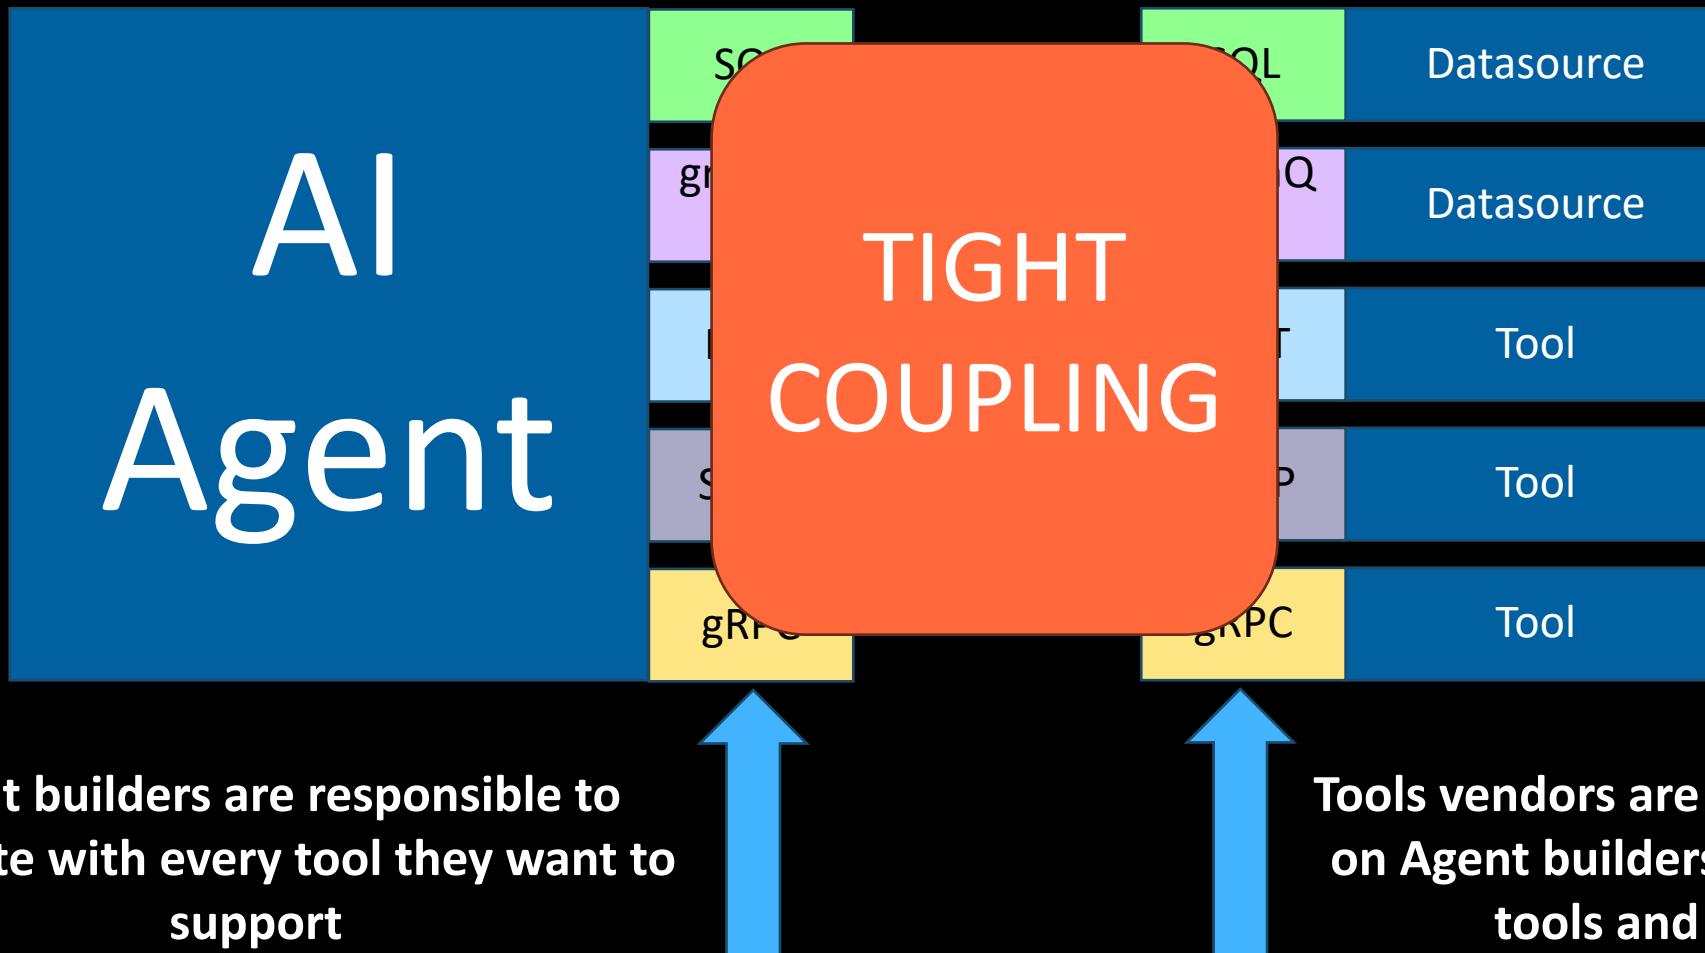
Communicating with multiple tools and resources



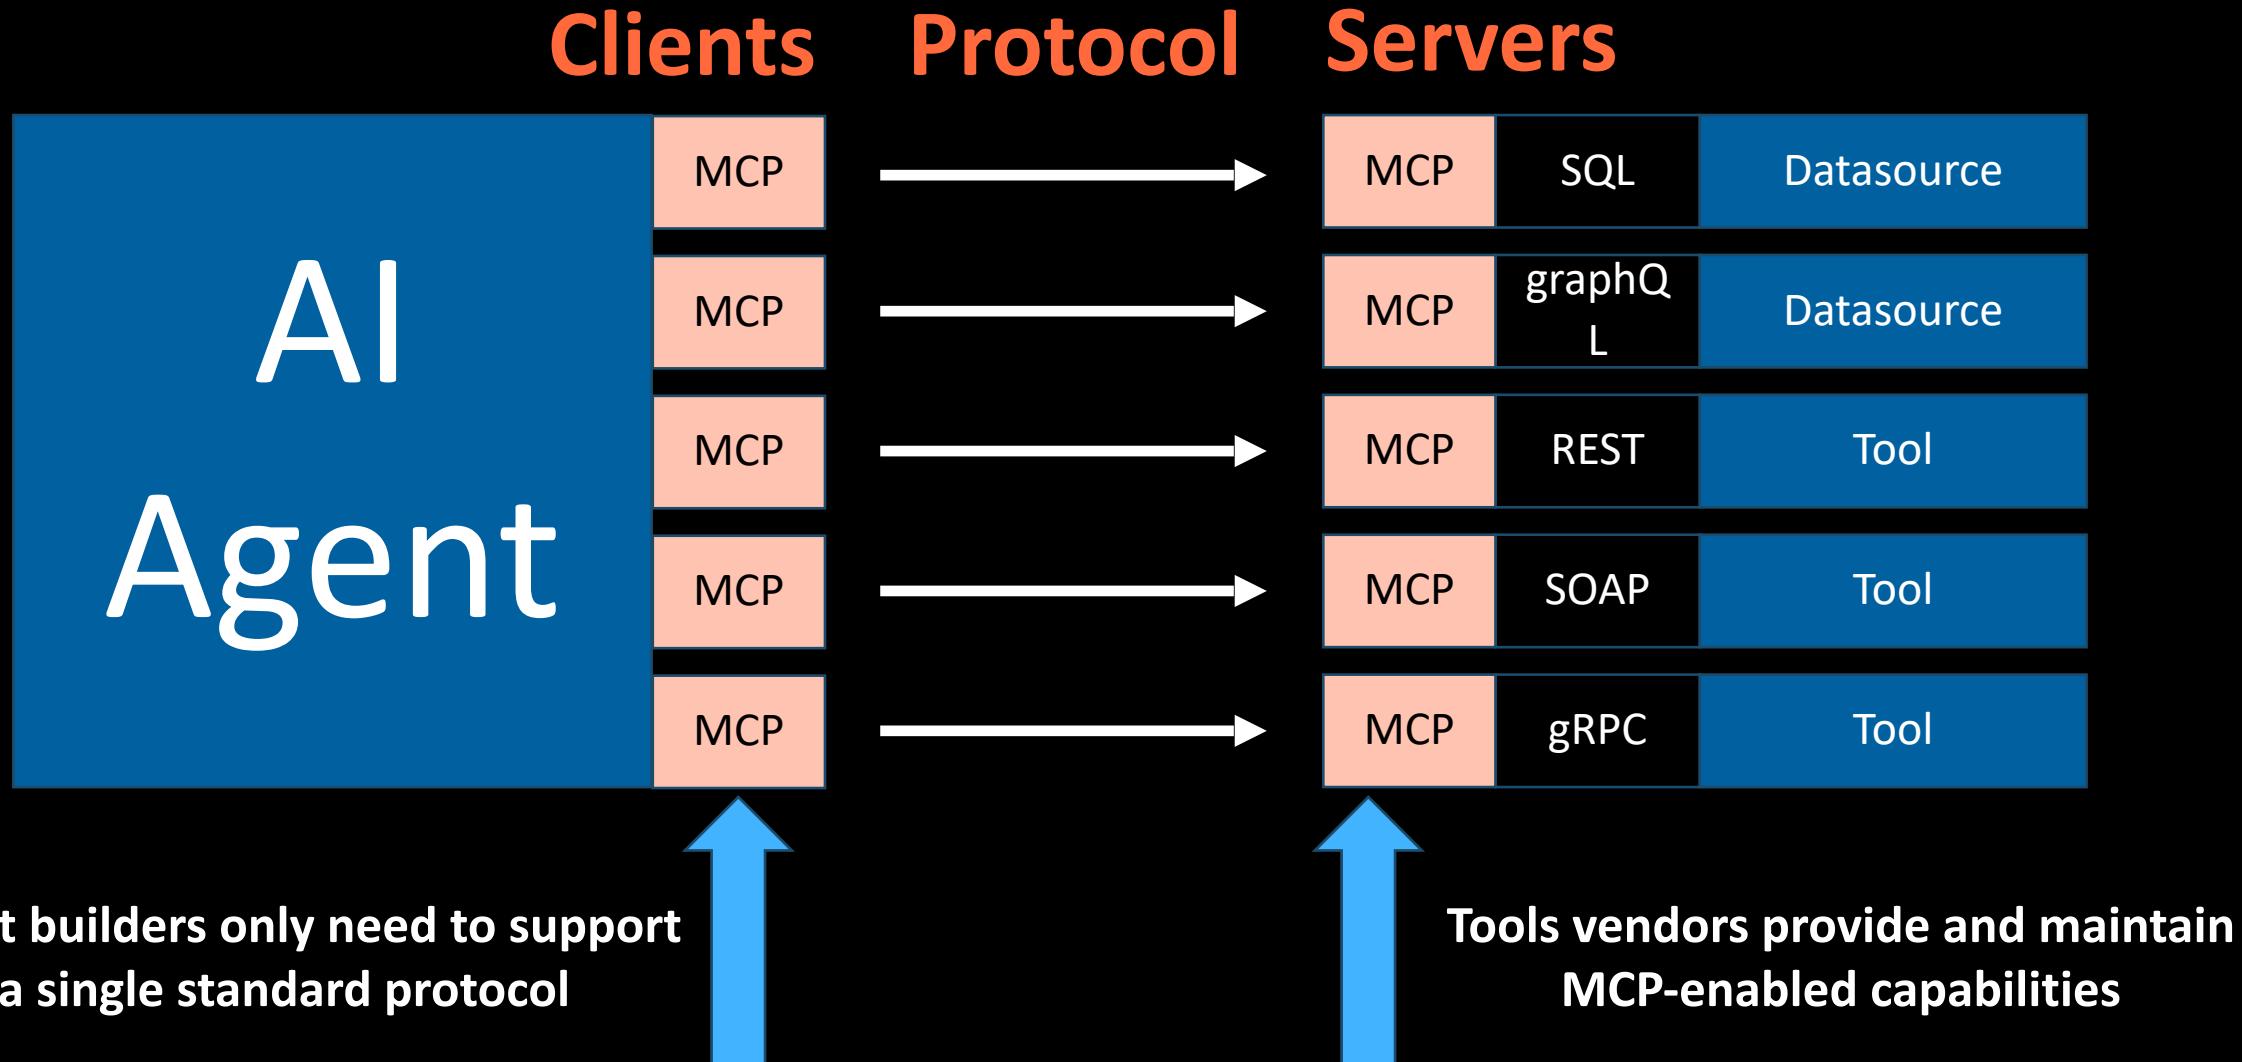
Agent builders are responsible to integrate with every tool they want to support

Tools vendors are highly dependent on Agent builders to support their tools and protocols

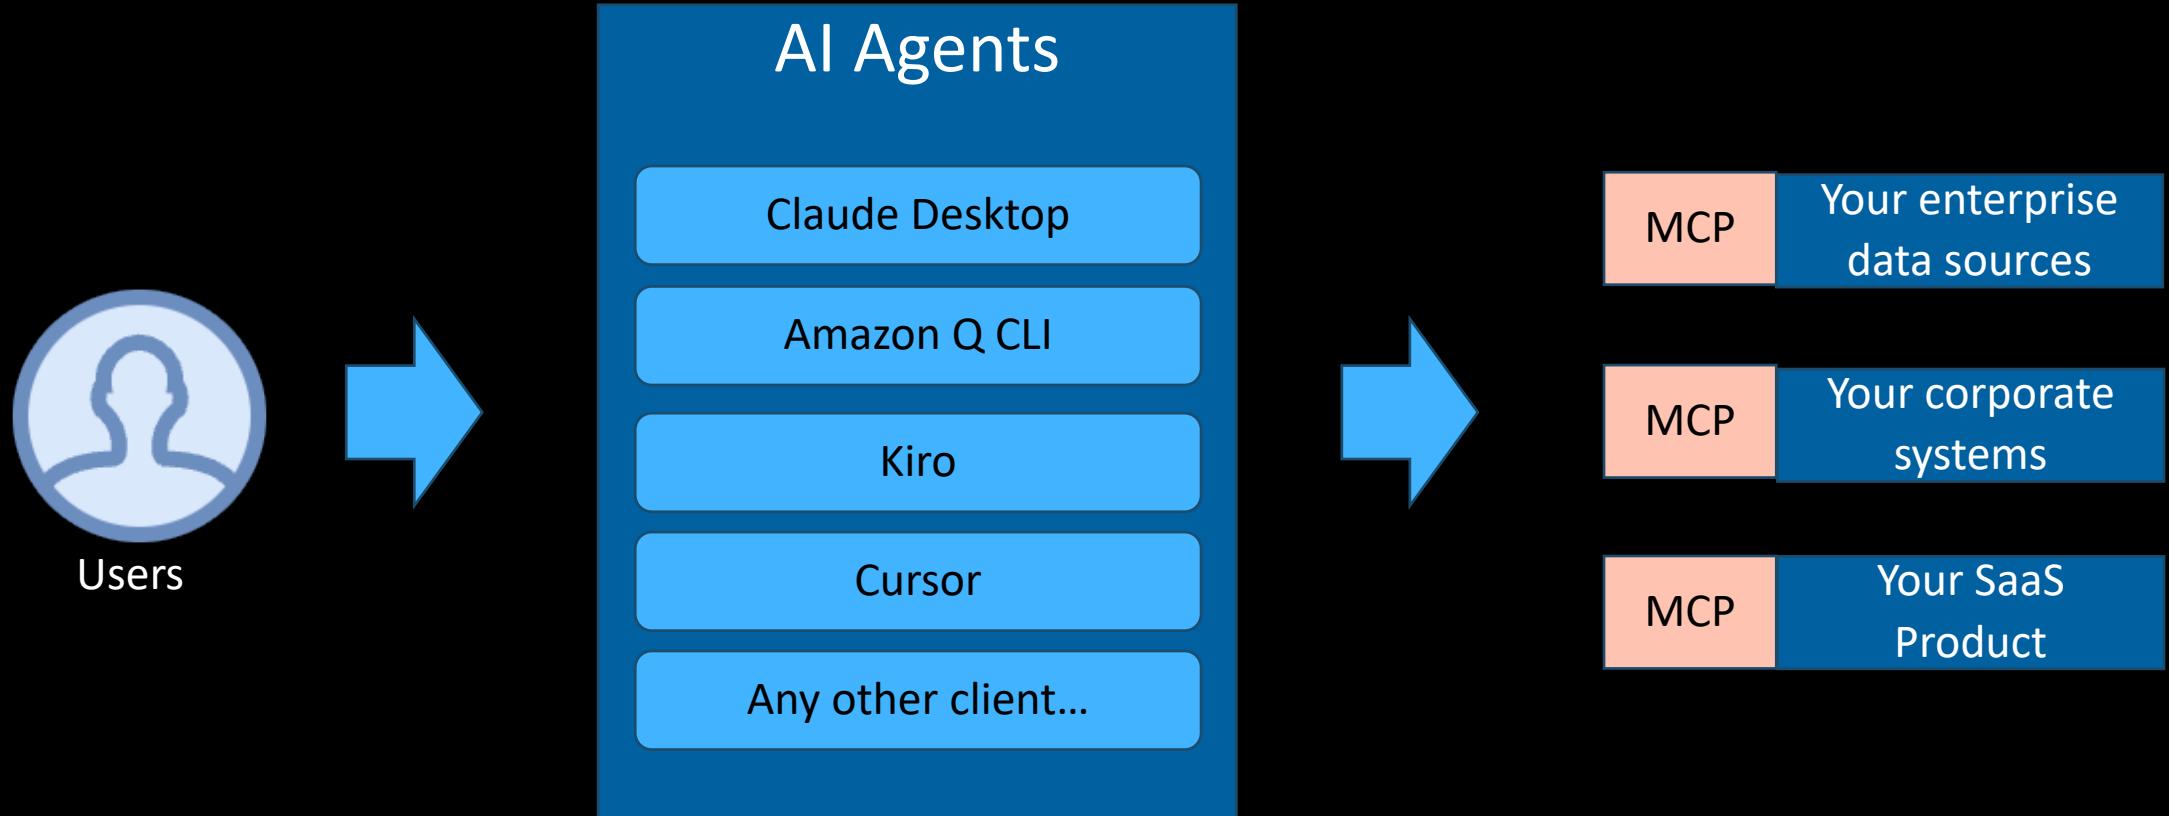
Communicating with multiple tools and resources



Communicating with multiple tools and resources



Communicating with multiple tools and resources



The Building Blocks

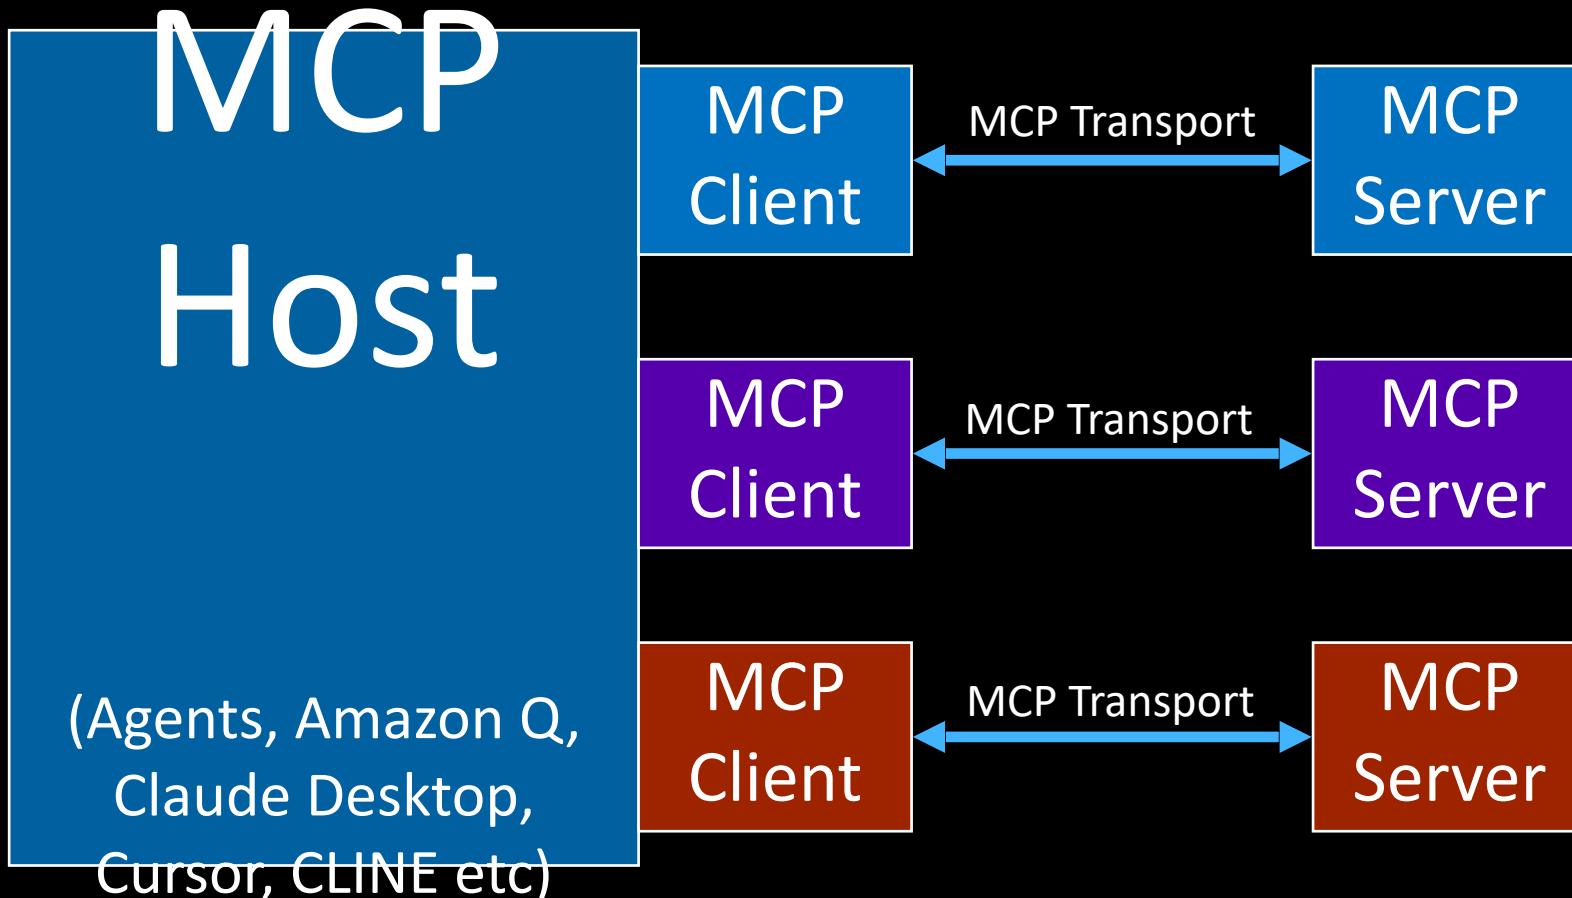
What are the MCP components

Model Context Protocol - Entities

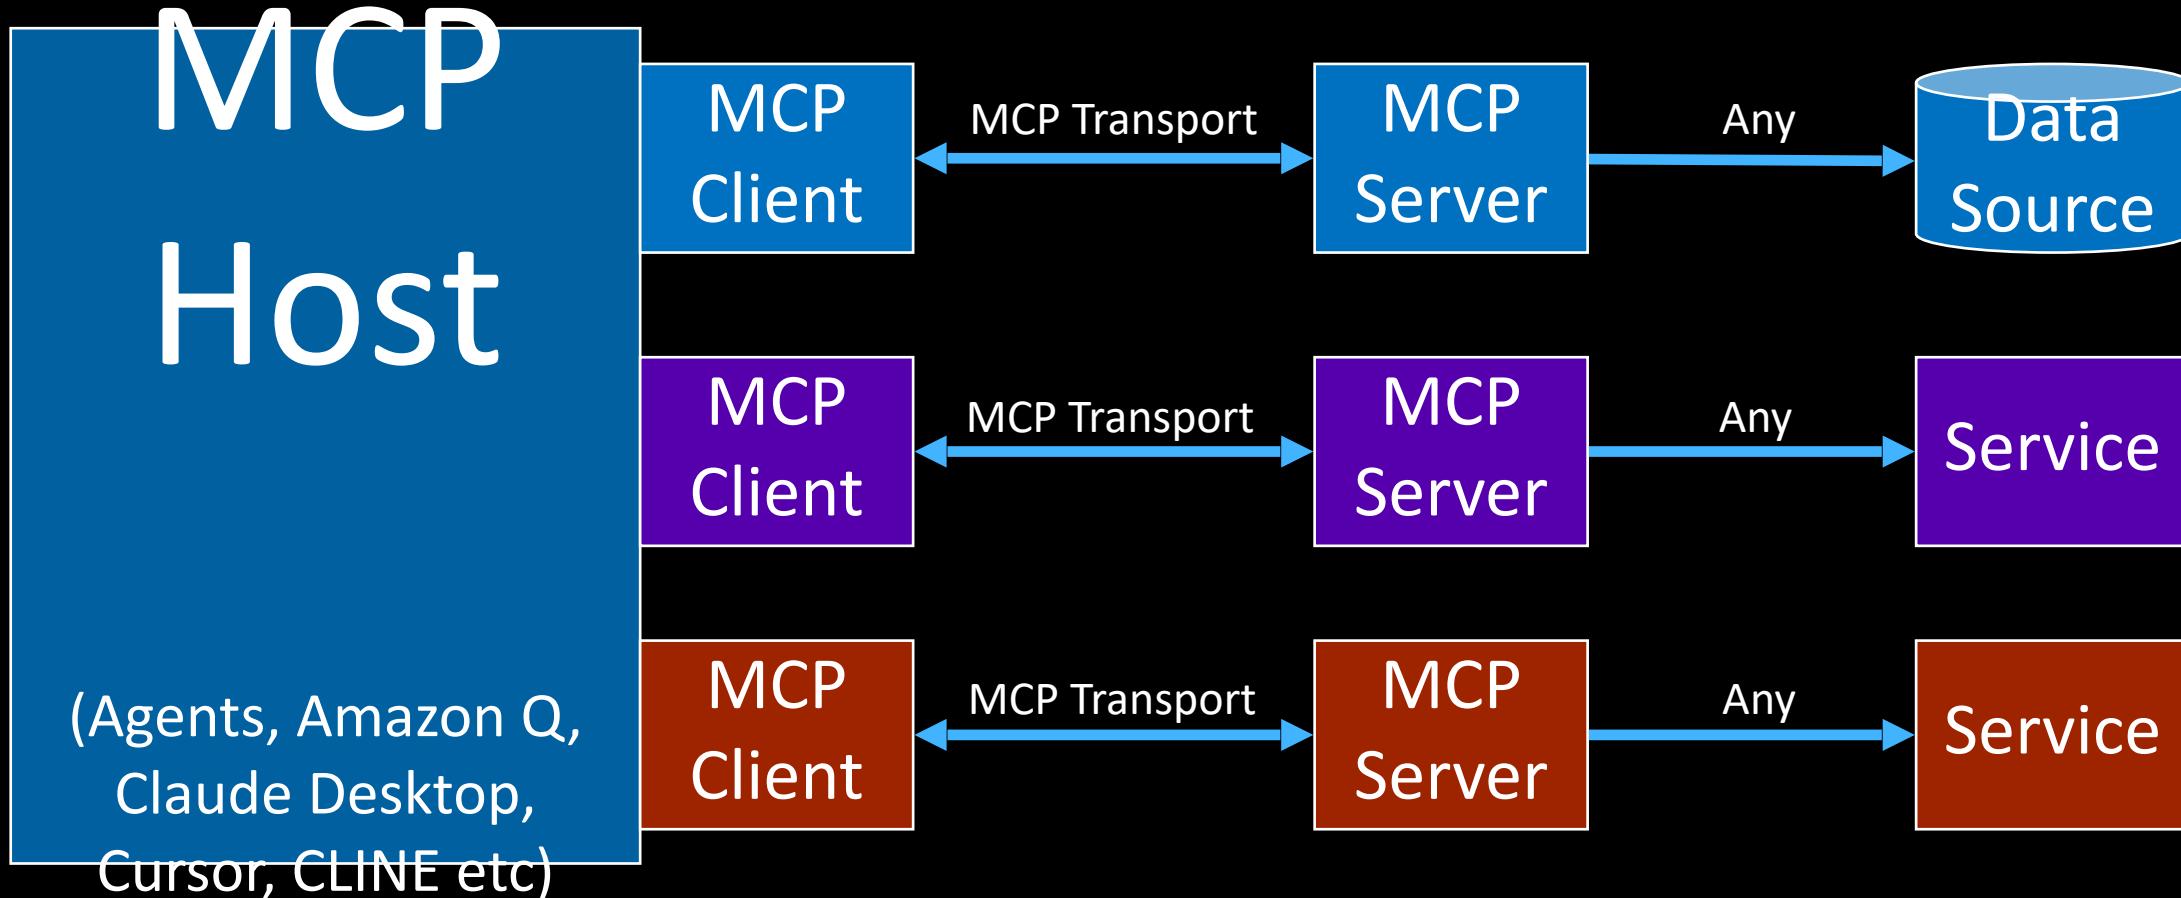
MCP
Host

(Agents, Amazon Q,
Claude Desktop,
Cursor, CLINE etc)

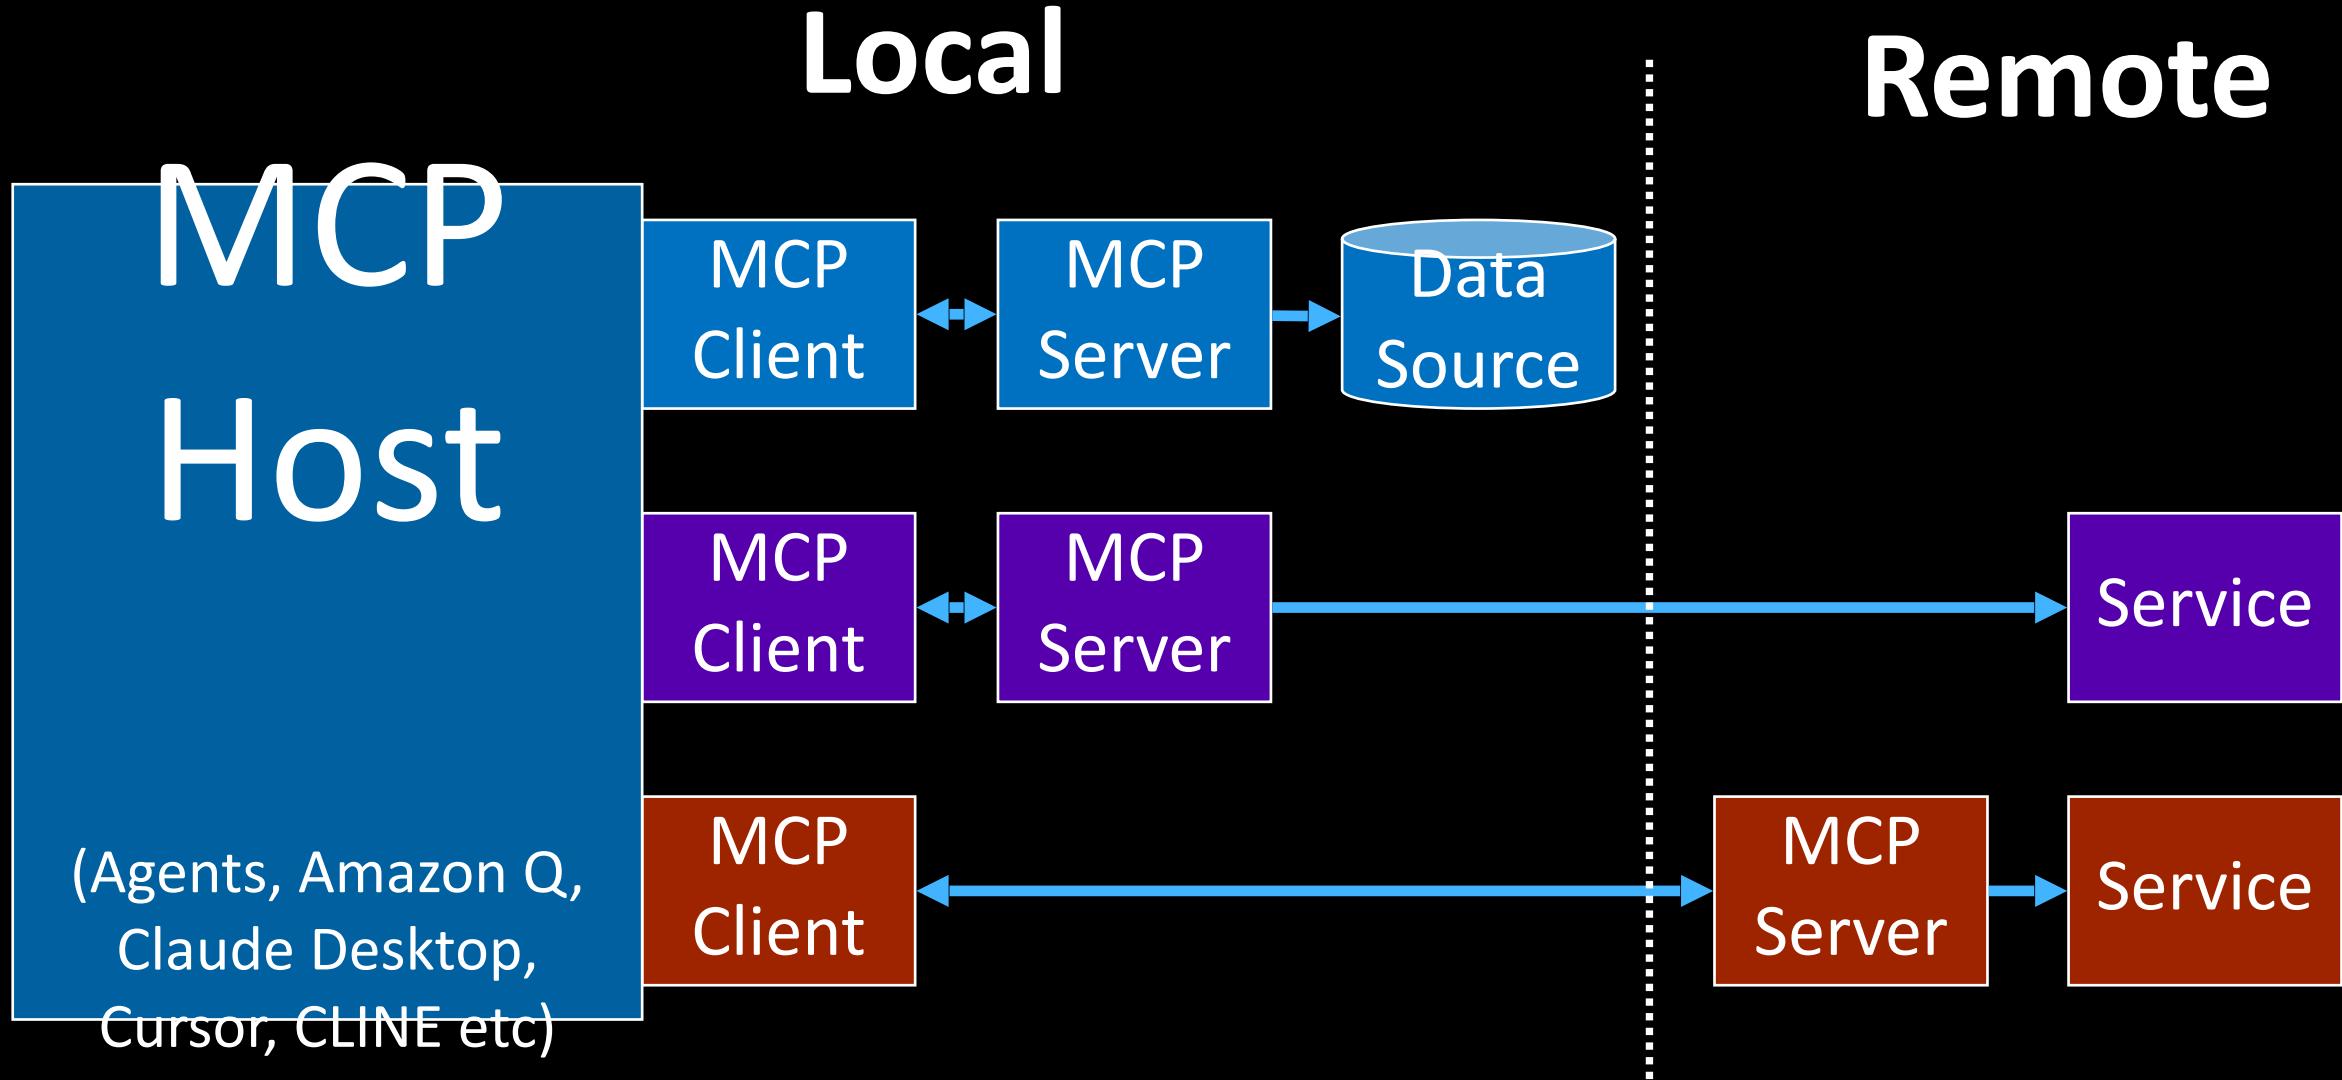
Model Context Protocol - Entities



Model Context Protocol - Entities



Model Context Protocol - Entities



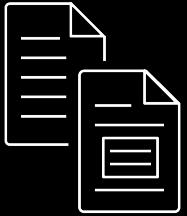
Model Context Protocol - Servers

- | | | | | | |
|---|--|---|---|--|---|
| <ul style="list-style-type: none">• AWS KB Retrieval - Re• Brave Search - Web a• EverArt - AI image gen• Everything - Reference• Fetch - Web content f• Filesystem - Secure fi• Git - Tools to read, sea• GitHub - Repository m• GitLab - GitLab API, et• Google Drive - File ac• Google Maps - Locatio• Memory - Knowledge• PostgreSQL - Read-o• Puppeteer - Browser a• Redis - Interact with R• Sentry - Retrieving an• Sequential Thinking -• Slack - Channel mana• Sqlite - Database inte• Time - Time and timez | <ul style="list-style-type: none">• ⚡ 21st.dev Magic - Create• ⚡ Adfin - The only platform reconciliations with Adfin.• ⚡ AgentQL - Enable AI age• ⚡ AgentRPC - Connect to a• ⚡ Aiven - Navigate your Aiv OpenSearch® services• ⚡ Apache IoTDB - MCP Se• ⚡ Apify - Actors MCP Ser• ⚡ APIMatic MCP - APIMatic processes OpenAPI files an• ⚡ Astra DB - Comprehensive database with a full range o• ⚡ Audience Insights - Marke demographic, cultural, influ• ⚡ Axiom - Query and analy• ⚡ Bankless Onchain - Que• ⚡ BICScan - Risk score / as• ⚡ Box - Interact with the In• ⚡ Browserbase - Automate and more)• ⚡ Chargebee - MCP Server• ⚡ Chroma - Embeddings, v• ⚡ Chronulus AI - Predict an | <ul style="list-style-type: none">• ⚡ CircleCI - Enable AI Ag• ⚡ ClickHouse - Query yo• ⚡ Cloudflare - Deploy, co Workers/KV/R2/D1)• ⚡ Codacy - Interact with your code.• ⚡ CodeLogic - Interact w architecture dependencie• ⚡ Comet Opik - Query ar LLMs in natural language• ⚡ Convex - Introspect an• ⚡ Dart - Interact with tas• ⚡ DevHub - Manage and• ⚡ E2B - Run code in secu• ⚡ EduBase - Interact with management, and contex• ⚡ Elasticsearch - Query• ⚡ eSignatures - Contrac• ⚡ Exa - Search Engine m• ⚡ FewSats - Enable AI Ag• ⚡ Fibery - Perform queri• ⚡ Financial Datasets - S• ⚡ Firecrawl - Extract web• ⚡ Fireproof - Immutable | <ul style="list-style-type: none">• ⚡ Gitee - Gitee API integr• ⚡ gotoHuman - Human-i to your gotoHuman inbox• ⚡ Grafana - Search dash• ⚡ Graphlit - Ingest anyth searchable Graphlit projec• ⚡ GreptimeDB - Provides GreptimeDB.• ⚡ Heroku - Interact with databases, and more.• ⚡ Hologres - Connect to• ⚡ Hyperbrowser - Hyper effortless, scalable brows• ⚡ IBM wflows - Tool platfo• ⚡ ForeverVM - Run Python• ⚡ Inbox Zero - AI persona• ⚡ Inkeep - RAG Search o• ⚡ Integration App - Inter• ⚡ JetBrains - Work on yo• ⚡ Kagi Search - Search t• ⚡ Keboola - Build robust• ⚡ Lara Translate - MCP S for language detection an• ⚡ Logfire - Provides acc• ⚡ Langfuse Prompt Man releasing prompts. | <ul style="list-style-type: none">• ⚡ Lingo.dev - Make your• ⚡ Mailgun - Interact with• ⚡ Make - Turn your Make• ⚡ Meilisearch - Interact &• ⚡ Metoro - Query and inte• ⚡ Milvus - Search, Query• ⚡ Momento - Momento C any scale.• ⚡ MotherDuck - Query ar• ⚡ Needle - Production-re• ⚡ Neo4j - Neo4j graph da memory• ⚡ Neon - Interact with the• ⚡ Notion - This project im• ⚡ OceanBase - MCP Serv• ⚡ Octagon - Deliver real-• ⚡ Oxylabs - Scrape webs data extraction.• ⚡ Paddle - Interact with th• ⚡ PayPal - PayPal's officia• ⚡ Perplexity - An MCP se in conversational AI.• ⚡ Qdrant - Implement ser• ⚡ Ramp - Interact with Ra• ⚡ Raygun - Interact with you | <ul style="list-style-type: none">• ⚡ Rember - Create spa• ⚡ Riza logo Riza - Arbitr• ⚡ Search1API - One API• ⚡ ScreenshotOne - Ren• ⚡ Semgrep - Enable AI• ⚡ SingleStore - Interact• ⚡ StarRocks - Interact w• ⚡ Stripe - Interact with• ⚡ Tavily - Search engine• ⚡ Thirdweb - Read/writ transaction execution, p• ⚡ Tinybird - Interact wit• ⚡ UnifAI - Dynamically s• ⚡ Unstructured - Set up Platform• Vectorize - Vectorize M extraction and text chur• ⚡ Verodat - Interact wit• ⚡ VeyraX - Single tool t• ⚡ Xero - Interact with th• ⚡ Zapier - Connect you• ZenML - Interact with yo |
|---|--|---|---|--|---|

<https://github.com/modelcontextprotocol/servers>

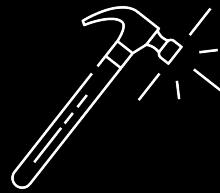


MCP server building blocks



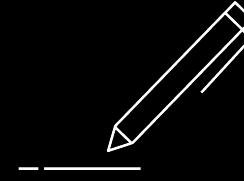
Resources

Enable **applications** to retrieve additional data



Tools

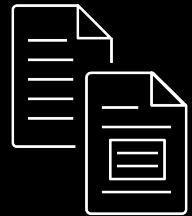
Enable **models** to perform actions



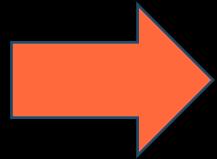
Prompts

Enable **users** with reusable prompt templates

Resources



Resources

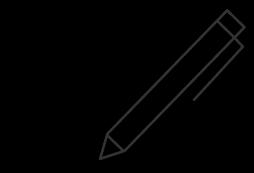


```
@mcp.resource("menu://{{shop_id}}")
def menu(shop_id: str) -> str:
    """Returns menu for a specified pizza shop"""
    menu = pizza_api.get_menu(shop_id)
    return menu
```



Tools

```
@mcp.tool()
def order_pizza(toppings: str, address: str) -> str:
    """Orders pizza"""
    order_id = pizza_api.submit_order(toppings, address)
    return f"Thank you! Your order ID is {order_id}"
```



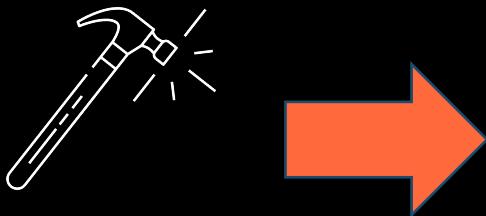
Prompts

```
@mcp.prompt(title="Get my usual")
def order_pizza_template(toppings: str) -> str:
    return f"""
        Order a pizza from AcmePie to my home address.
        I want thin crust, extra cheese, spicy sauce.
        Add following toppings: {toppings}
        """
```

Tools



Resources



Tools



Prompts



```
@mcp.resource("menu://{{shop_id}}")
def menu(shop_id: str) -> str:
    """Returns menu for a specified pizza shop"""
    menu = pizza_api.get_menu(shop_id)
    return menu
```

```
@mcp.tool()
def order_pizza(toppings: str, address: str) -> str:
    """Orders pizza"""
    order_id = pizza_api.submit_order(toppings, address)
    return f"Thank you! Your order ID is {order_id}"
```

```
@mcp.prompt(title="Get my usual")
def order_pizza_template(toppings: str) -> str:
    return f"""
        Order a pizza from AcmePie to my home address.
        I want thin crust, extra cheese, spicy sauce.
        Add following toppings: {toppings}
        """
```

Prompts



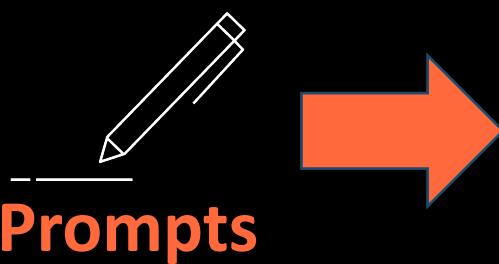
Resources

```
@mcp.resource("menu://{{shop_id}}")
def menu(shop_id: str) -> str:
    """Returns menu for a specified pizza shop"""
    menu = pizza_api.get_menu(shop_id)
    return menu
```



Tools

```
@mcp.tool()
def order_pizza(toppings: str, address: str) -> str:
    """Orders pizza"""
    order_id = pizza_api.submit_order(toppings, address)
    return f"Thank you! Your order ID is {order_id}"
```



Let's build a solution

Building your first MCP server **for** AWS



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Building your first MCP Server

A

Give me a list of Lambda functions in my account that are using nodejs runtime. Group by nodejs version. Tell me which ones are using deprecated runtimes.

Building your first MCP Server

A Give me a list of Lambda functions in my account that are using nodejs runtime. Group by nodejs version. Tell me which ones are using deprecated runtimes.

I can't directly list Lambda functions in your AWS account since I don't have access to your AWS resources. However, I can guide you on how to retrieve this information yourself.

To list Lambda functions using Node.js runtimes in your account, you can use the AWS CLI or AWS Management Console.

Here's an AWS CLI command that will help:

bash

```
aws lambda list-functions --query "Functions[?Runtime.contains(@, 'nodejs')].{Name:Fun
```

Building your first MCP Server



I want to use AI assistant to
ask questions about Lambda
functions in my account

Building your first MCP Server

Do not implement the low-level wire protocol yourself.

Use official open-source SDKs.

-  Python SDK
-  TypeScript SDK
-  Java SDK
-  Kotlin SDK
-  C# SDK

<https://modelcontextprotocol.io/introduction>



Building your first MCP Server – Create Tools

```
const toolName = "list-functions";
```

Self-explanatory names



Building your first MCP Server – Create Tools

```
const toolName = "list-functions";  
  
const toolDescription =  
  "1. This tool can be used for retrieving a list of AWS Lambda functions. \  
  2. This tool should be used to answer questions about what functions are available \  
  in user's account, as well as questions about specific properties of these functions. \  
  3. This tool should not be used when you need information about one specific function. \  
  If you need information about one specific function, use the get-function tool instead. \  
  4. This tool supports optional pagination via the marker parameter. If you see a 'marker' \  
  property in the tool response, you MUST use it to make another request to the tool to \  
  continue building the list. When marker is not available, it means you've reached the \  
  last page of results. When you're using this tool for the first time, you obviously do not have \  
  marker yet, so supply an empty string as a value instead \  
  5. This tool returns a JSON object with two elements. The first element is an array of \  
  functions. The second element is a marker that can be used for further pagination. \  
";
```

Descriptions are crucial



Building your first MCP Server – Create Tools

```
const toolName = "list-functions";

const toolDescription =
  "1. This tool can be used for retrieving a list of AWS Lambda functions. \
  2. This tool should be used to answer questions about what functions are available \
  in user's account, as well as questions about specific properties of these functions. \
  3. This tool should not be used when you need information about one specific function. \
  If you need information about one specific function, use the get-function tool instead. \
  4. This tool supports optional pagination via the marker parameter. If you see a 'marker' \
  property in the tool response, you MUST use it to make another request to the tool to \
  continue building the list. When marker is not available, it means you've reached the \
  last page of results. When you're using this tool for the first time, you obviously do not have \
  marker yet, so supply an empty string as a value instead \
  5. This tool returns a JSON object with two elements. The first element is an array of \
  functions. The second element is a marker that can be used for further pagination. \
  ";


```

Descriptions are crucial



Building your first MCP Server – Create Tools

```
const toolName = "list-functions";

const toolDescription =
    "1. This tool can be used for retrieving a list of AWS Lambda functions. \
     2. This tool should be used to answer questions about what functions are available \
        in user's account, as well as questions about specific properties of these functions. \
     3. This tool should not be used when you need information about one specific function. \
        If you need information about one specific function, use the get-function tool instead. \
     4. This tool supports optional pagination via the marker parameter. If you see a 'marker' \
        property in the tool response, you MUST use it to make another request to the tool to \
        continue building the list. When marker is not available, it means you've reached the \
        last page of results. When you're using this tool for the first time, you obviously do not have \
        marker yet, so supply an empty string as a value instead \
     5. This tool returns a JSON object with two elements. The first element is an array of \
        functions. The second element is a marker that can be used for further pagination. \
    ";

const toolParamsSchema = {
    marker: z.string().optional().describe(
        "Pagination marker. Send empty string for the first \
         request, when previous marker is not yet available."
    )
};
```

Parameters schema allows AI assistants to know which parameters tools support



Building your first MCP Server – Create Tools

```
const ASYNC_IMPL = async ({ marker }) => {
  const command = new ListFunctionsCommand({
    Marker: marker || null,
    MaxItems: 50
  });

  const lambdaResponse = await lambdaClient.send(command);

  const result = {
    functions: lambdaResponse.Functions,
    marker: lambdaResponse.NextMarker,
  };

  return {
    content: [
      {
        type: "text",
        text: JSON.stringify(result)
      }
    ]
  }
}
```

Uses AWS SDK and Lambda client to retrieve the list of functions. Uses next page marker, if available.

Building your first MCP Server – Create Tools

```
const ASYNC_IMPL = async ({ marker }) => {
  const command = new ListFunctionsCommand({
    Marker: marker || null,
    MaxItems: 50
  });

  const lambdaResponse = await lambdaClient.send(command);

  const result = {
    functions: lambdaResponse.Functions,
    marker: lambdaResponse.NextMarker,
  };

  return {
    content: [
      {
        type: "text",
        text: JSON.stringify(result)
      }
    ]
}

}
```

Return the list of functions and next page marker



Building your first MCP Server – Create Server

```
const server = new McpServer({  
    name: "AWS Lambda Operations MCP Server",  
    version: "0.0.1"  
}, {
```

Define server name
and version



Building your first MCP Server – Create Server

```
const server = new McpServer({  
  name: "AWS Lambda Operations MCP Server",  
  version: "0.0.1"  
, {  
  capabilities: {  
    tools: {},  
    resources: {}  
,  
  instructions: 'Use this MCP server to list AWS Lambda functions'  
});
```

Detailed instructions are crucial. They allow LLM understand why, when, and how they should be using your server.



Building your first MCP Server – Create Server

```
const server = new McpServer({
  name: "AWS Lambda Operations MCP Server",
  version: "0.0.1"
}, {
  capabilities: {
    tools: {},
    resources: {}
  },
  instructions: 'Use this MCP server to list AWS Lambda functions'
});
```

```
server.tool(
  listFunctions.toolName,
  listFunctions.toolDescription,
  listFunctions.toolParamsSchema,
  listFunctions.toolCallback);
```

Add your tools. Do not skip descriptions and parameters schemas.



Building your first MCP Server – Use in AI Assistant

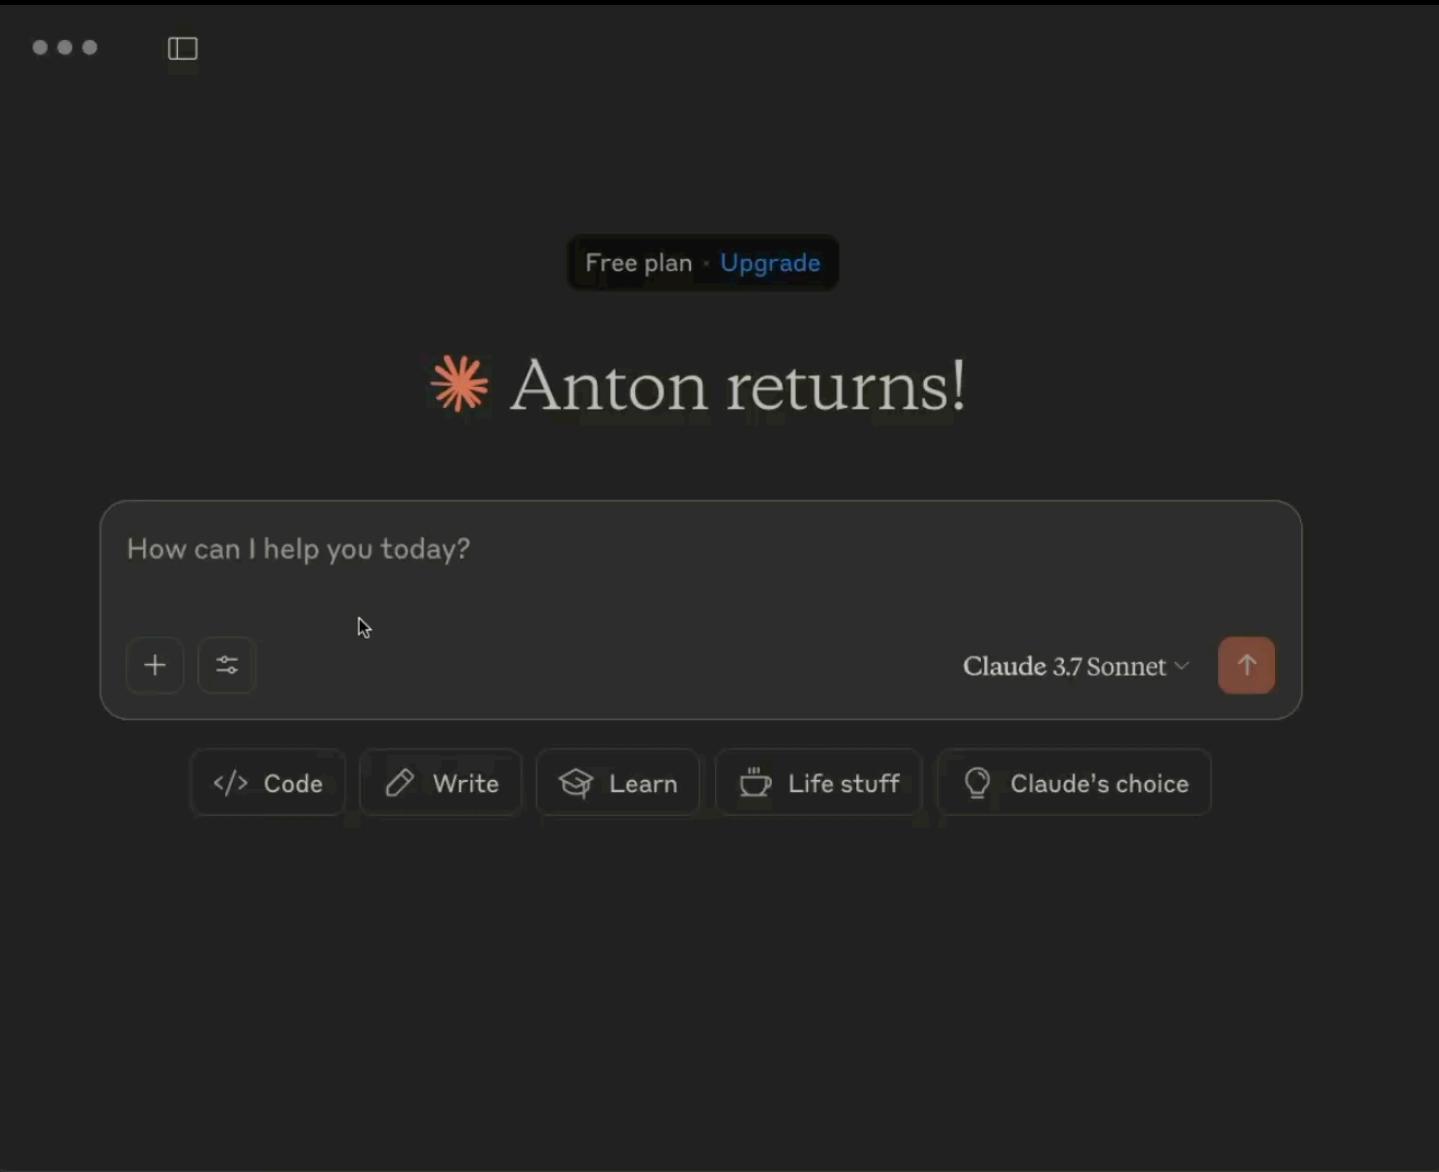
claude_desktop_config.json or mcp.json

```
{  
  "mcpServers": {  
    "lambda-ops": {  
      "command": "node",  
      "args": [  
        "/{full_path_to}/server.js"  
      ]  
    }  
  }  
}
```

Configure AI assistant to use
your MCP Server



Building your first MCP Server – Use in AI Assistant



Building your first MCP Server – Use in AI Assistant

The screenshot shows the AWS Lambda console interface. At the top, there are three dots, a refresh icon, and the text "Node.js Lambda Functions by Runtime Version". On the right side, there are two orange buttons with upward arrows and a plus sign. Below this, the title "Deprecated Runtimes" is displayed in bold. A message states: "According to AWS documentation, the following Node.js runtimes are deprecated:". A numbered list follows: 1. **nodejs16.x** - This runtime is deprecated and should be migrated to a newer version. You have 4 functions using this runtime:

- observability-demo-function
- observability-demo-function-poller
- audit-layer-demo-nodejs
- ModernizingCloudSecurityS-LogRetentionaae0aa3c5b4d-nrxTmLqngB1C

 All other Node.js functions in your account are using supported runtimes (nodejs18.x, nodejs20.x, and nodejs22.x). A callout box contains the text: "For best practices, I recommend upgrading your nodejs16.x functions to at least nodejs18.x or preferably to the latest nodejs22.x runtime to ensure continued support and security updates." At the bottom, there is a red star icon, a "Reply to Claude..." input field, and a "Claude 3.7 Sonnet" dropdown menu. The AWS logo is in the bottom left corner.

Building your first MCP Server – Use in AI Assistant

All other Node.js functions in your account are using supported runtimes (nodejs18.x, nodejs20.x, and nodejs22.x).

For best practices, I recommend upgrading your nodejs16.x functions to at least nodejs18.x or preferably to the latest nodejs22.x runtime to ensure continued support and security updates.

Building your first MCP Server



I want to use AI assistant to
safely update functions to the
latest runtime version

Building your first MCP Server



I want to use AI assistant to
safely update functions to the
latest runtime version

AWS Lambda Operations MCP Server

Tools

- get-runtime-info
- list-functions
- get-function
- invoke-function
- update-function-runtime

Update Function Configuration

```
const toolName = "update-function-runtime";

const toolDescription =
  "1. This tool can be used for updating the runtime version used by a Lambda function. \
  You MUST supply two parameters as explained below. \
  2. A 'functionName' parameter value MUST be supplied by the caller. \
  3. A 'runtime' parameter value MUST be supplied by the caller. \
  4. After updating a function runtime, you SHOULD invoke the function \
  in order to validate it still works. ALWAYS ask the user if they want to check if function is still \
  working. Use invoke-function tool to test if the function can still be invoked. If invocation fails \
  rollback to the previous runtime version automatically. \
";

const toolParamsSchema = {
  functionName: z.string().describe("Name of the Lambda function to retrieve information about."),
  runtime: z.string().describe("The new runtime version to be used by the Lambda function.")
}
```

The result

The Transport

How does it work under-the-hood?



Message Types

Request

```
{  
  "jsonrpc": "2.0",  
  "id": 0,  
  "method": "initialize",  
  "params": {  
    "protocolVersion": "2024-11-05",  
    "capabilities": {  
      "sampling": {},  
      "roots": {}  
    },  
    "clientInfo": {  
      "name": "mcp-inspector",  
      "version": "0.8.2"  
    }  
  }  
}
```

Response

```
{  
  "jsonrpc": "2.0",  
  "id": 0,  
  "result": {  
    "protocolVersion": "2024-11-05",  
    "capabilities": {  
      "tools": {}  
    },  
    "serverInfo": {  
      "name": "Demo",  
      "version": "1.0.0"  
    },  
    "instructions": "This is a demo server"  
  }  
}
```

Notification

```
{  
  "jsonrpc": "2.0",  
  "method": "notifications/initialized"  
}
```

Clients and servers can send
messages in either direction



Transport Types



Stdio (local)

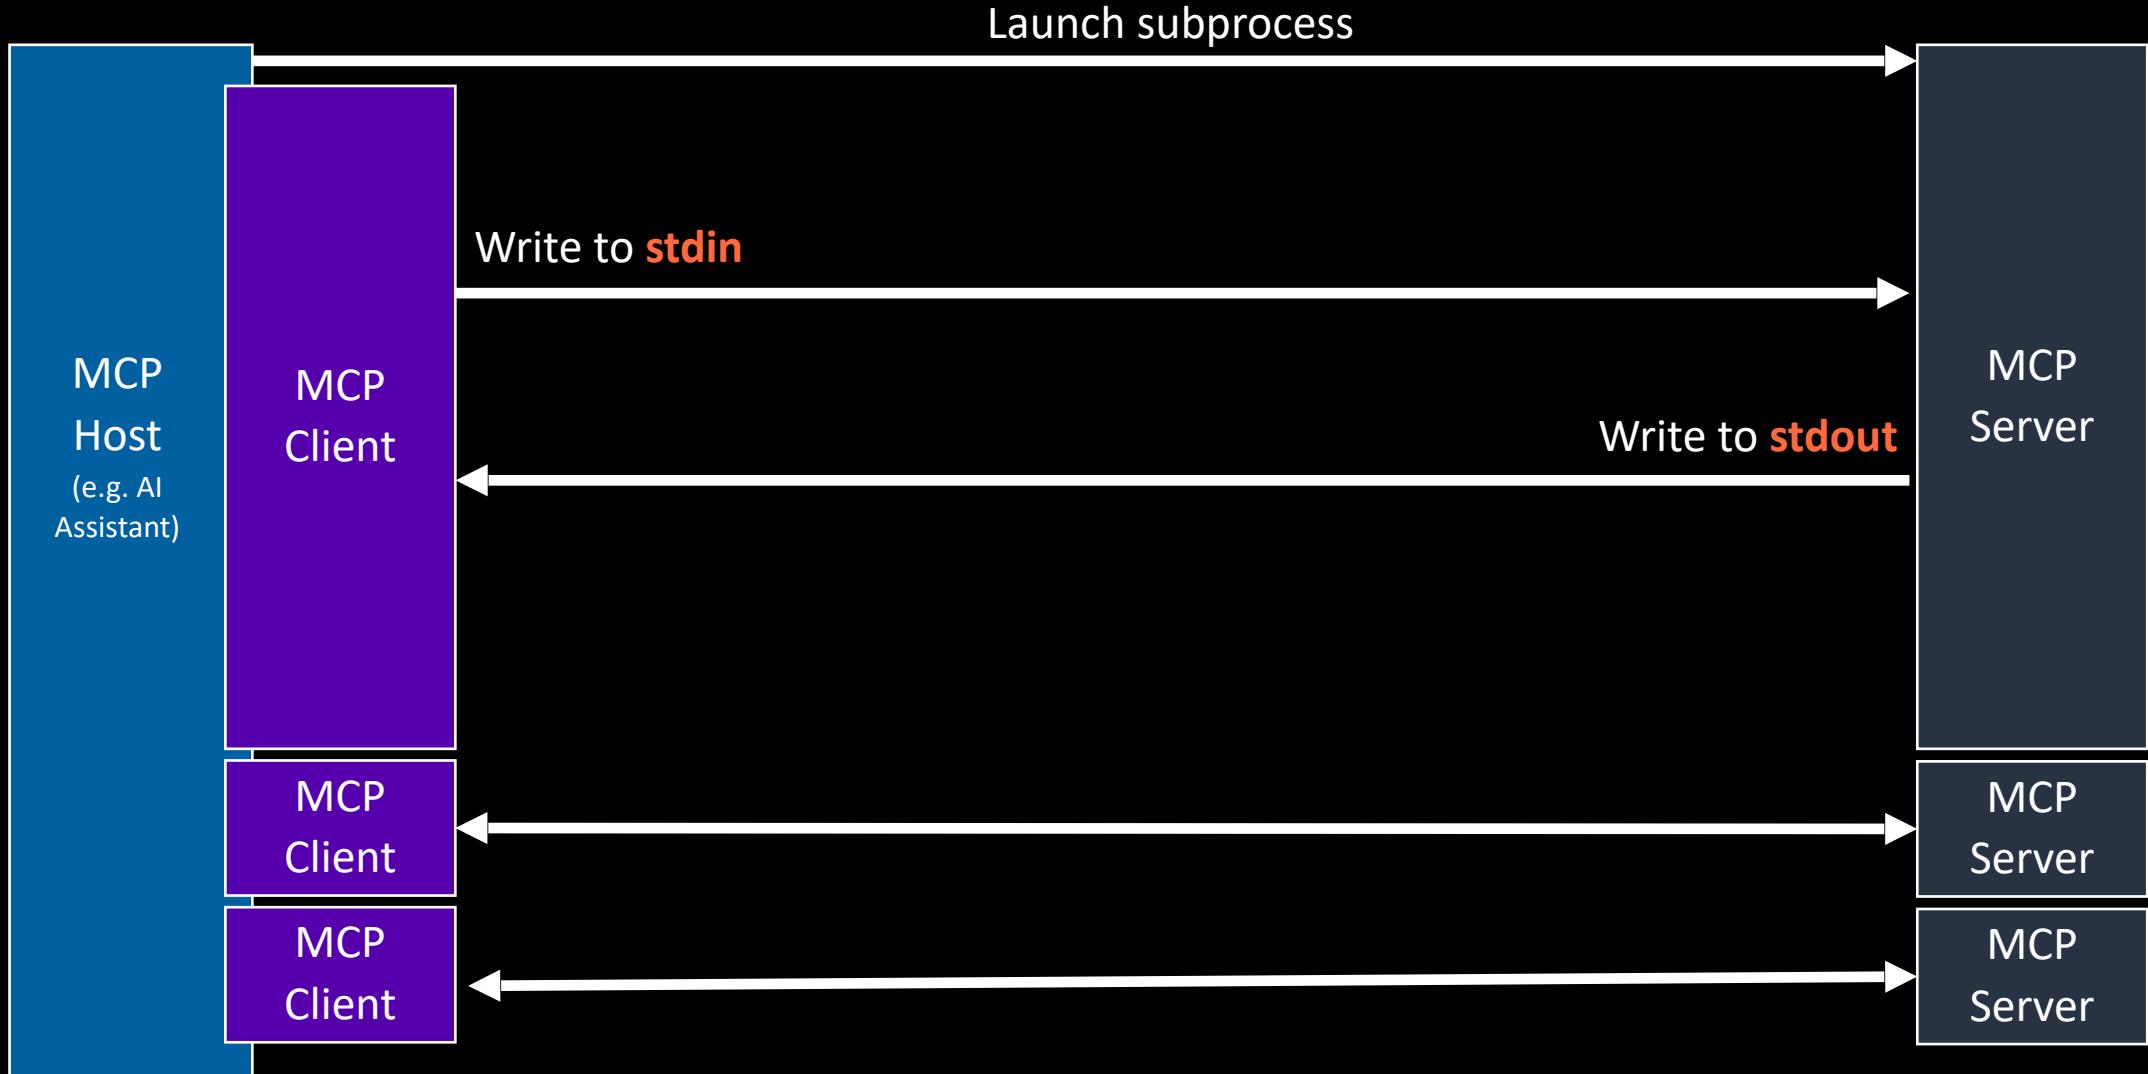
Communicates through standard input/output streams. Use for integrations with local MCP Servers and command-line tools.

Streamable HTTP (remote)

Communicates over HTTP, with optional persistent SSE connection. Use for integrations with remote MCP Servers.

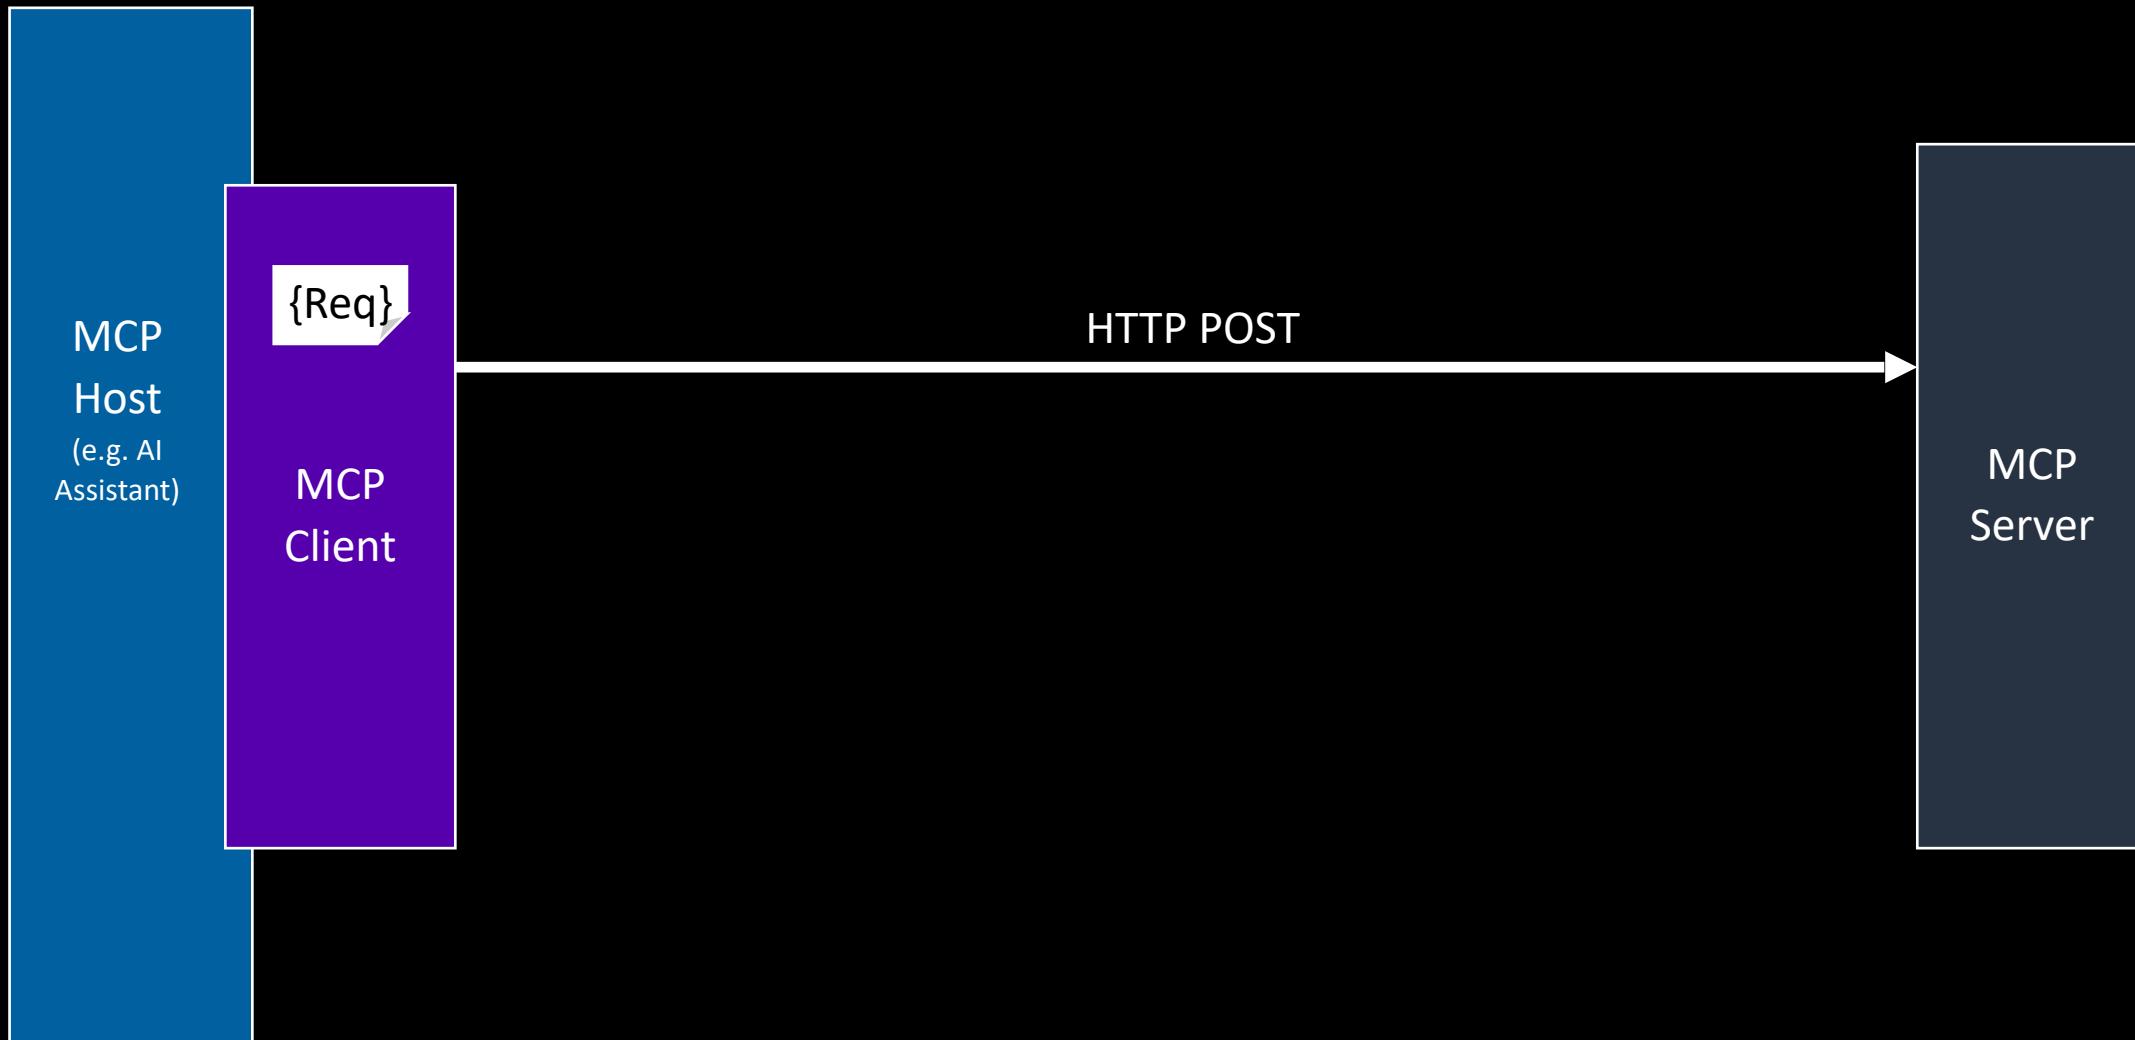
<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Stdio Transport



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



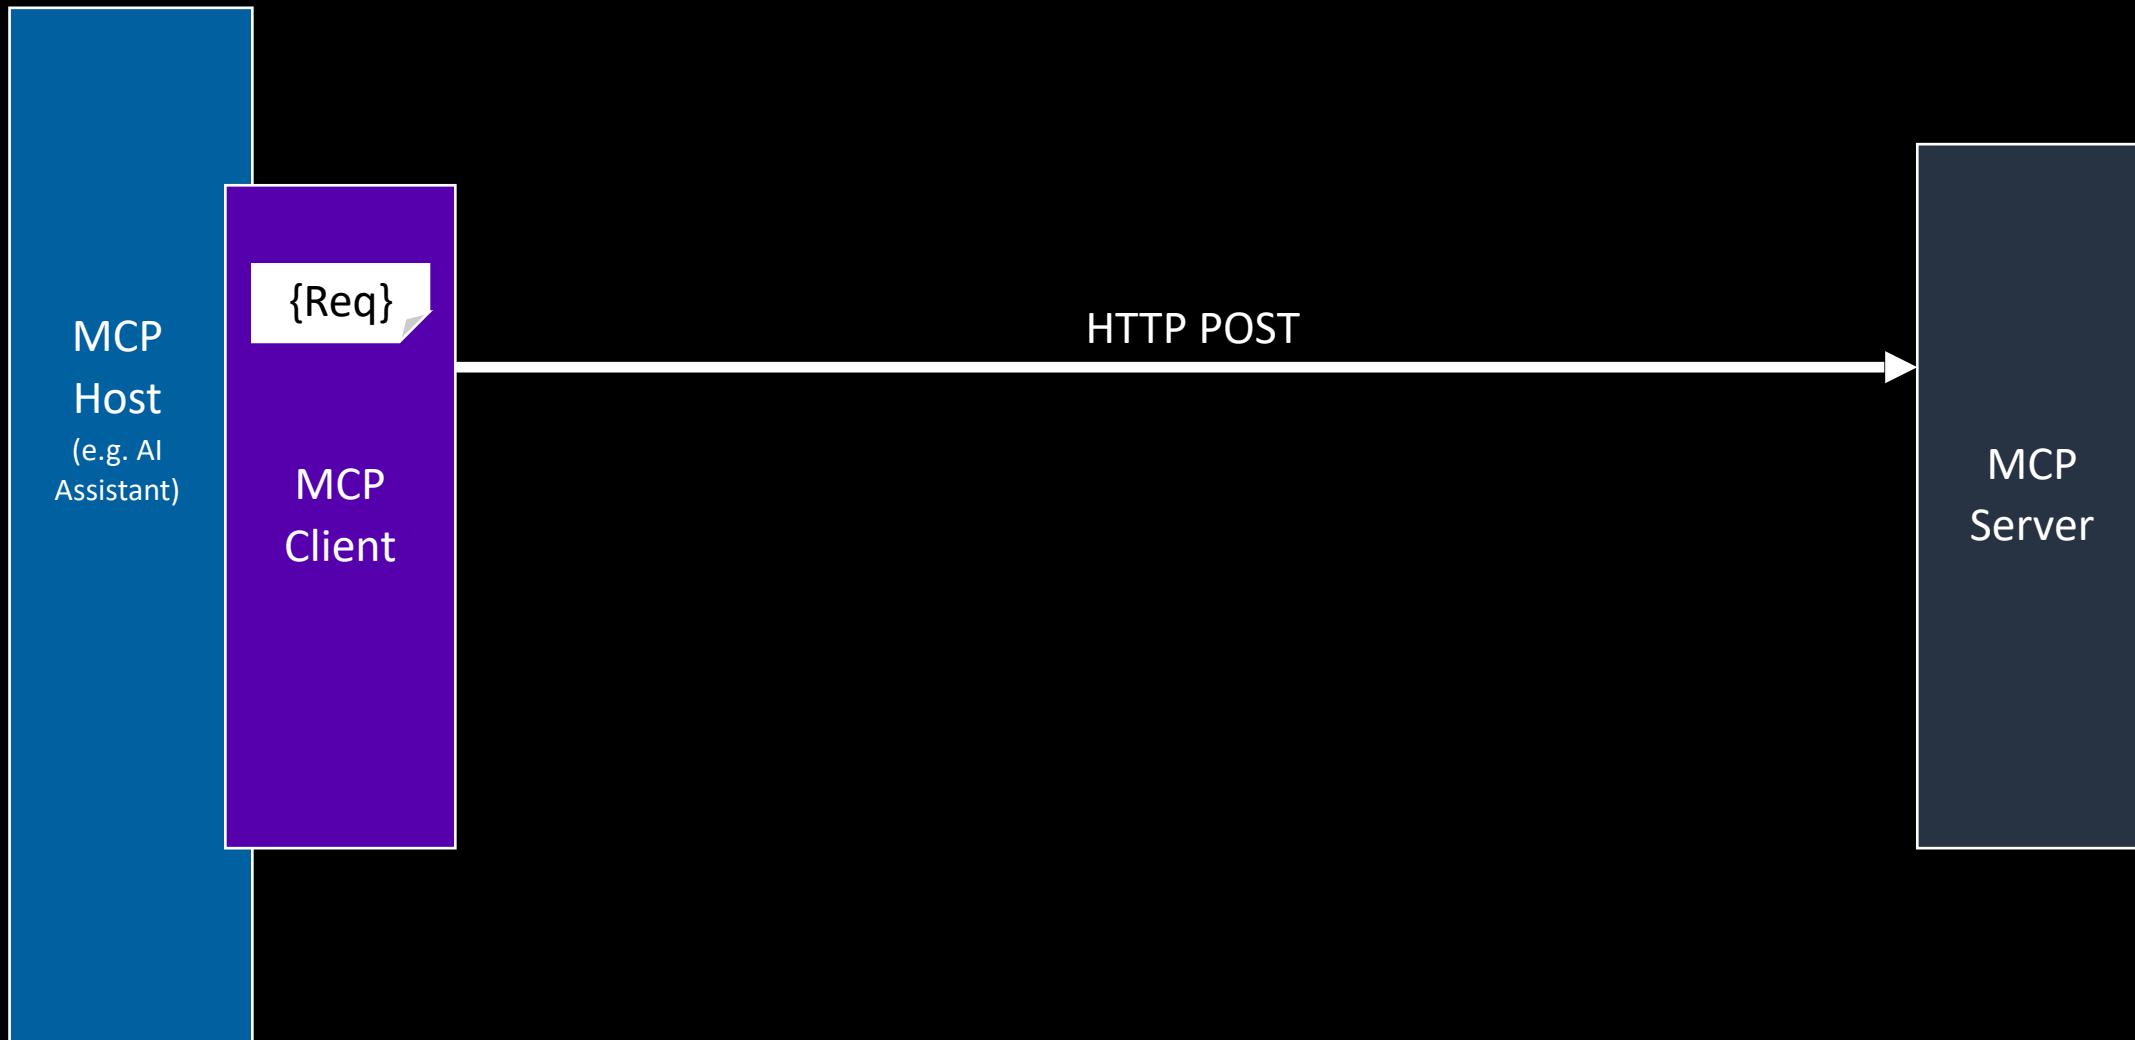
<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



Streamable HTTP (with SSE connection)



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



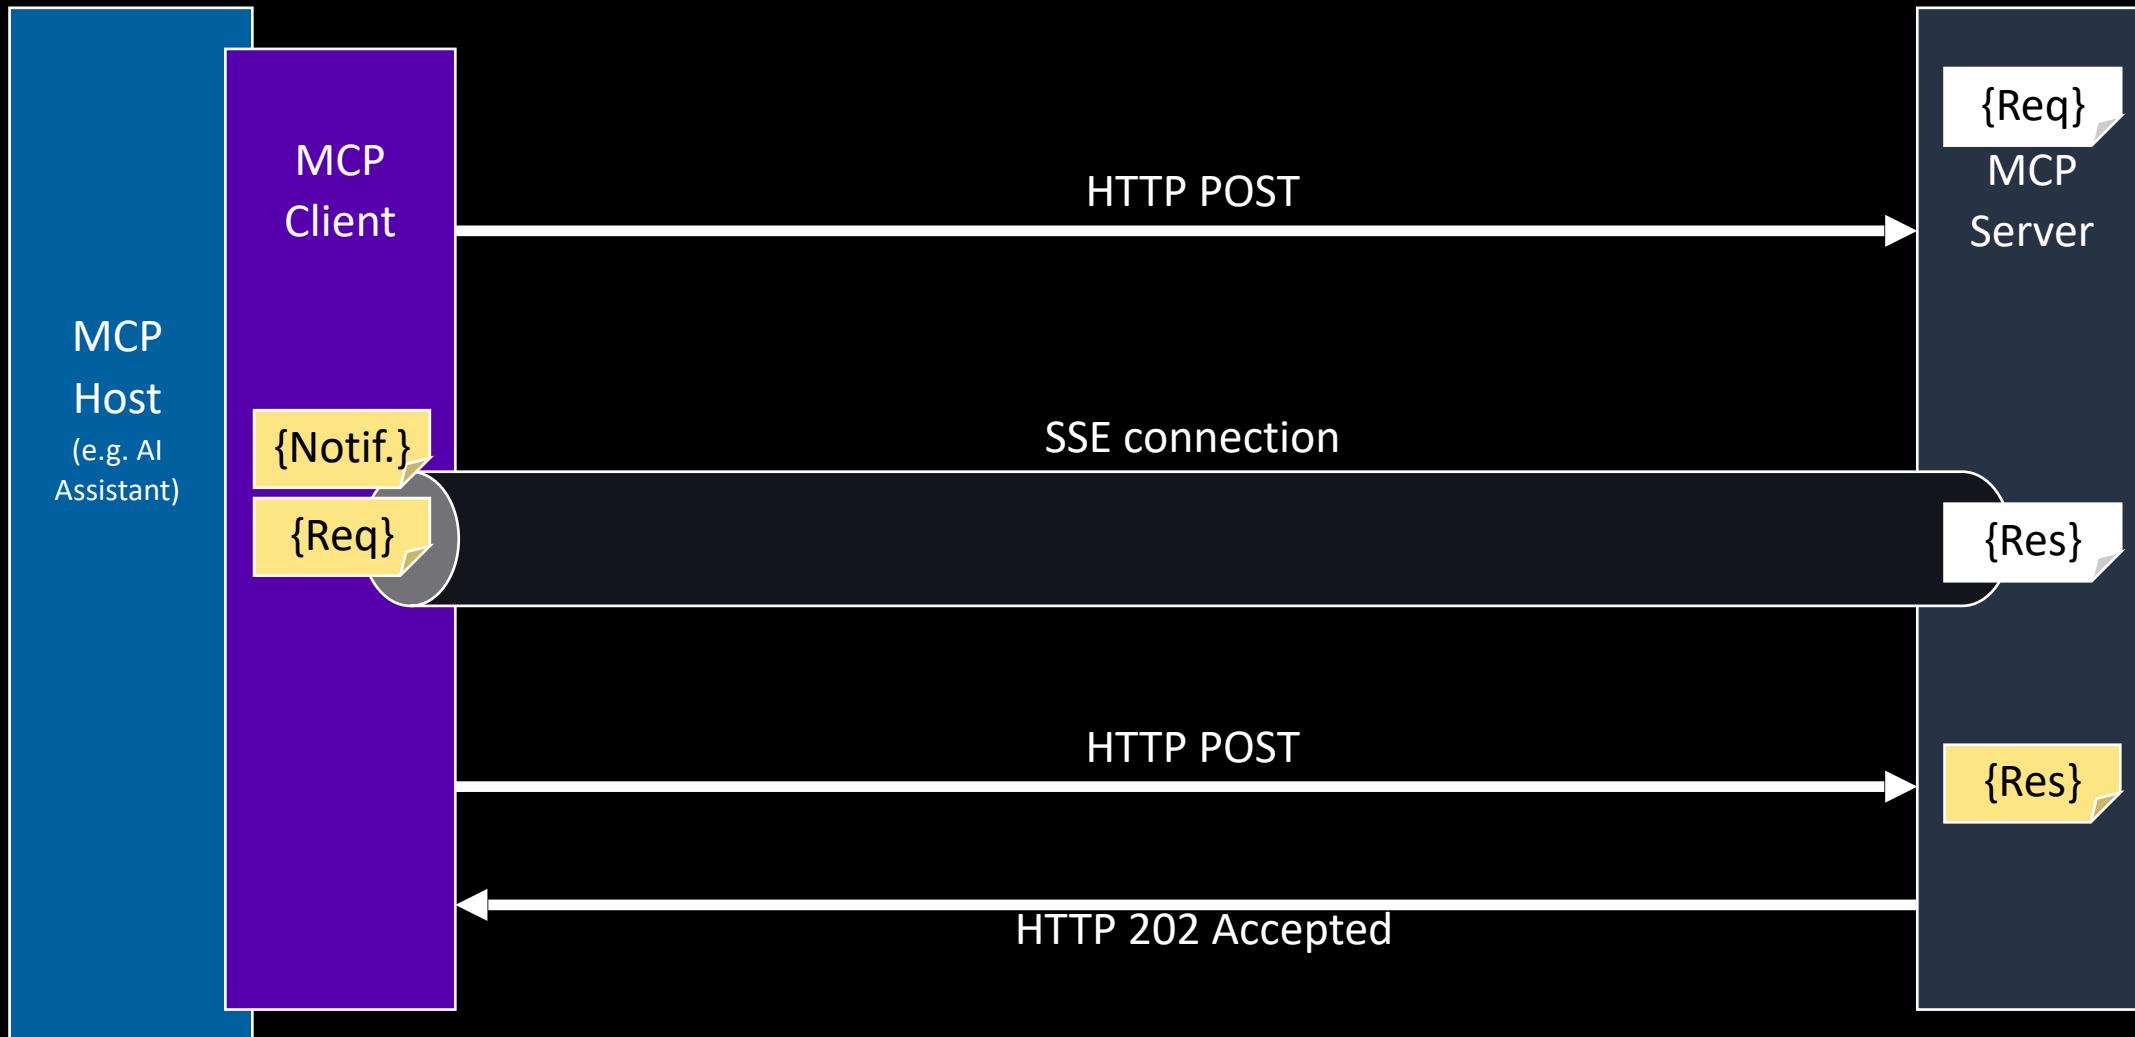
<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



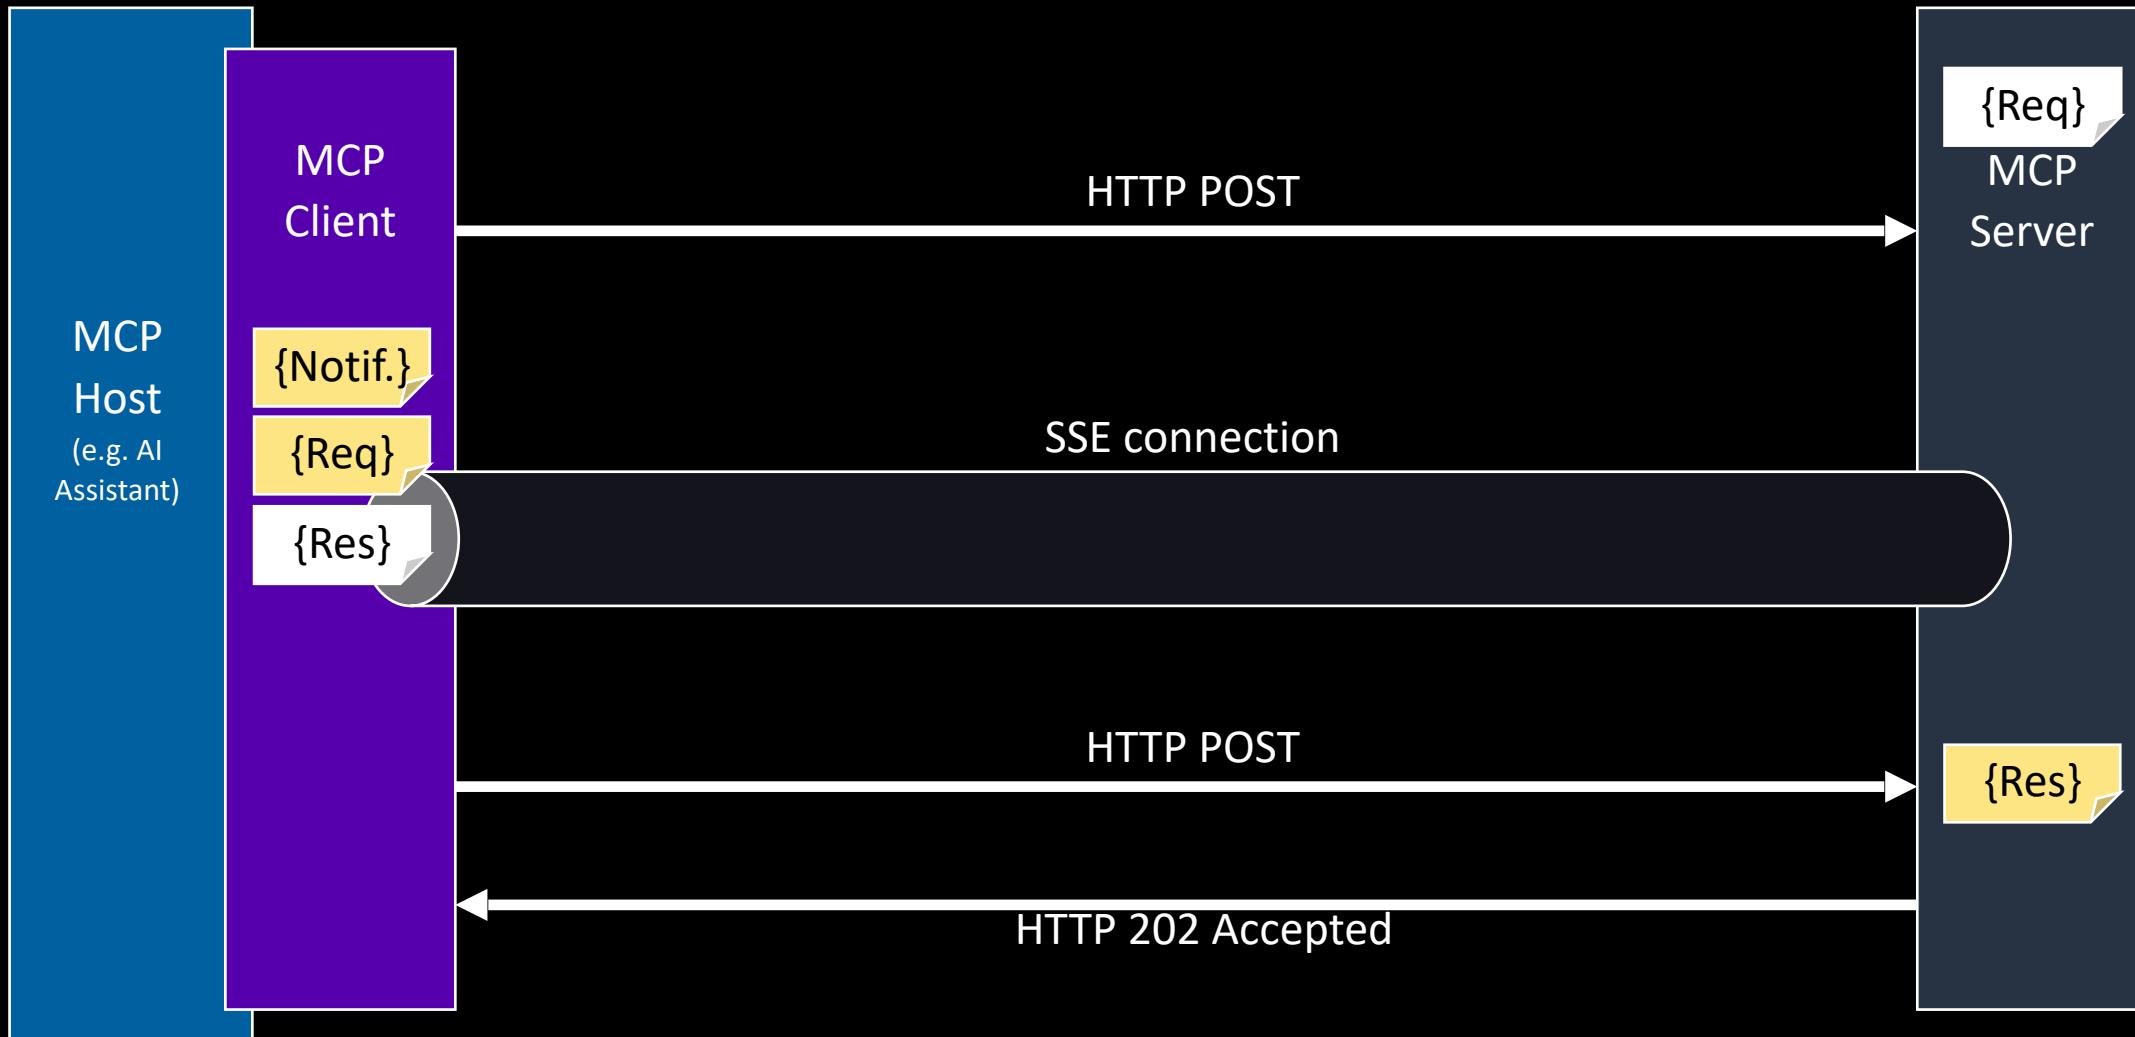
<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



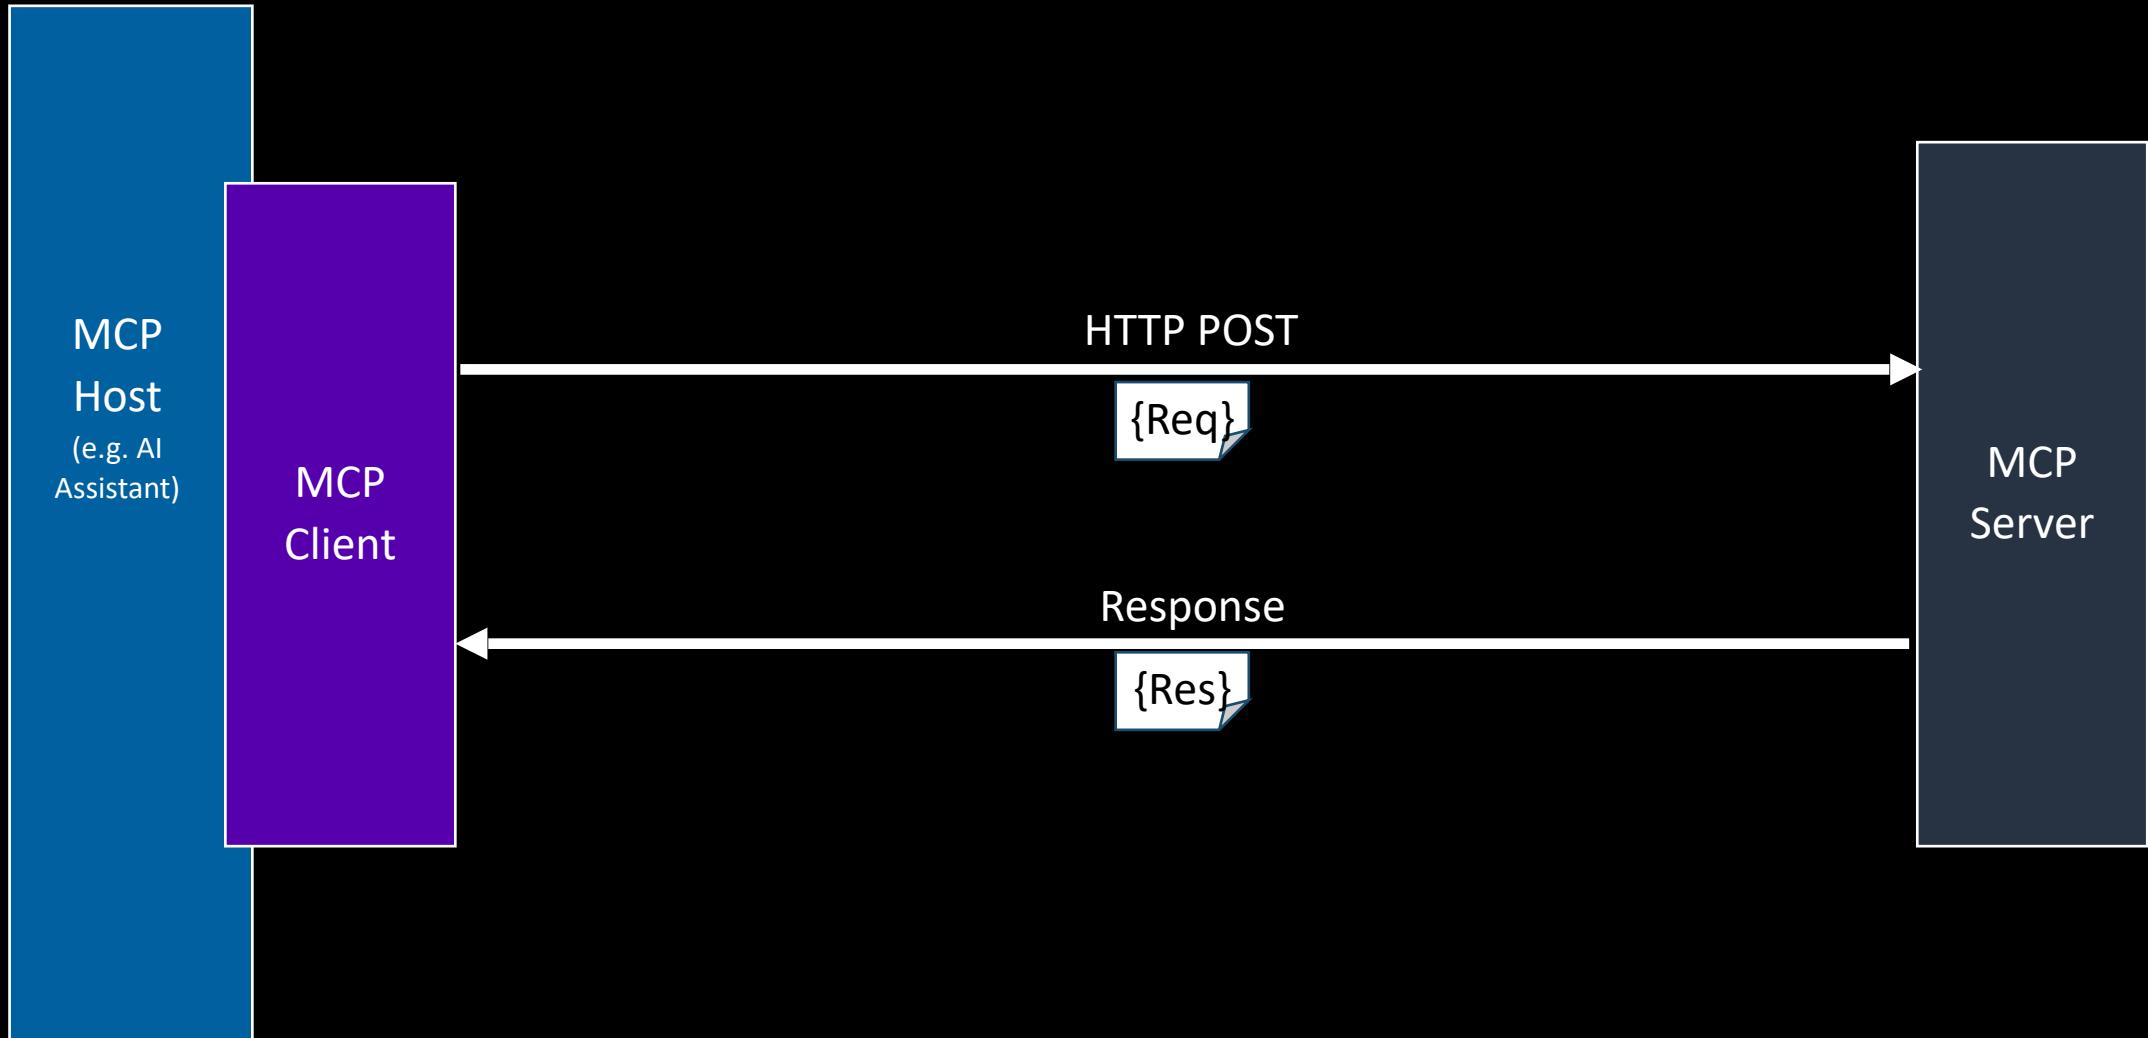
<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with SSE connection)



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Streamable HTTP (with JSON response)



<https://modelcontextprotocol.io/specification/2025-06-18/basic/transports>

Remote MCP Servers

Running your MCP servers **on** AWS



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Deploying to AWS



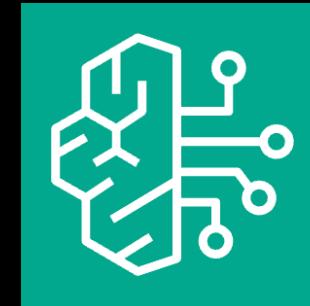
AWS Lambda



Amazon Elastic
Container Service (ECS)



Amazon Elastic
Kubernetes Service (EKS)



Amazon Bedrock
AgentCore (preview)

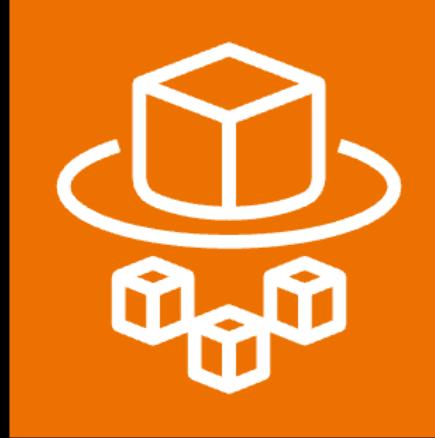
(or ANY other compute type, your Agent is just a Python app)

Building Serverless MCP Servers on AWS



AWS Lambda

**Stateless /
connectionless**



Amazon ECS on AWS Fargate

**Persistent stateful
connections with SSE or
stateless/connectionless**

Stateless MCP Servers

with AWS Lambda

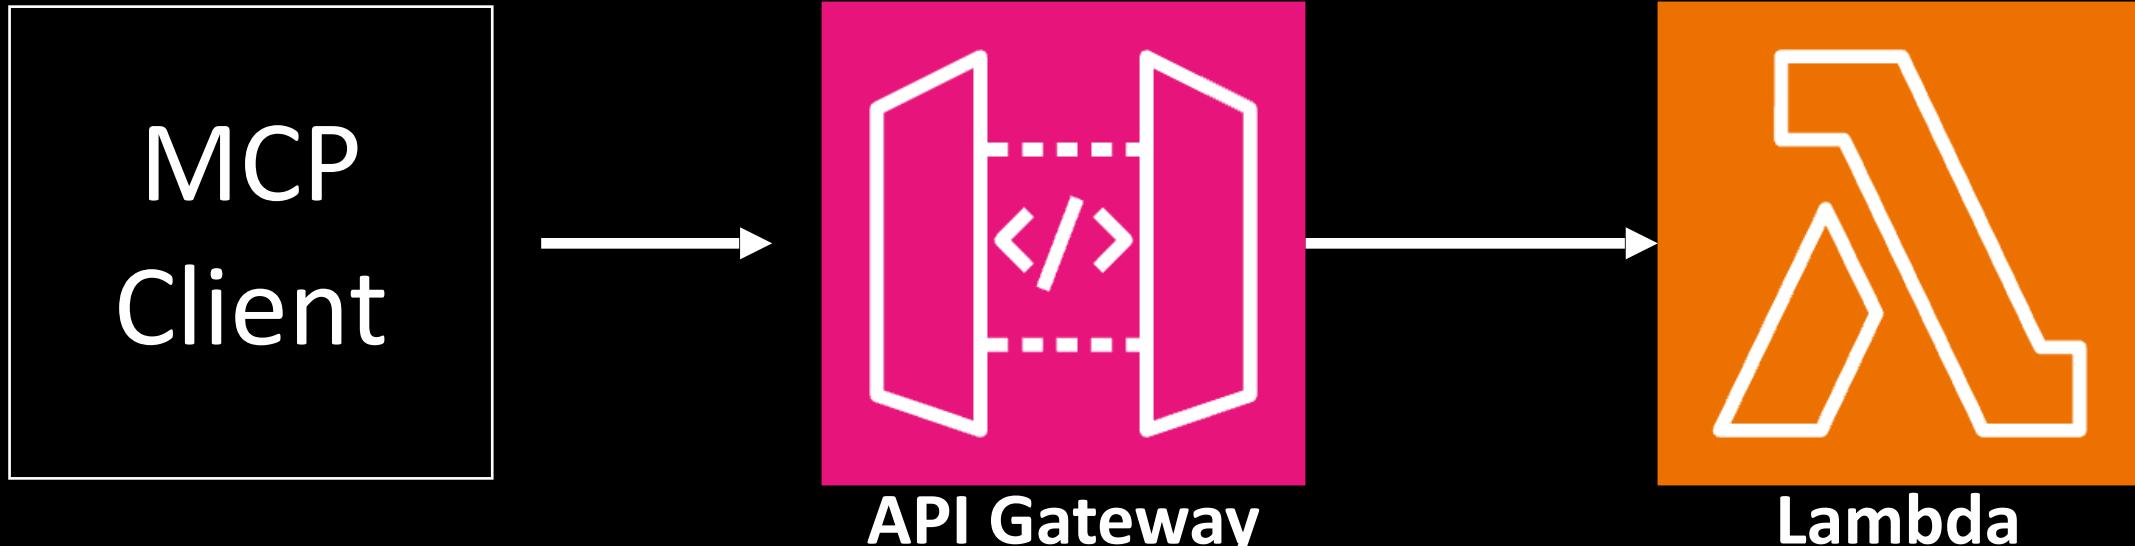


© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Building Serverless MCP Servers on AWS

```
@mcp.tool()  
def ping() -> str:  
    return "Pong!"
```

Building a stateless remote MCP Server on AWS Lambda



**Exposing MCP
functionality to clients**

**Running a serverless
MCP Server**

Building a stateless remote MCP Server on AWS Lambda

```
# Stateful server, SSE enabled
mcp = FastMCP("MyAwesomeServer")

# Stateless server, SSE enabled
mcp = FastMCP("MyAwesomeServer",
| | | | stateless_http=True)

# Stateless server, SSE disabled
mcp = FastMCP("MyAwesomeServer",
| | | | stateless_http=True,
| | | | json_response=True)
```

Define your MCP Server configuration

```
# Add a simple tool to demonstrate the server
@mcp.tool()
def ping() -> str:
|   return "Pong!"

# Run server with streamable_http transport
if __name__ == "__main__":
|   mcp.run(transport="streamable-http")
```



Building a stateless remote MCP Server on AWS Lambda

```
# Stateful server, SSE enabled
mcp = FastMCP("MyAwesomeServer")

# Stateless server, SSE enabled
mcp = FastMCP("MyAwesomeServer",
               stateless_http=True)

# Stateless server, SSE disabled
mcp = FastMCP("MyAwesomeServer",
               stateless_http=True,
               json_response=True)
```

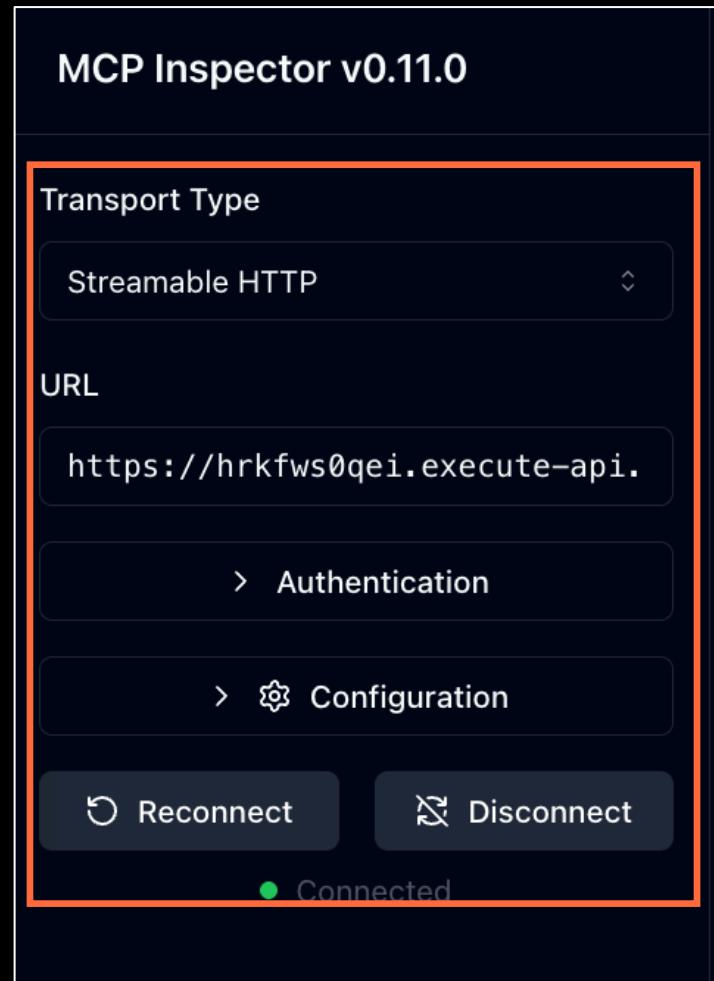
Define your MCP Server configuration

```
# Add a simple tool to demonstrate the server
@mcp.tool()
def ping() -> str:
    return "Pong!"

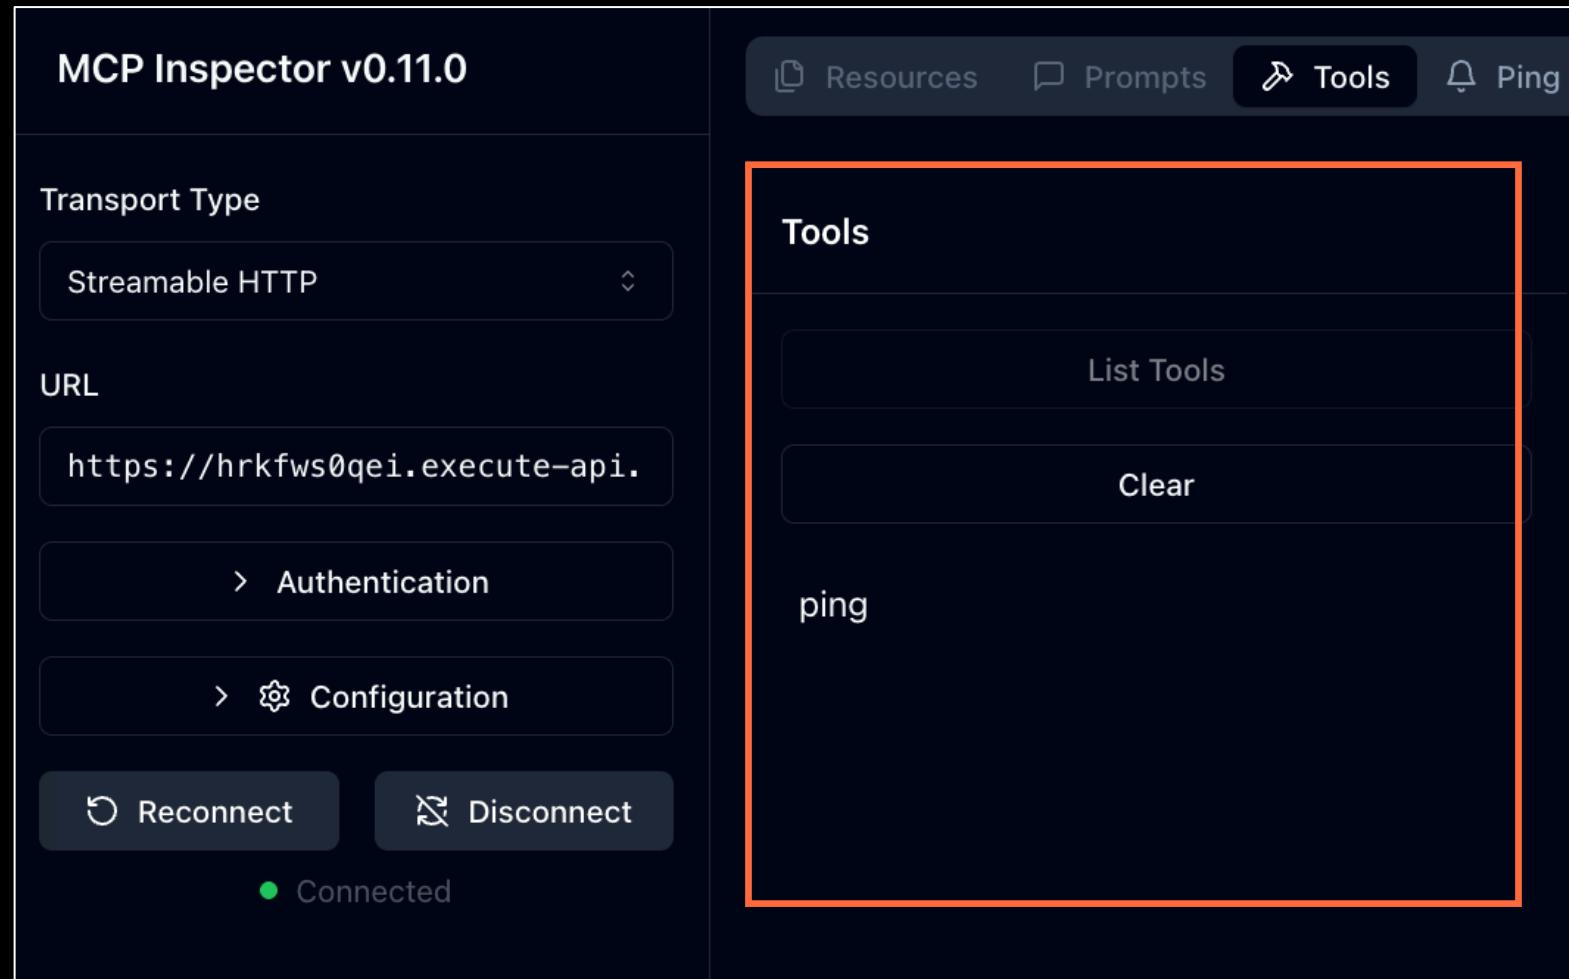
# Run server with streamable_http transport
if __name__ == "__main__":
    mcp.run(transport="streamable-http")
```



Building a stateless remote MCP Server on AWS Lambda



Building a stateless remote MCP Server on AWS Lambda



Building a stateless remote MCP Server on AWS Lambda

The screenshot shows the MCP Inspector v0.11.0 application window. On the left, the main interface includes fields for Transport Type (Streamable HTTP), URL (<https://hrkfws0qei.execute-api.>), and sections for Authentication and Configuration. At the bottom, there are Reconnect and Disconnect buttons, and a status indicator showing Connected with a green dot. On the right, a navigation bar at the top has tabs for Resources, Prompts, Tools (selected), Ping, Sampling, and Roots. Below the navigation bar, the 'Tools' section lists 'ping' and 'List Tools'. A red box highlights the 'ping' entry, which has a 'Run Tool' button. To the right of the 'ping' entry, a box displays the 'Tool Result: Success' message: "pong! Log Stream Name=2025/05/01/ [LATEST]f013ae0ff6f74f74b6691c37d59d58cf".

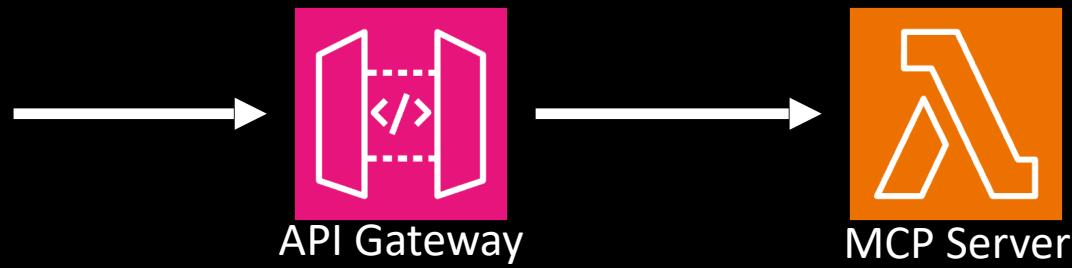
Building a stateless remote MCP Server on AWS

But what about observability?
Security? Governance? CI/CD?

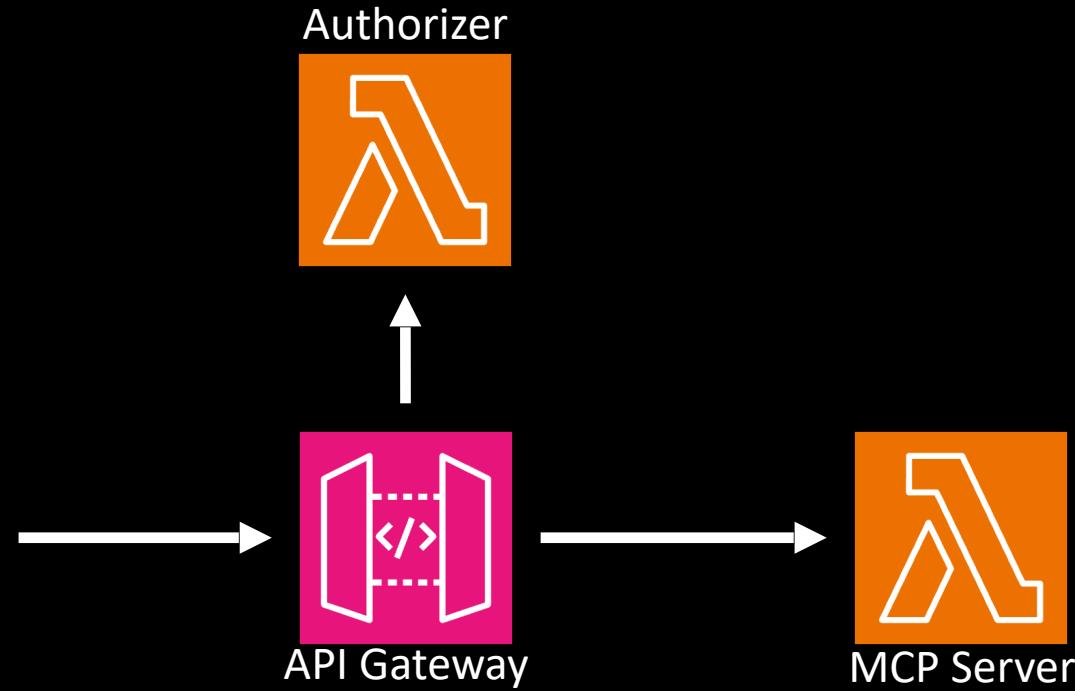
**Everything you know about building
serverless applications and APIs on
AWS is still applicable!**

(and will further improve as MCP spec evolves)

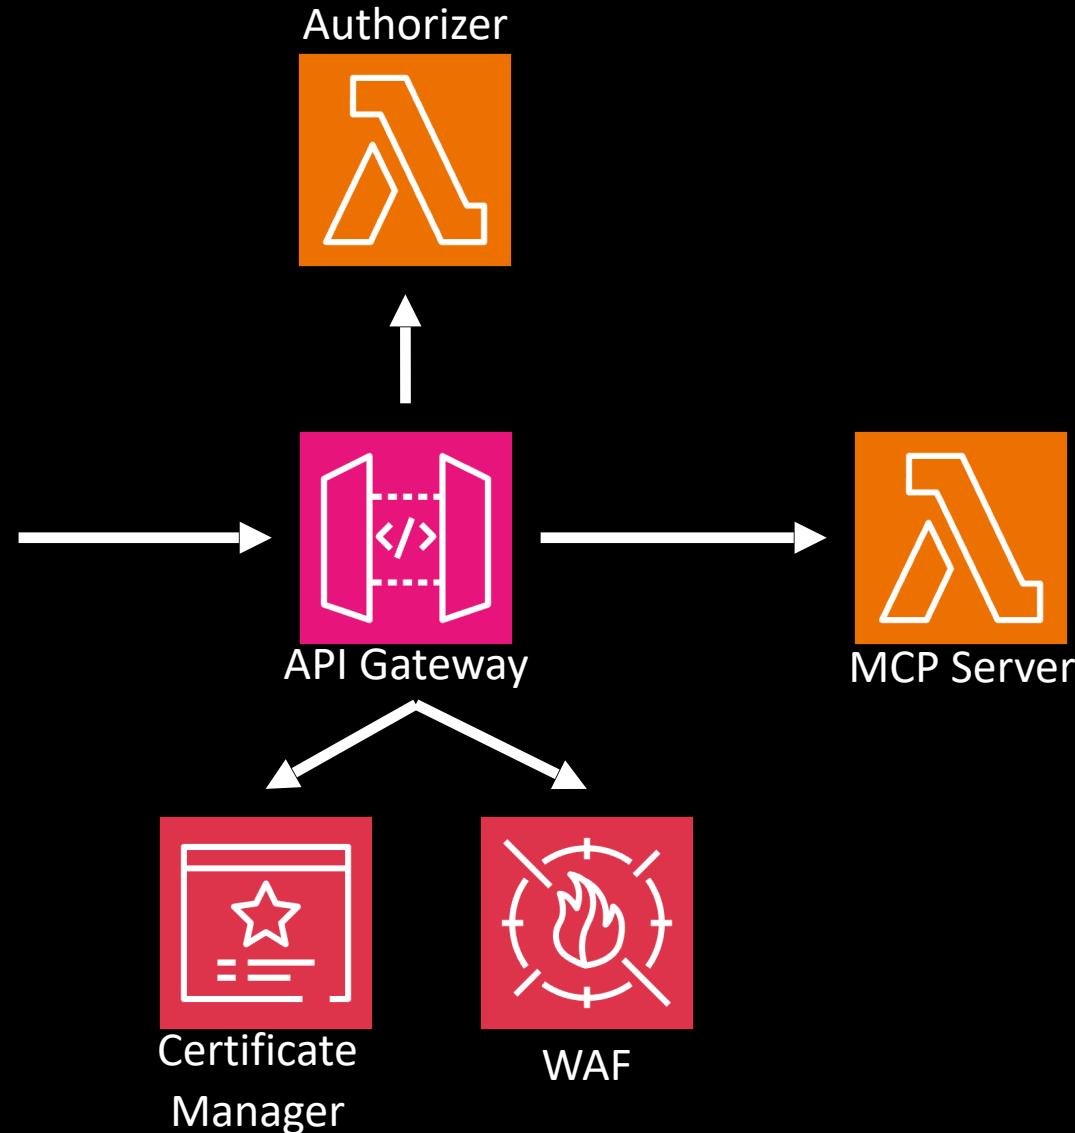
Building a stateless remote MCP Server on AWS Lambda



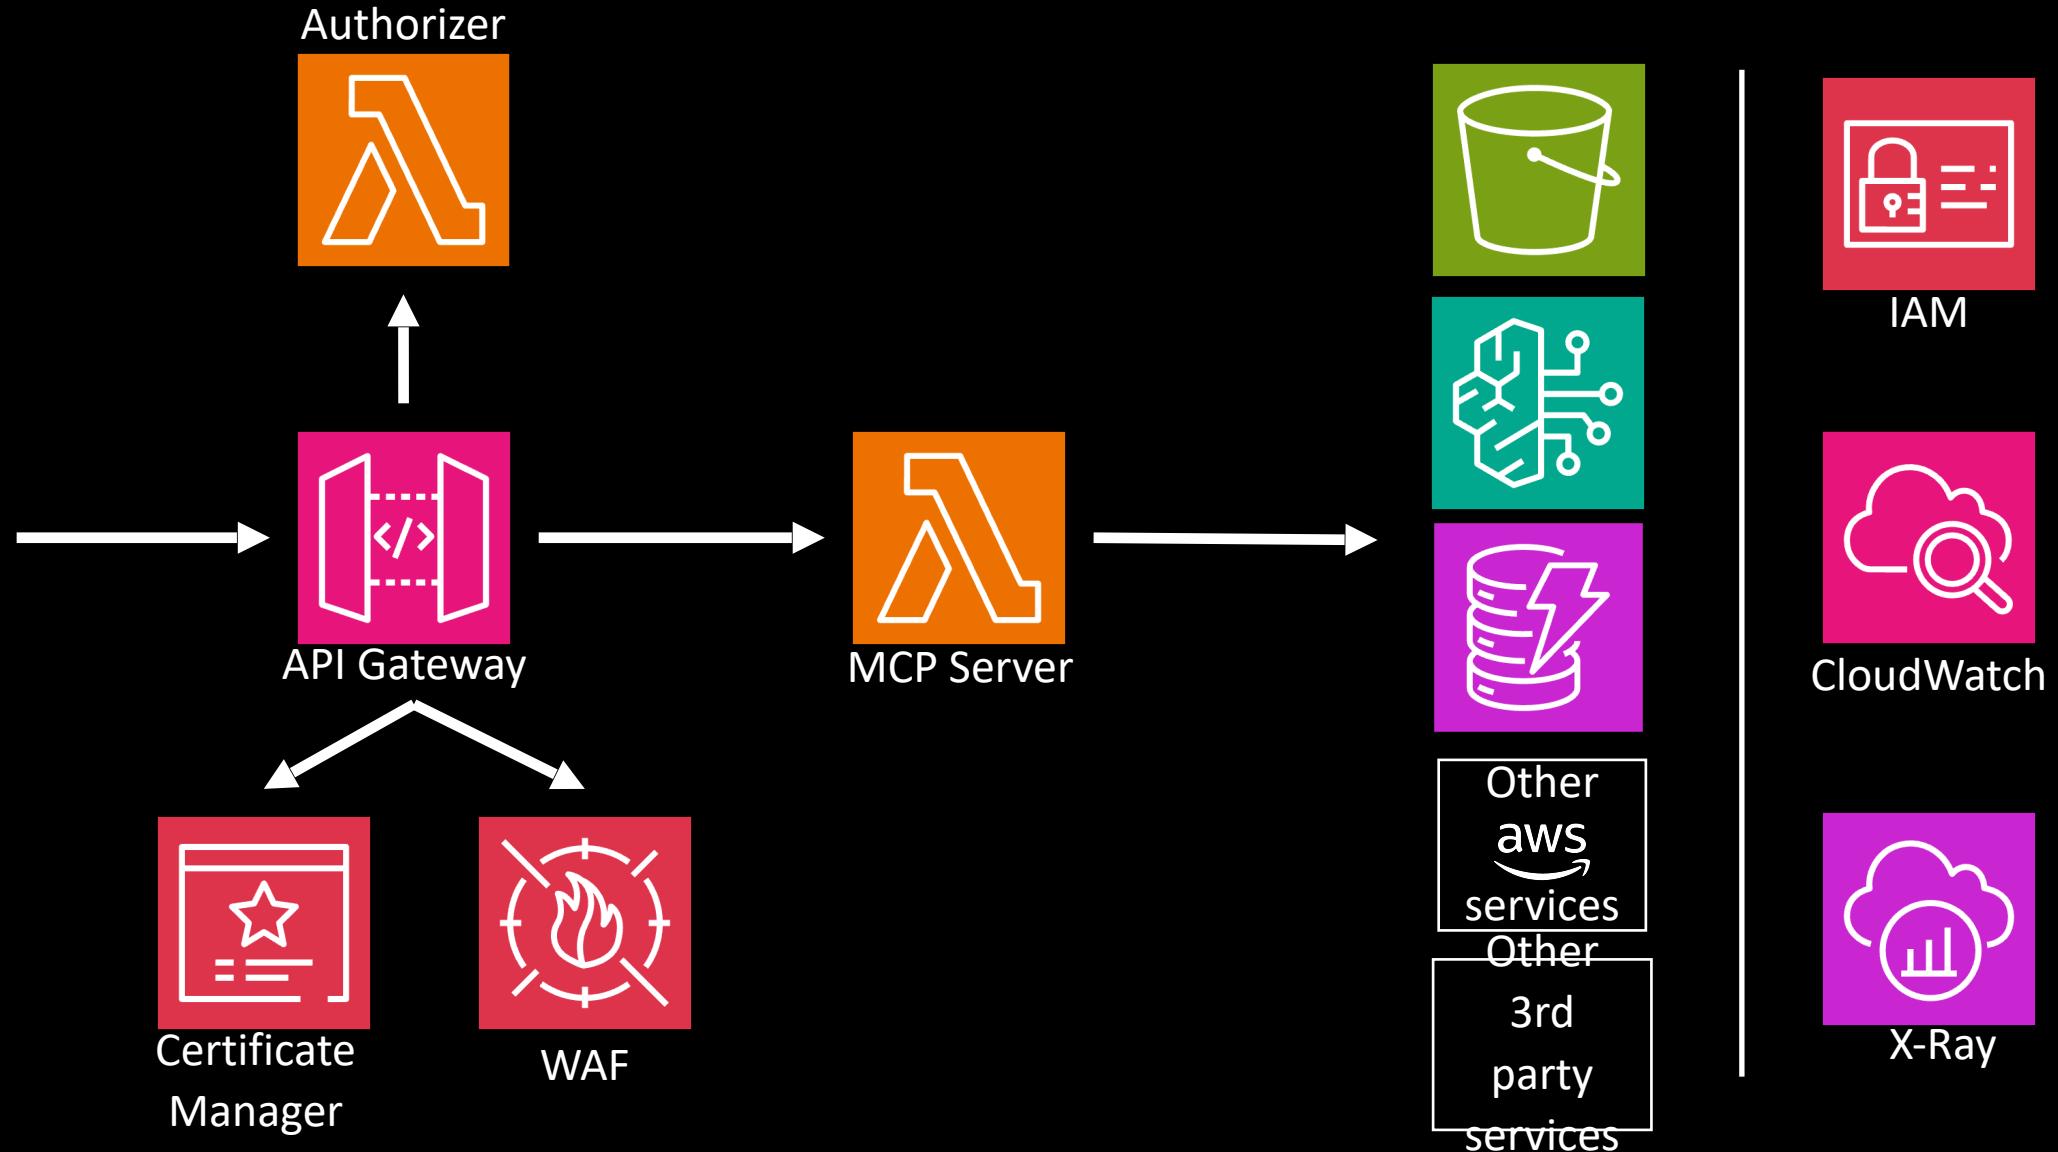
Building a stateless remote MCP Server on AWS Lambda



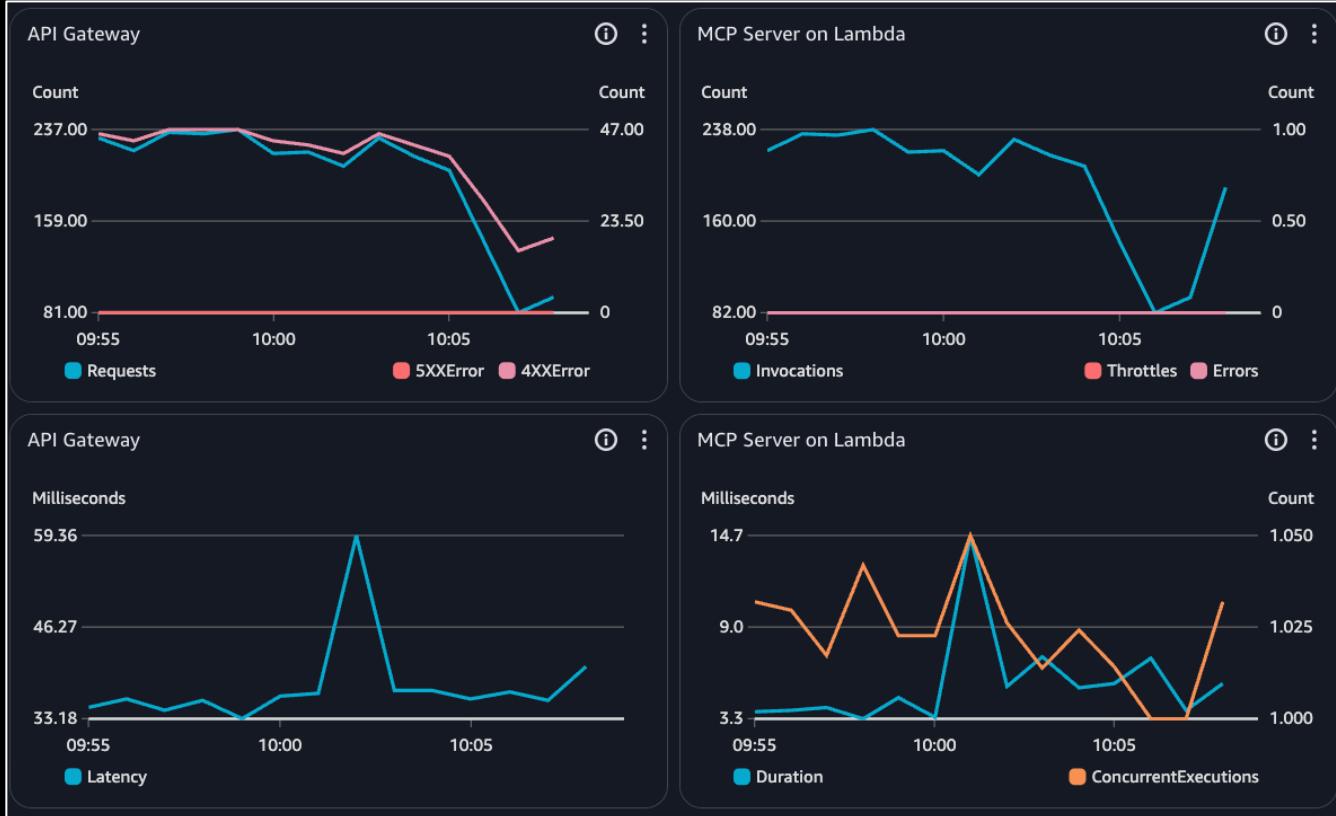
Building a stateless remote MCP Server on AWS Lambda



Building a stateless remote MCP Server on AWS Lambda



Building a stateless remote MCP Server on AWS Lambda



```
▶ ⚒ 2025-05-02T10:13:45.444-05:00 START RequestId: 42f25ab8-ed1c-4da6-96ef-e1e36c9fb33b Version: $LATEST
▶ ⚒ 2025-05-02T10:13:45.446-05:00 DEBUG [transport-http-no-session.js:27:requestLogger] > POST /mcp
▶ ⚒ 2025-05-02T10:13:45.446-05:00 DEBUG [transport-http-no-session.js:29:requestLogger] { method: 'notifications/initialized', jsonrpc: '2.0' }
▶ ⚒ 2025-05-02T10:13:45.450-05:00 END RequestId: 42f25ab8-ed1c-4da6-96ef-e1e36c9fb33b
```

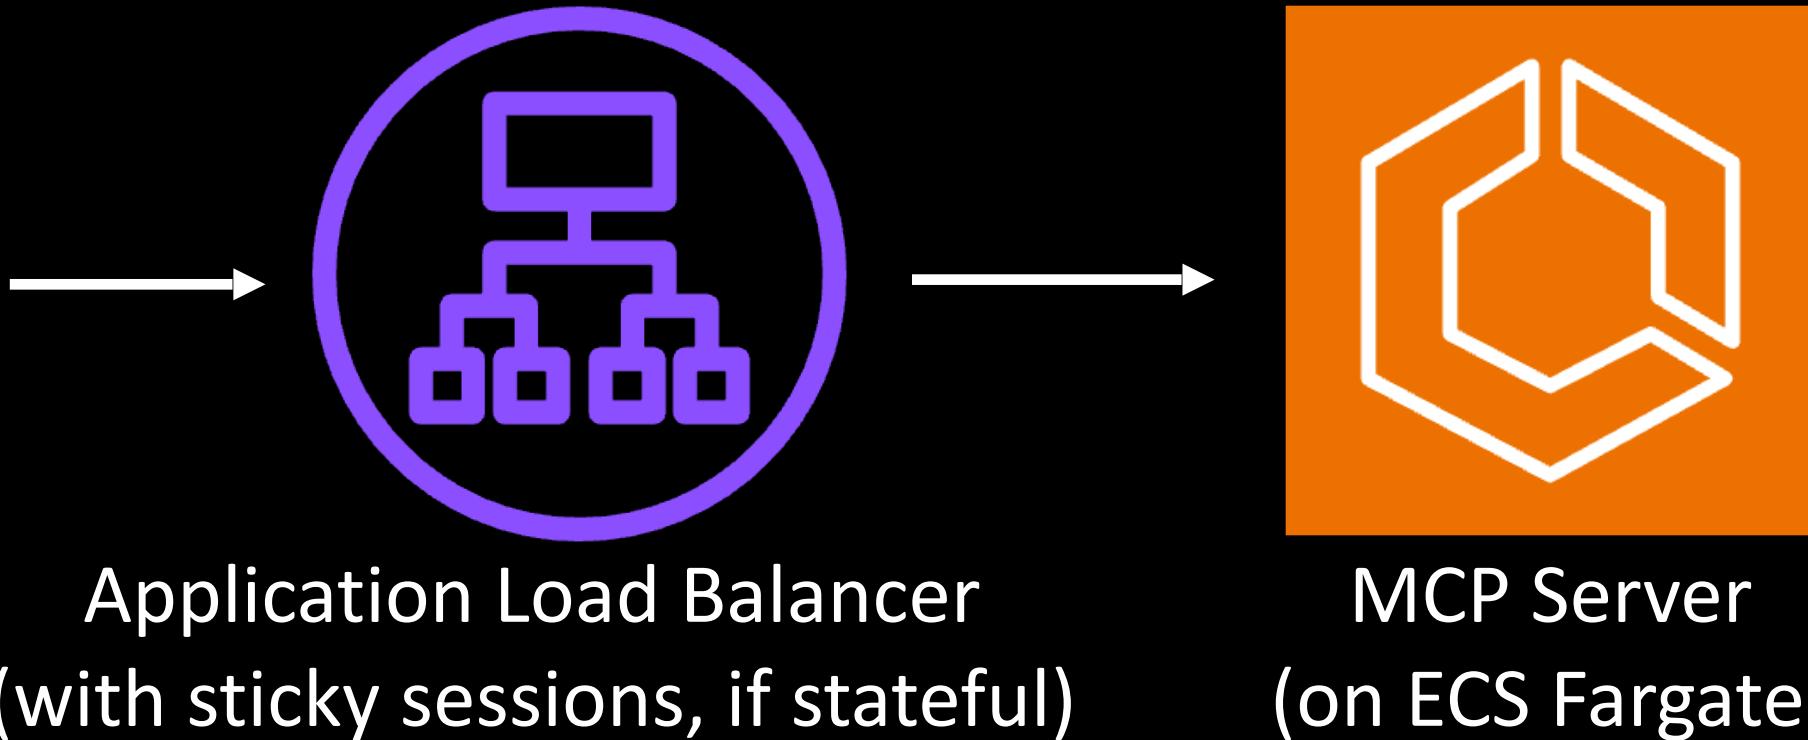
Stateful MCP Server

with Amazon ECS

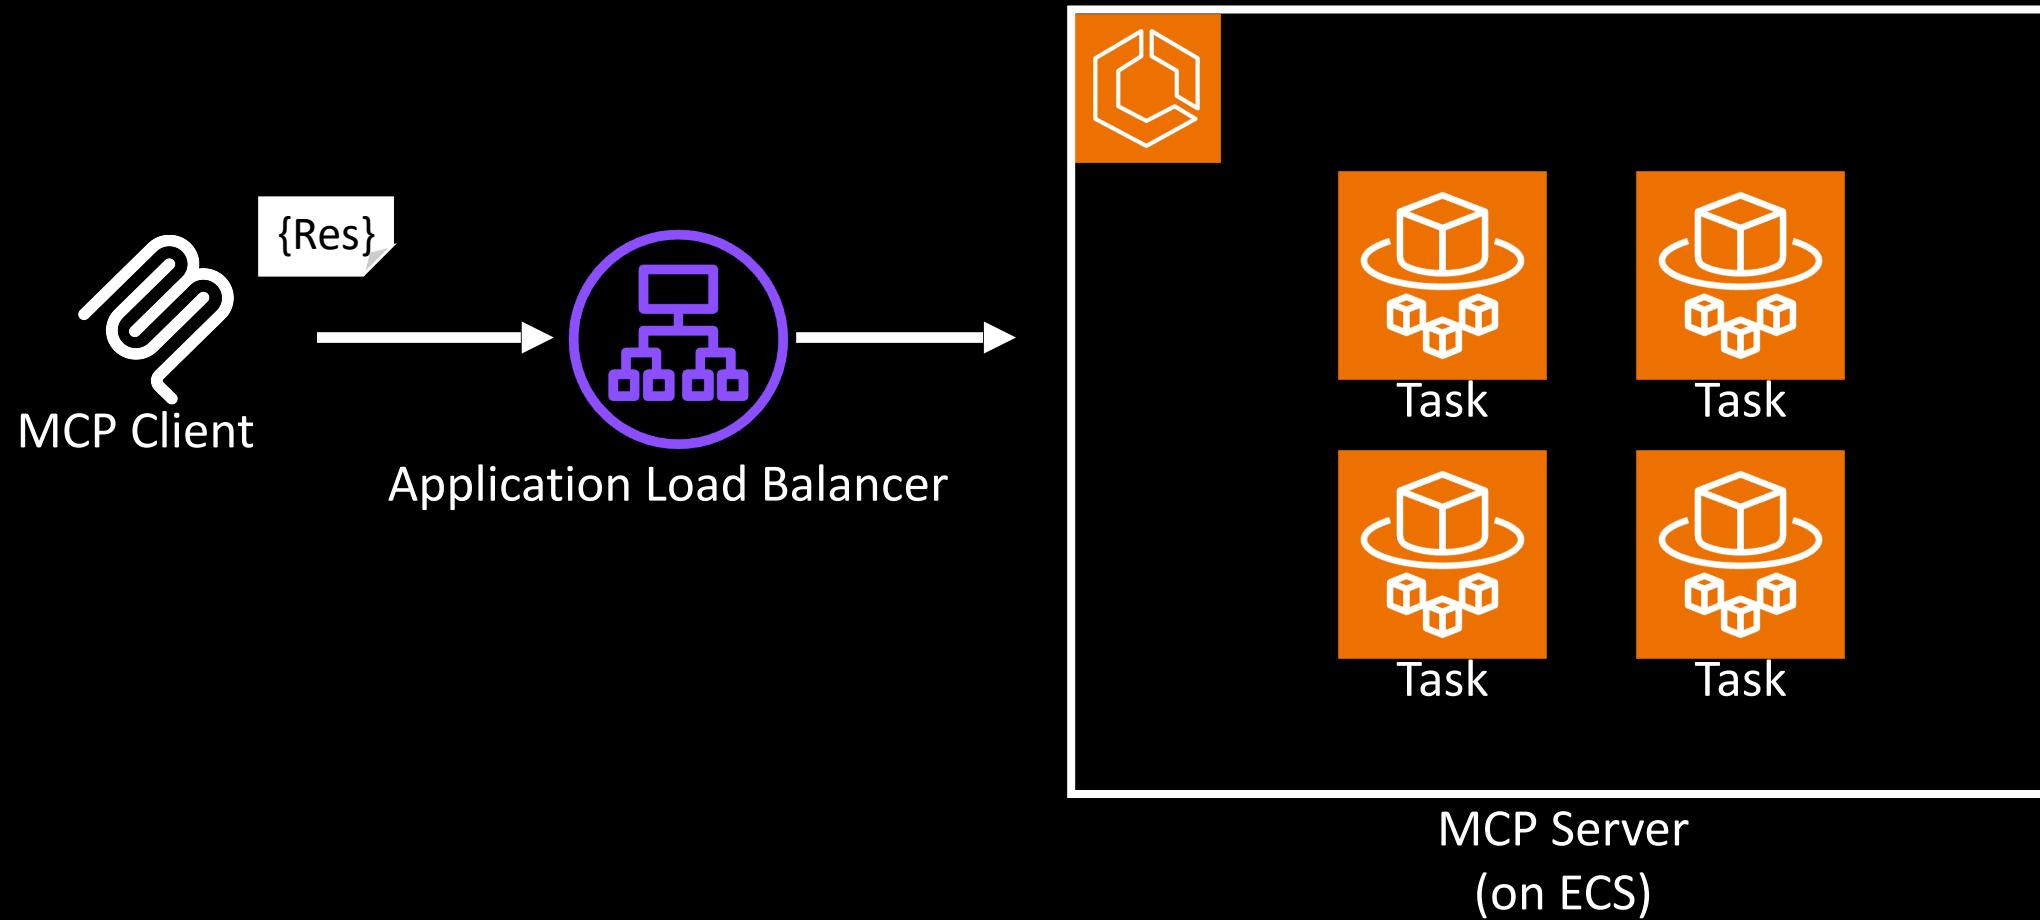


© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

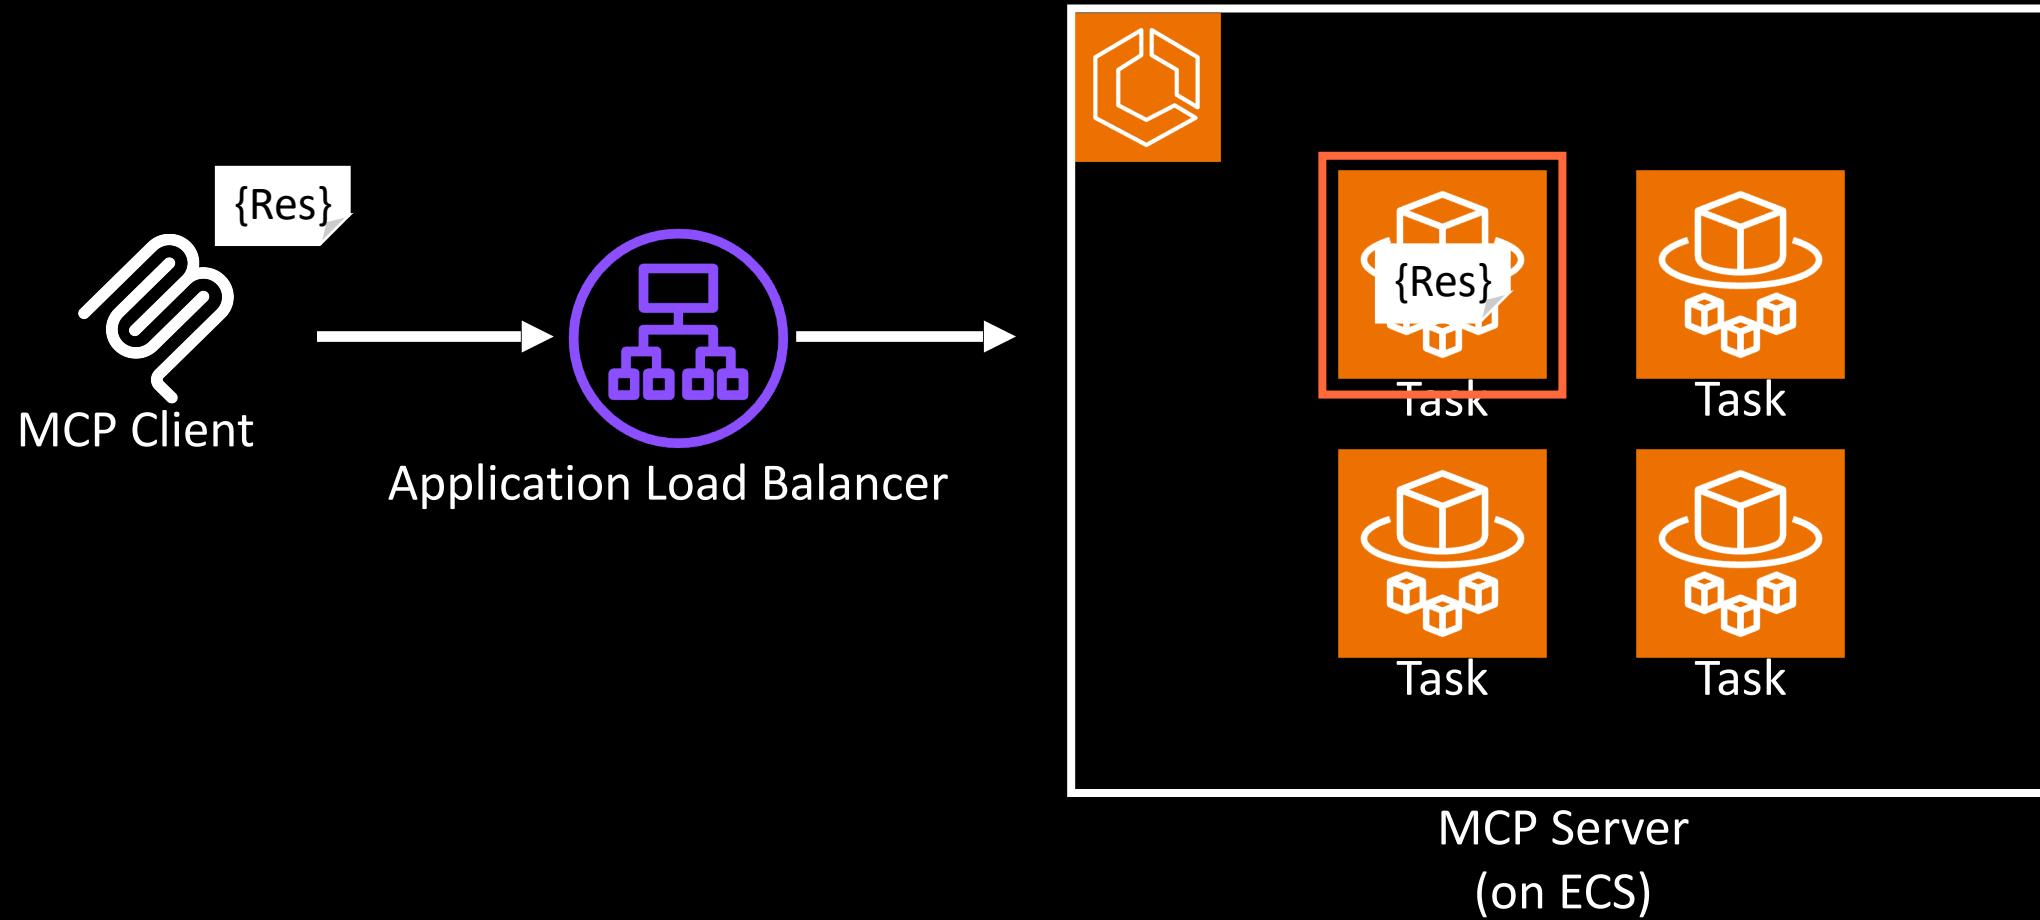
Building a stateful remote MCP Server on ECS



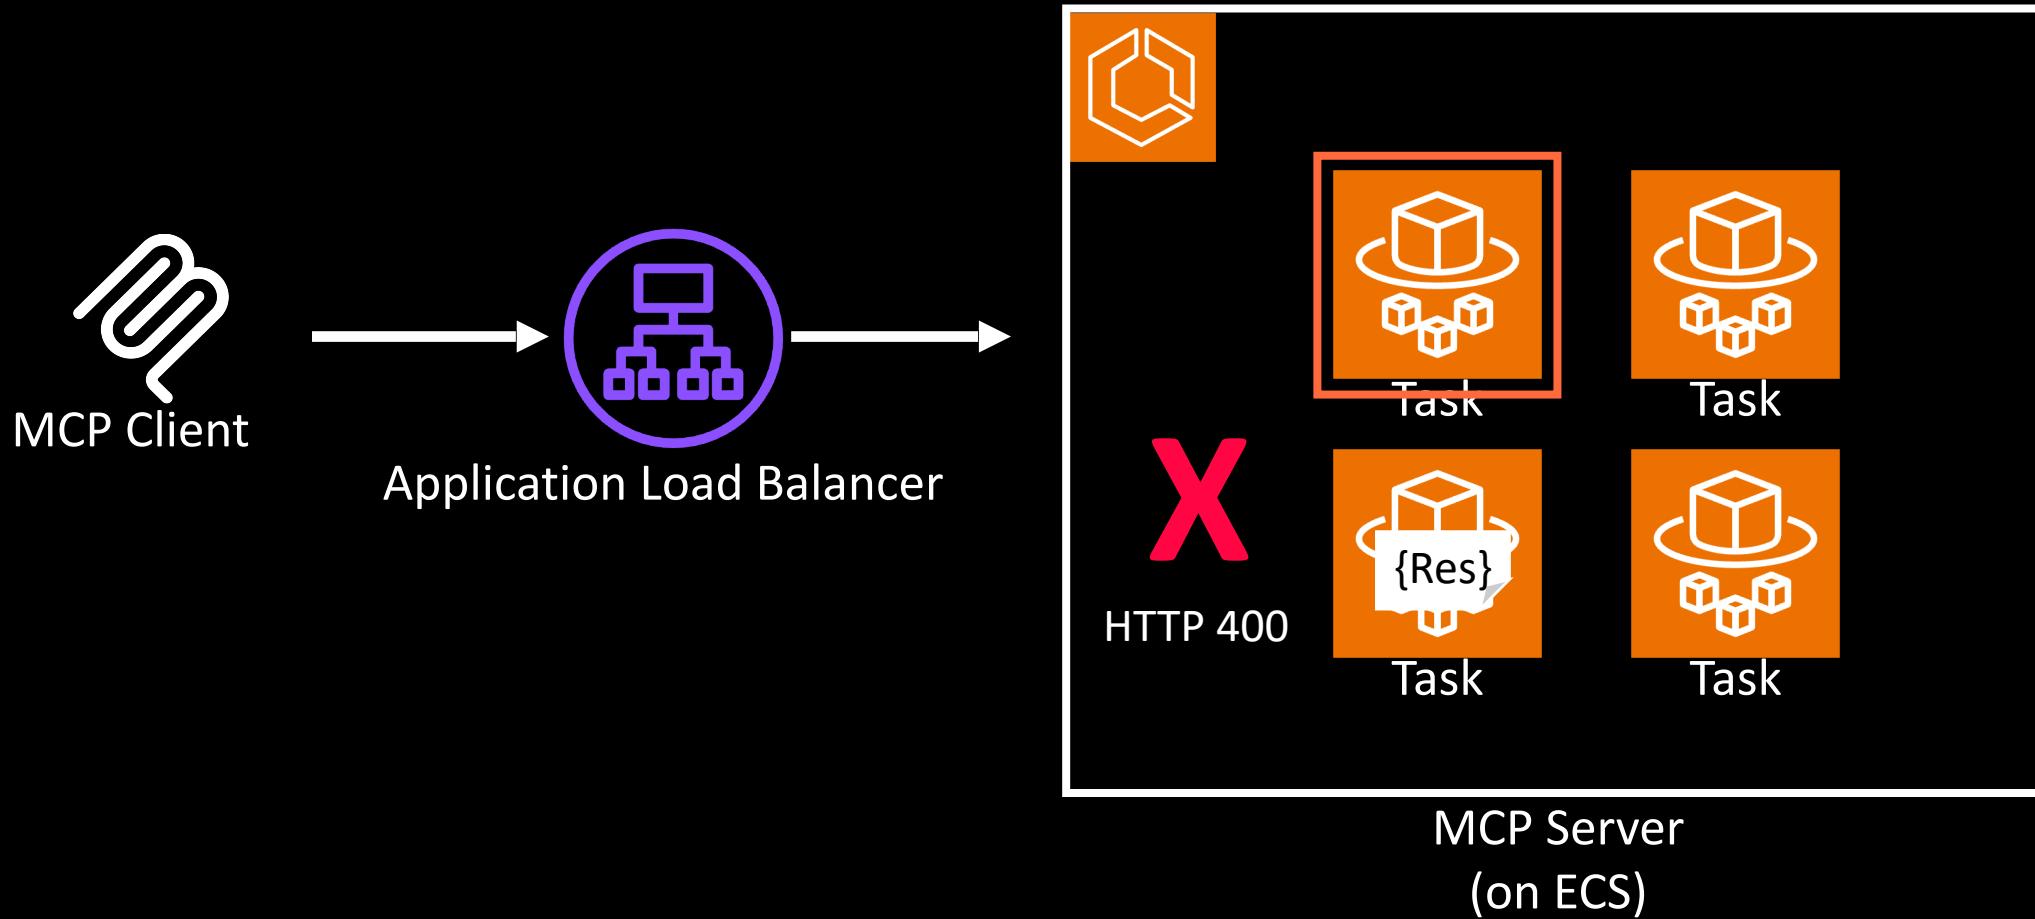
Building a stateful remote MCP Server on ECS



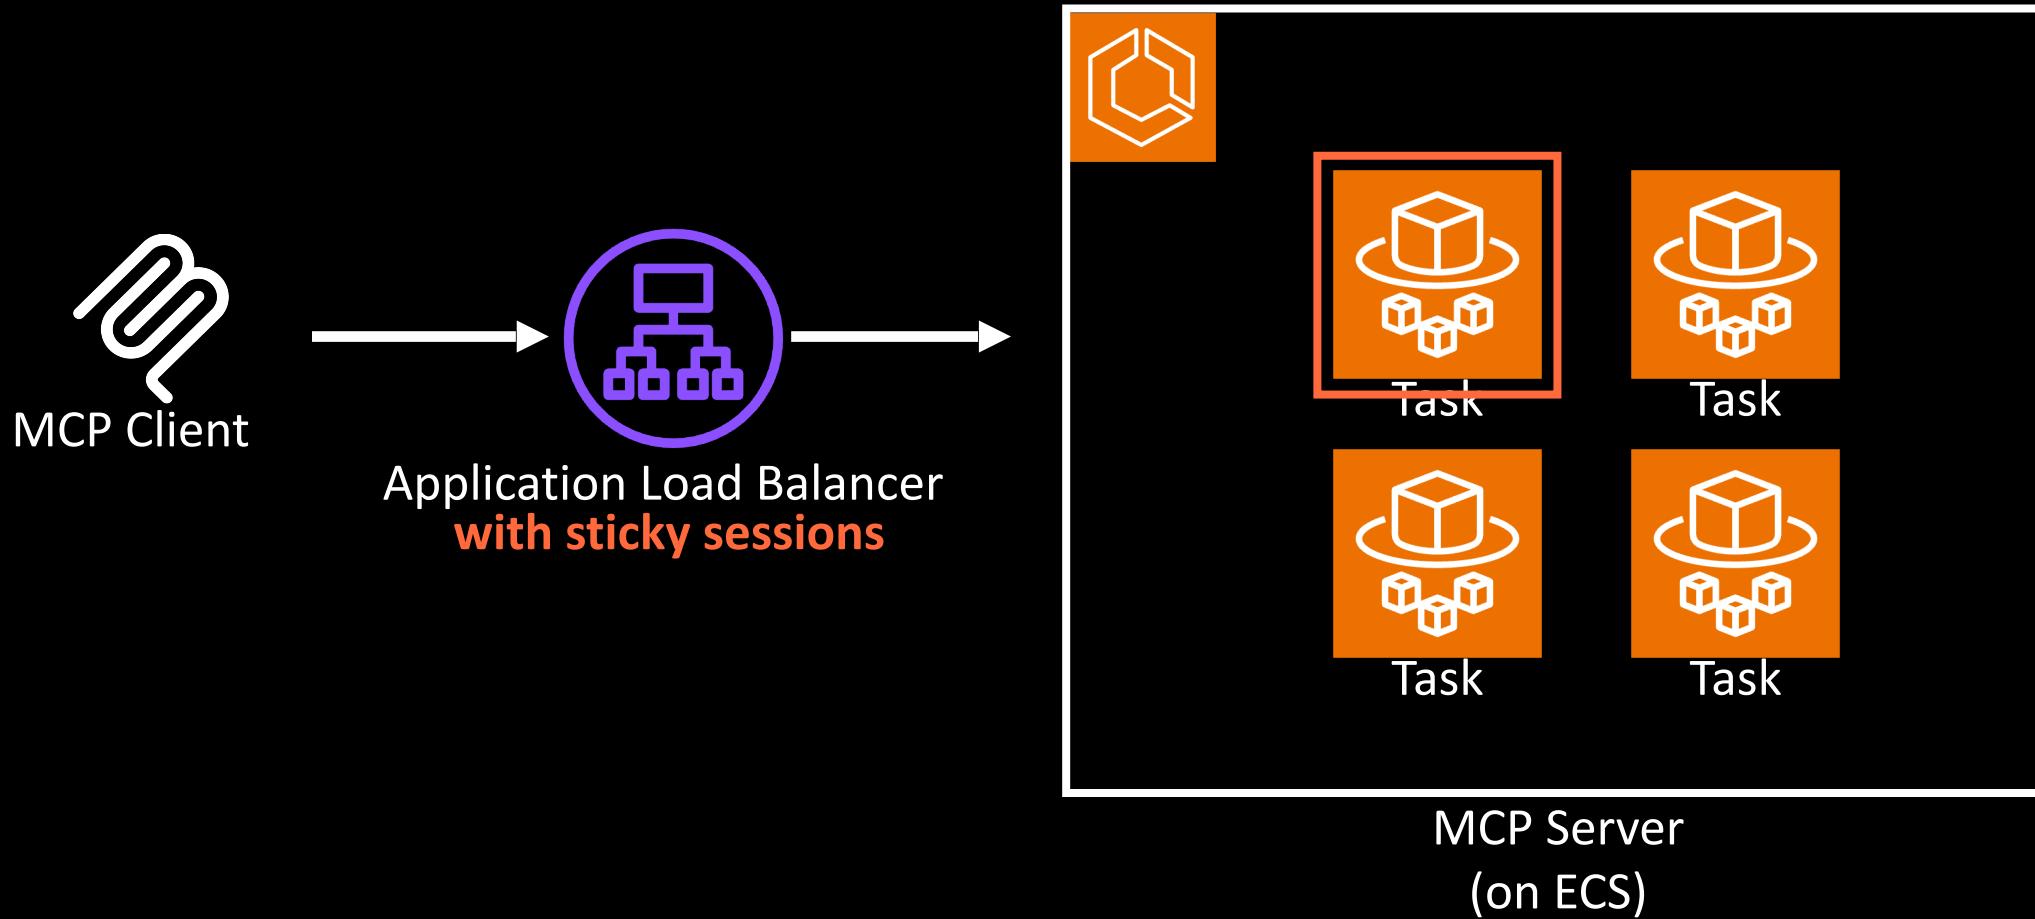
Building a stateful remote MCP Server on ECS



Building a stateful remote MCP Server on ECS



Building a stateful remote MCP Server on ECS



Building a stateful remote MCP Server on ECS

```
resource "aws_lb_target_group" "mcp_server" {
    name          = local.project_name
    port          = local.ecs_task_container_port
    protocol     = "HTTP"
    vpc_id        = aws_vpc.main.id
    target_type   = "ip"
    deregistration_delay = 60

    stickiness {
        enabled      = true
        type         = "lb_cookie"
        cookie_duration = 86400 # 1 day
    }
}
```

Enable sticky sessions



Building a stateful remote MCP Server on ECS

```
# Stateful server, SSE enabled  
mcp = FastMCP("MyAwesomeServer")
```

```
# Stateless server, SSE enabled  
mcp = FastMCP("MyAwesomeServer",  
| | | | stateless_http=True)
```

```
# Stateless server, SSE disabled  
mcp = FastMCP("MyAwesomeServer",  
| | | | stateless_http=True,  
| | | | json_response=True)
```

```
# Add a simple tool to demonstrate the server  
@mcp.tool()  
def ping() -> str:  
| return "Pong!"
```

```
# Run server with streamable_http transport  
if __name__ == "__main__":  
| mcp.run(transport="streamable-http")
```

Define your MCP Server configuration



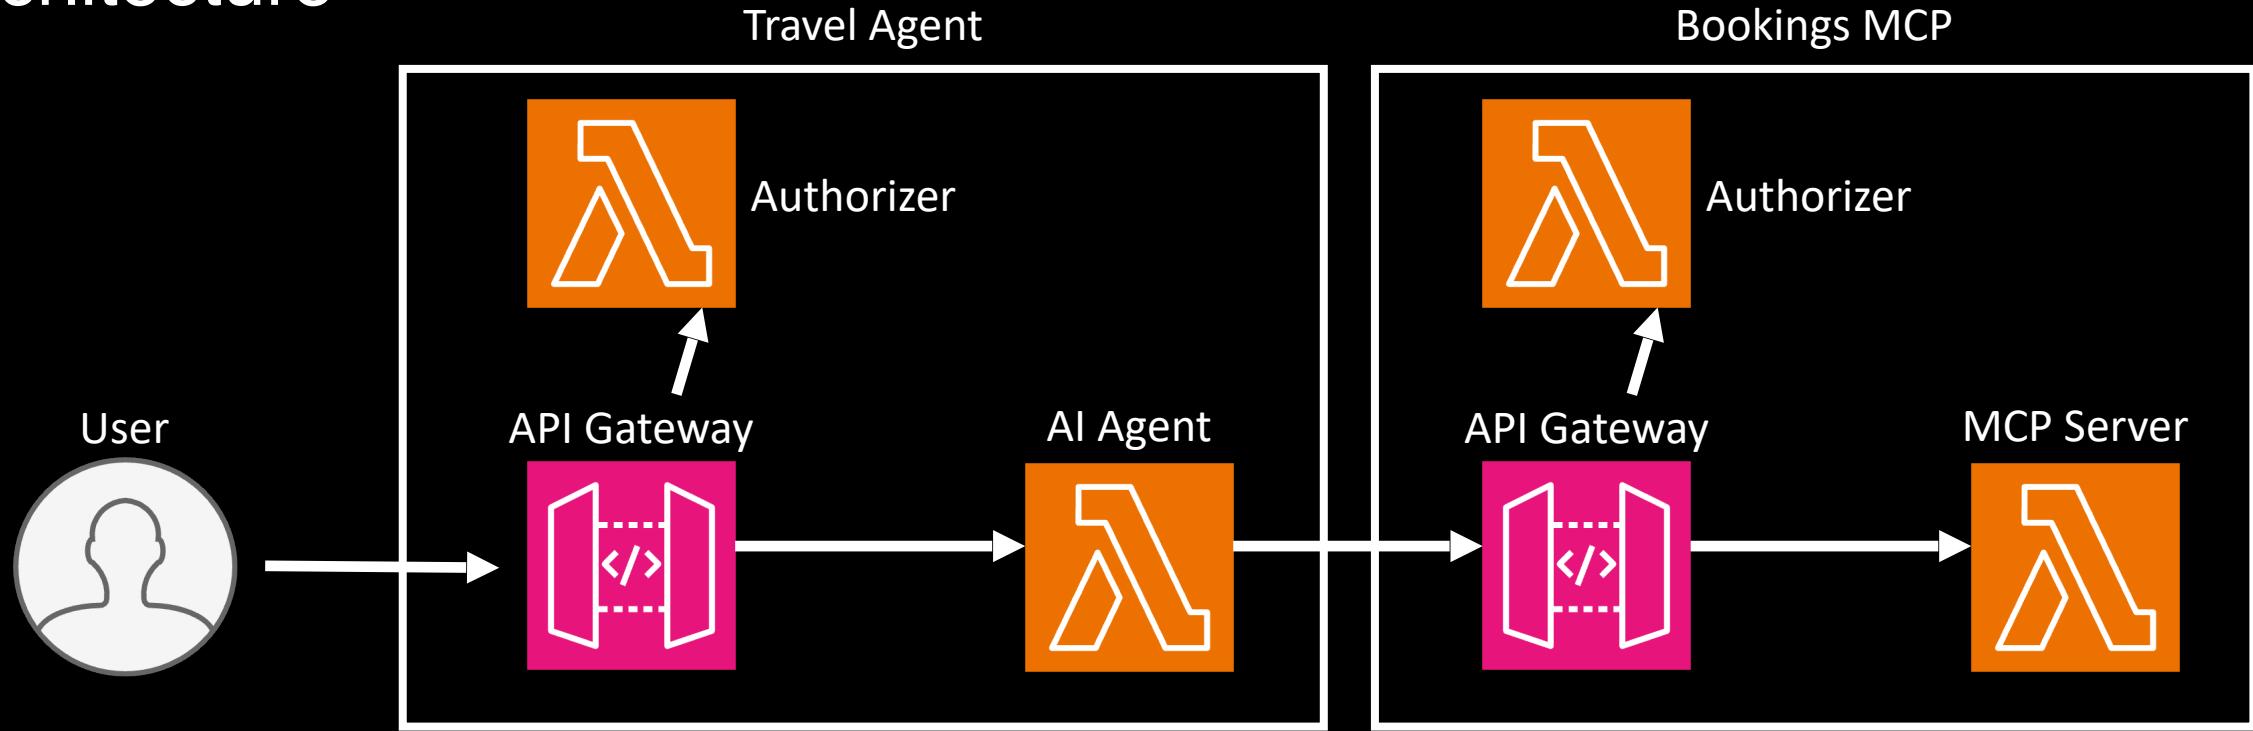
Let's see some demos

Building MCP Servers on AWS

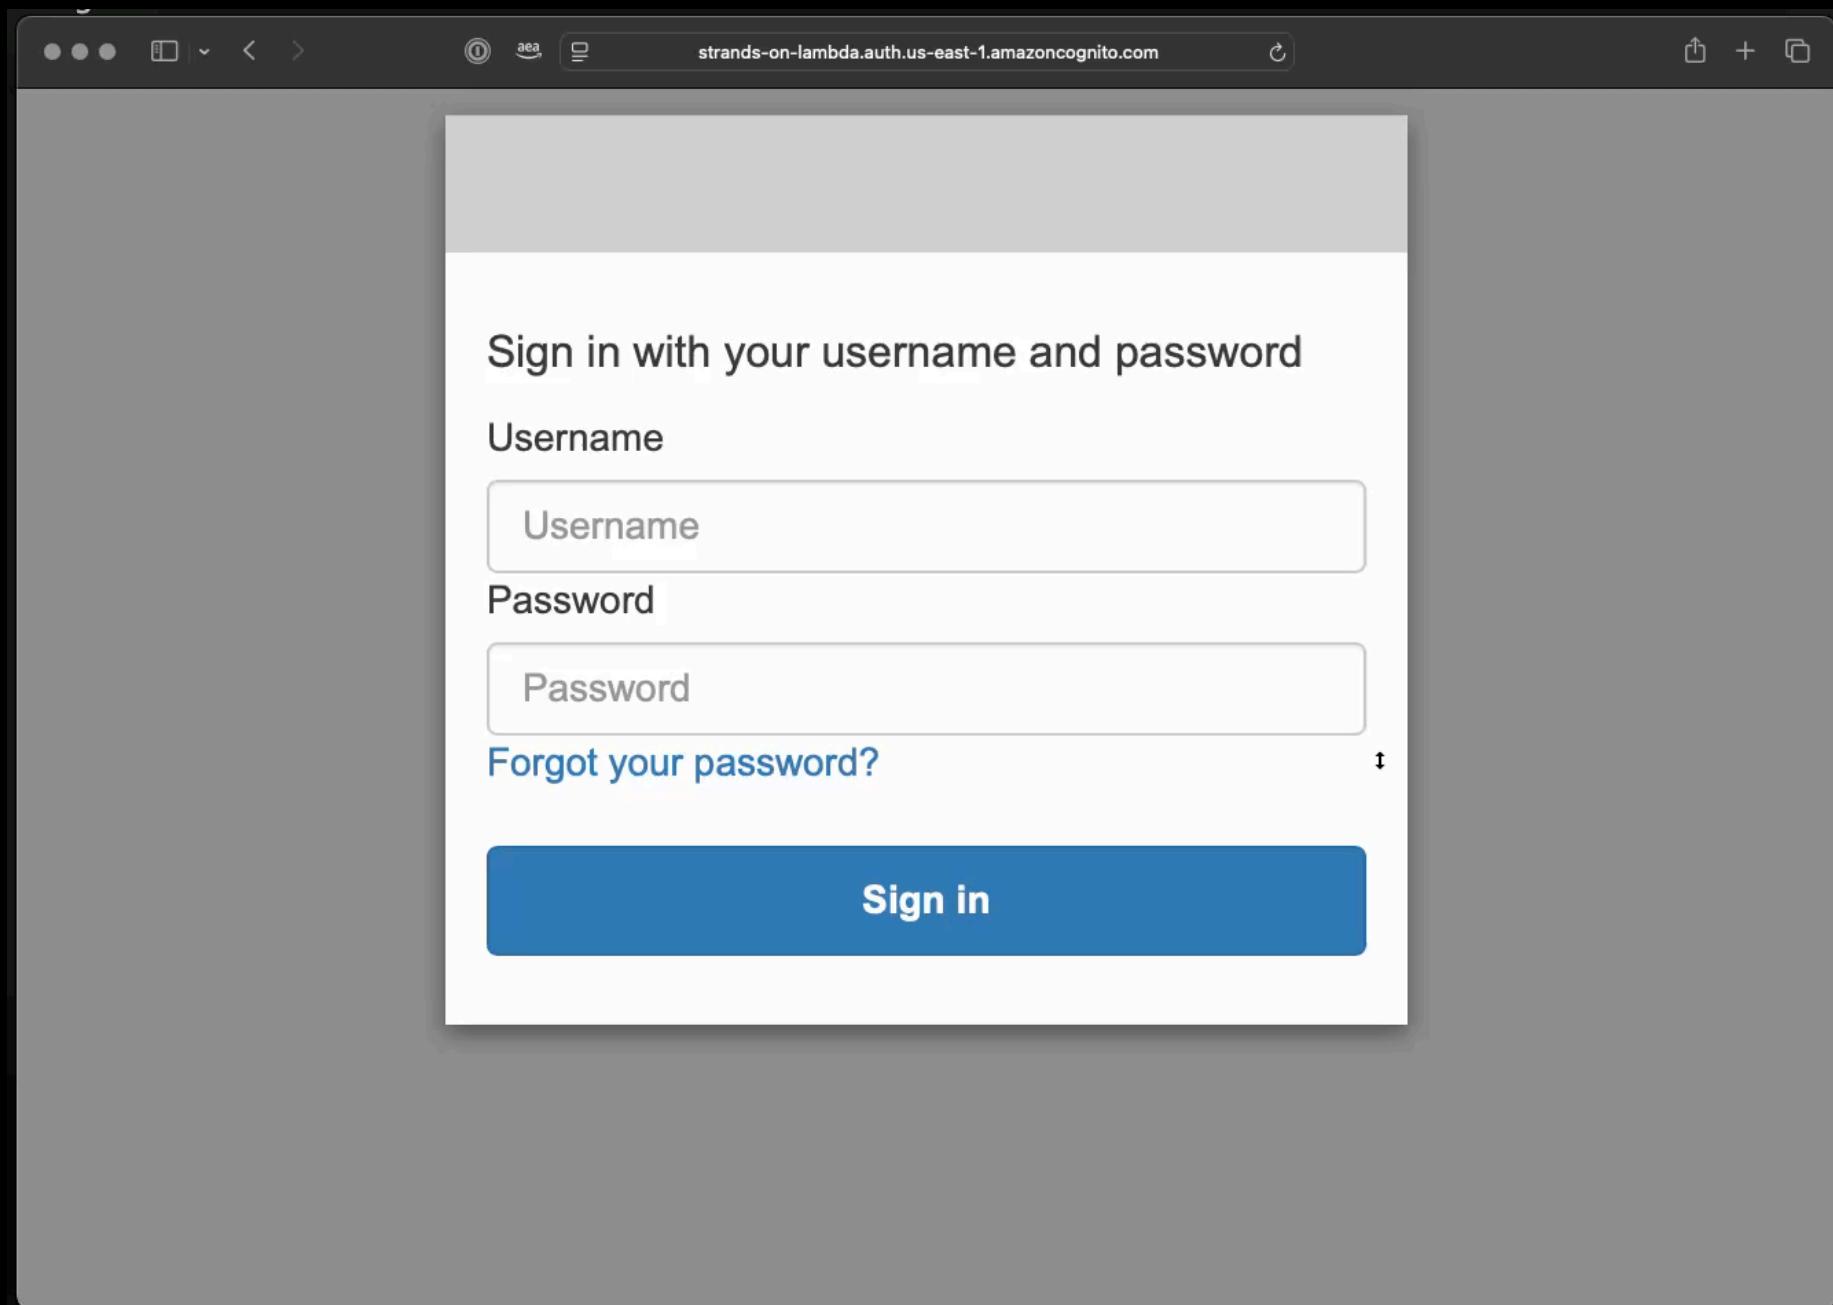


© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Architecture



- Get Travel Policies
- Get Available Cars
- Get Available Hotels
- Book Car
- Book Hotel
-



✳️ Hi Max, how are you?

How can I help you today?



Research

Claude Sonnet 4 ▾



Write

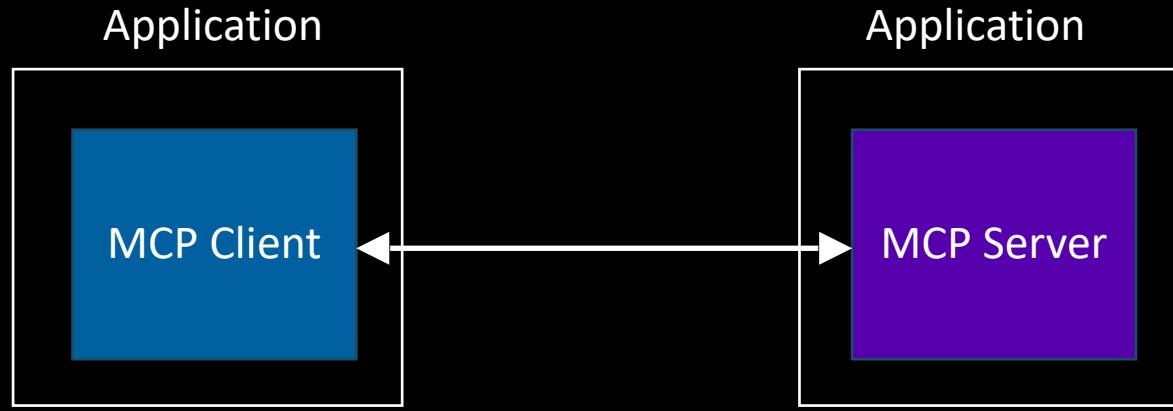
Learn

Code

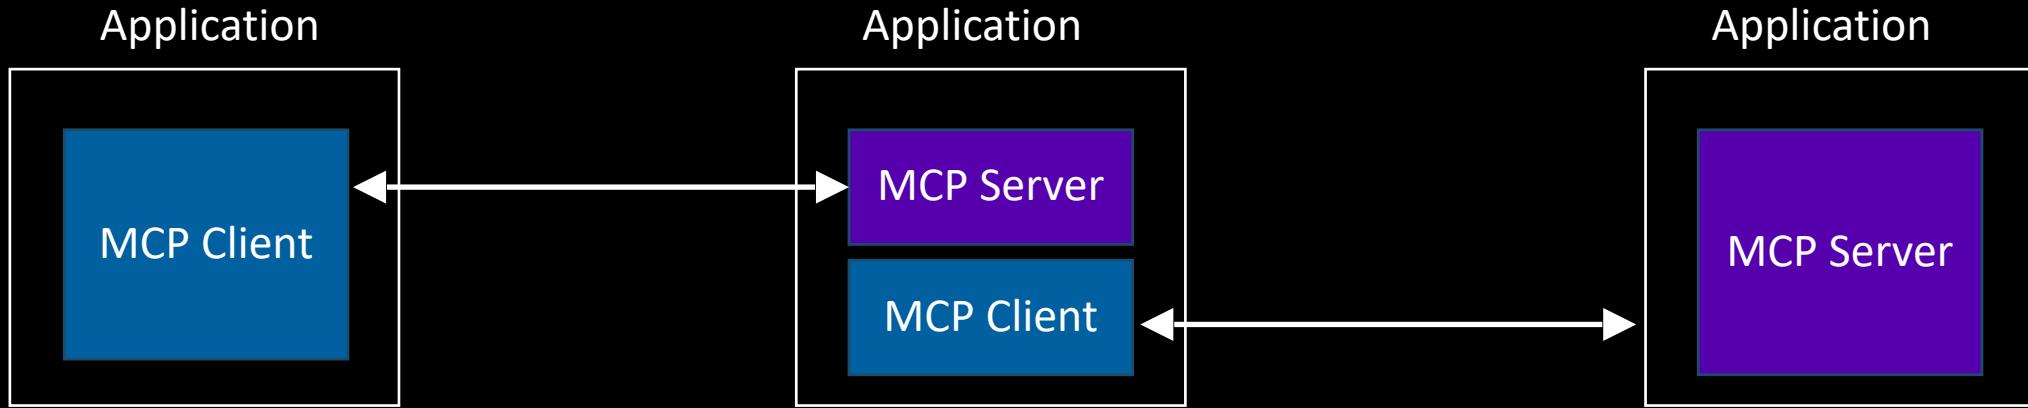
Life stuff

Advanced patterns

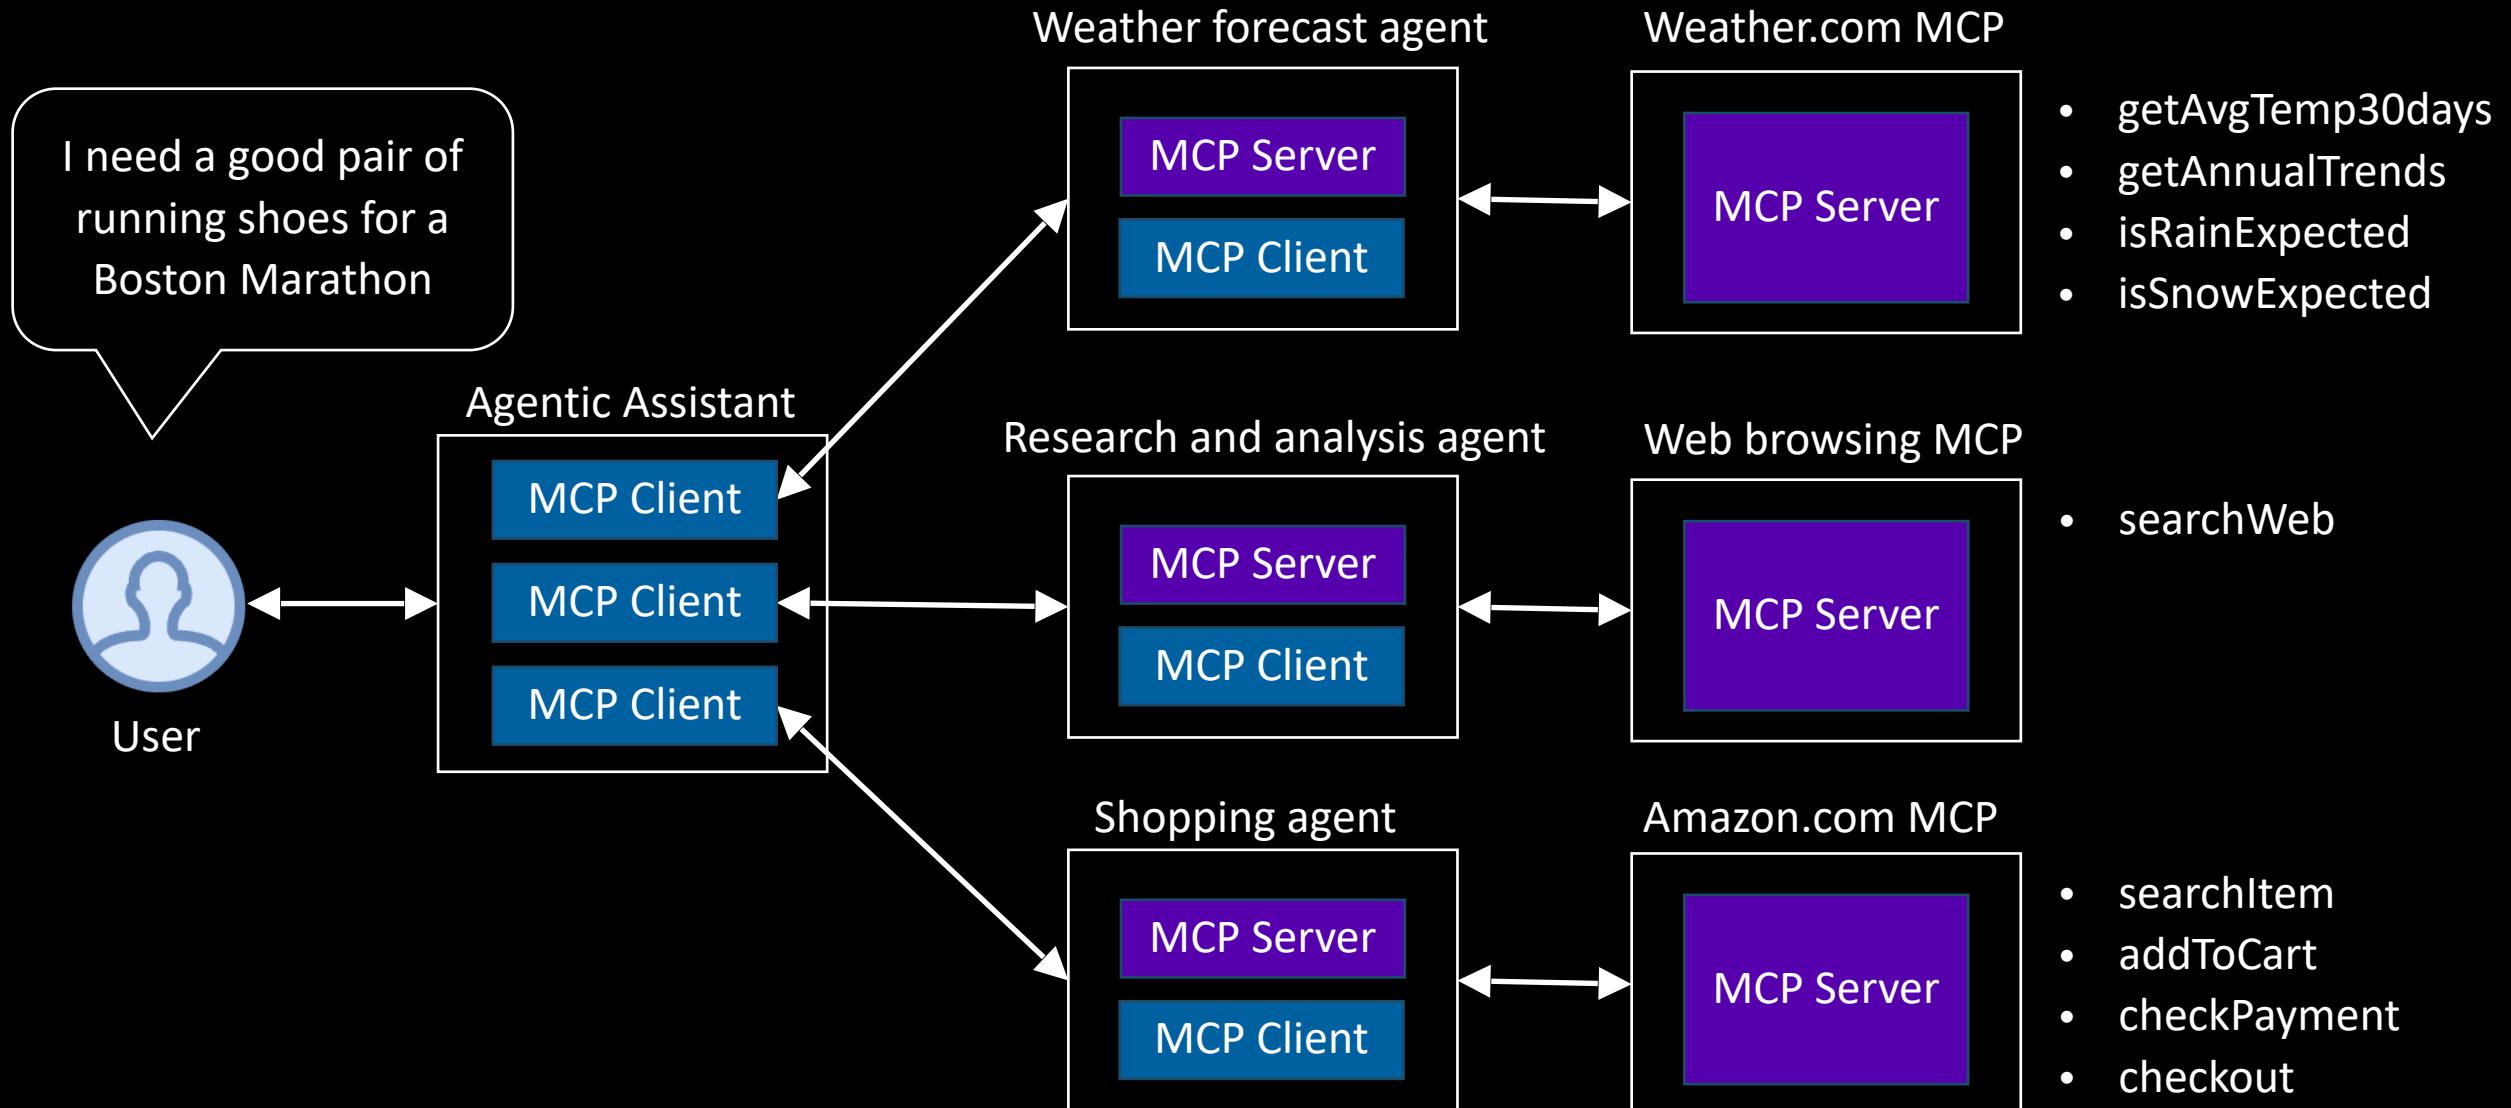
Composability



Composability



Agent-to-Agent communications



In conclusion

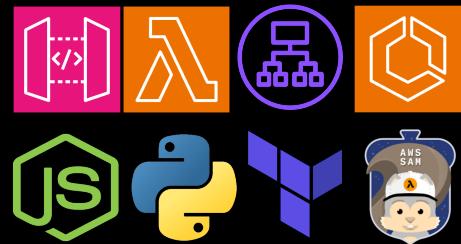
Building MCP Servers on AWS



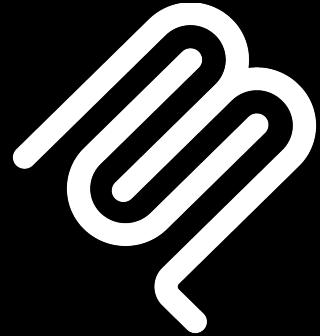
Serverless MCP Servers on Lambda and ECS



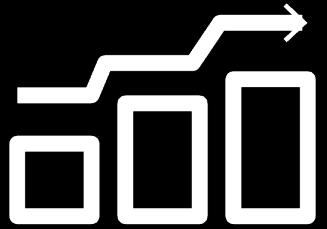
<https://github.com/aws-samples/sample-serverless-mcp-servers>



Best practices



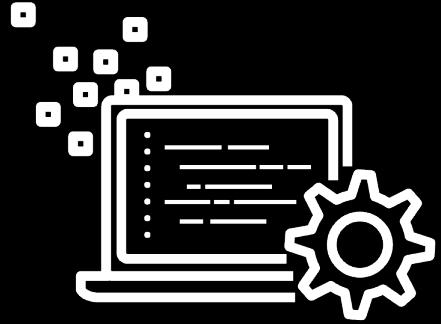
Use official MCP SDKs
(and follow the MCP
spec evolution)



Design for scalability
and performance



Implement robust
network and application
security



Establish
comprehensive
monitoring and logging

Continue your Serverless Gen AI / MCP journey



- [Sample Serverless MCP Servers](#)
- [Resources for Generative AI and Serverless](#)
- [Deploy MCP servers in AgentCore Runtime](#)
- CDK Construct for Serverless MCP Servers ([1](#), [2](#), [3](#))
- [AWS Labs MCP Servers](#) (20+ MCP Servers!!!)
- [Introducing AWS MCP Servers for code assistants](#)
- [Extend the Amazon Q Developer CLI with MCP](#)
- [Running MCP-Based Agents on AWS](#)



- [Latest MCP Spec](#)
- [Getting started with MCP](#)
- [MCP Registry](#)
- [A list of maintained MCP Servers](#)



Thank you!



Anton Aleksandrov
Principal Solutions Architect
AWS Serverless